




Advanced CV methods

Tracking patches

Matej Kristan

Laboratorij za Umetne Vizualne Spoznavne Sisteme,
Fakulteta za računalništvo in informatiko,
Univerza v Ljubljani

Consider motion of patches of pixels




Select a region of interest
in the first frame.



Assuming the **object will not move by too much** in consecutive frames, **re-localize** the object (target) in each frame.



k-1



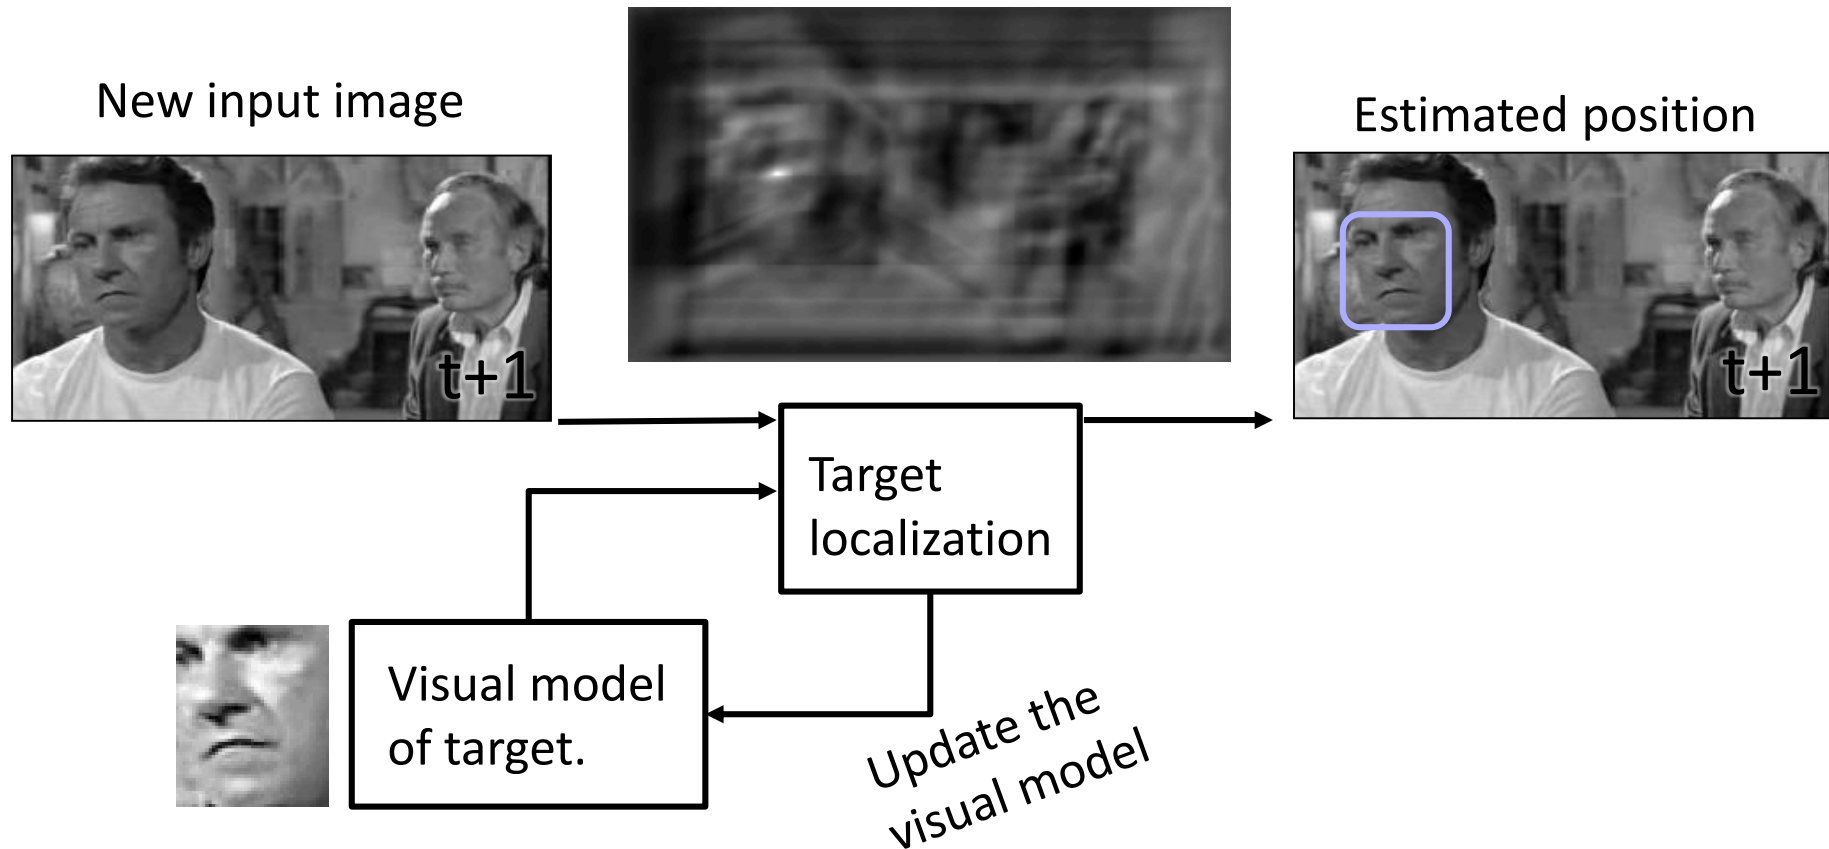
k



k+1

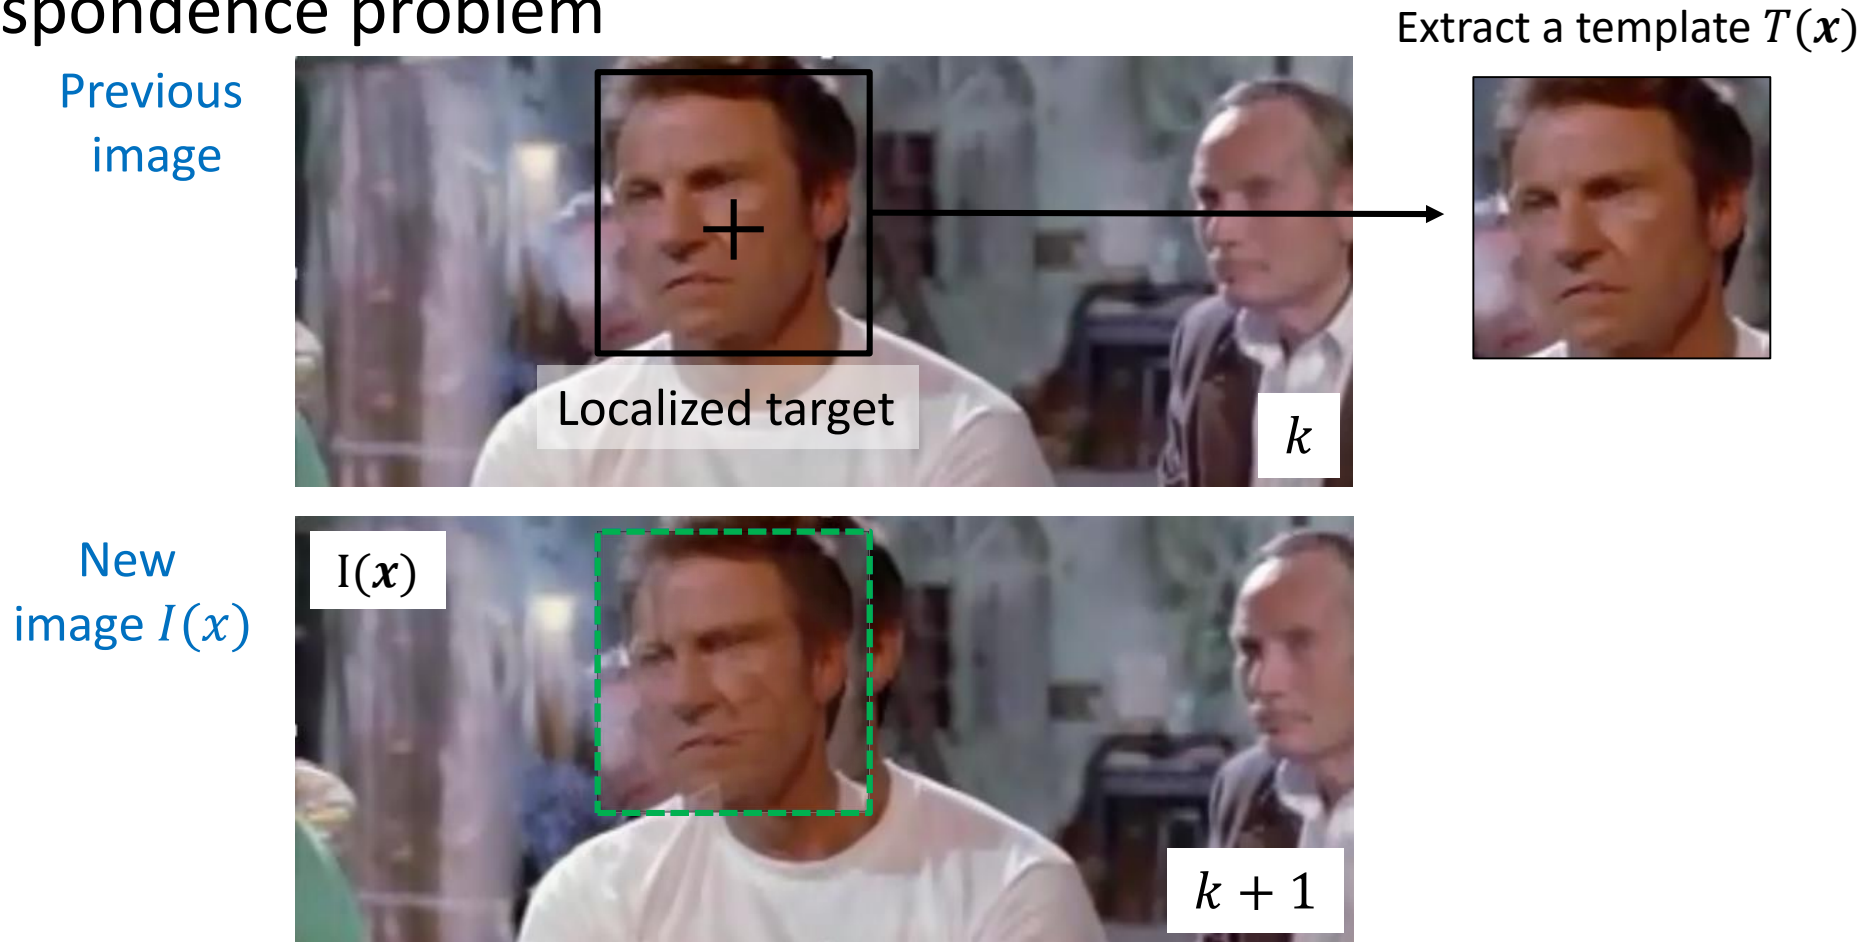
A high-level view of tracking

- Assume some **model** of the target (e.g., template)
- Assume **estimated position** in the previous time-step



Target localization

- Correspondence problem

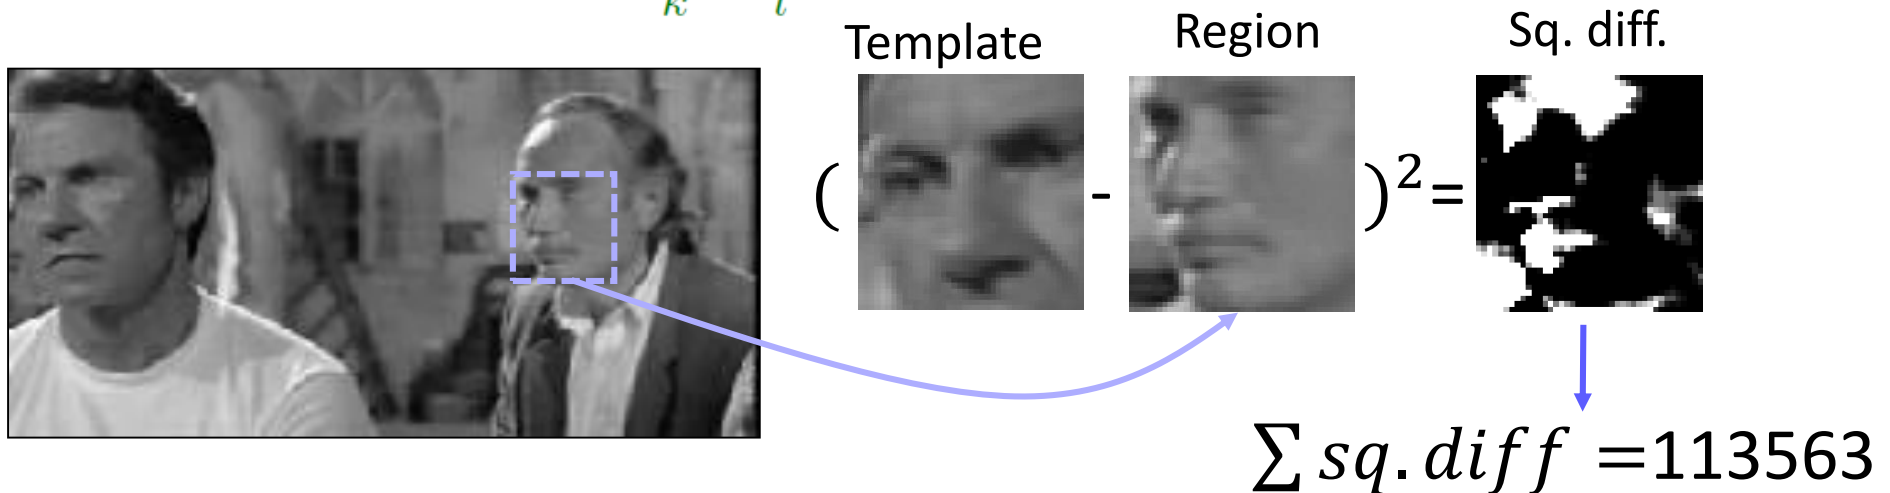


The goal is to align a template image $T(x)$ to an input image $I(x)$.
How to **measure the quality of the alignment**?

Similarity measure


- Quantify the similarity between the visual model and the target region
- Straight-forward: compare pixel intensities
- “Sum of squared differences” (SSD)

$$ssd(x, y) = \sum_k \sum_l (T(k, l) - I(x + k, y + l))^2$$




Naïve localization


- Greedy approach: calculate the SSD for all displacements and select the point where similarity is maximal – the distance is the smallest!



SSD map



SSD map




But usually we can assume our starting position is “close” to the right one!

Can we do better?

- How would we find the bottom of a valley?




1. Decide which way is up/down
2. Move downward by some step
3. Continue until the bottom is reached



Mathematically: Known as "the Gradient descent"


Gradient descent

- Gradient points toward *increase* of f .




Gradient descent

- Move in the opposite direction of the gradient




Gradient descent

- Move in the opposite direction of the gradient



Gradient descent

- A simple recursive algorithm



① Start at some x_0

② For $k=1:\infty$

- $\Delta_k = -\frac{\partial f}{\partial x} \big|_{x_{k-1}} \cdot \alpha$

- $x_k = x_{k-1} + \Delta_k$


- If $|f(x_k) - f(x_{k-1})| < \epsilon$
exit loop

Straight-forward in n-D


- A 2D example

$$\mathbf{x} = [x_1, x_2]^T$$

$$\nabla f = \frac{\partial f}{\partial \mathbf{x}} = \begin{bmatrix} \frac{\partial f(x_1, x_2)}{\partial x_1} \\ \frac{\partial f(x_1, x_2)}{\partial x_2} \end{bmatrix}$$




2D similarity map




Visualize as a 3D surface

Straight-forward in n-D

- Initialize \mathbf{x}_0
- Iterate: $\mathbf{x}_k = \mathbf{x}_{k-1} - \alpha \nabla f|_{\mathbf{x}_{k-1}}$





2D similarity map



Visualize as a 3D surface

The tools we've got so far

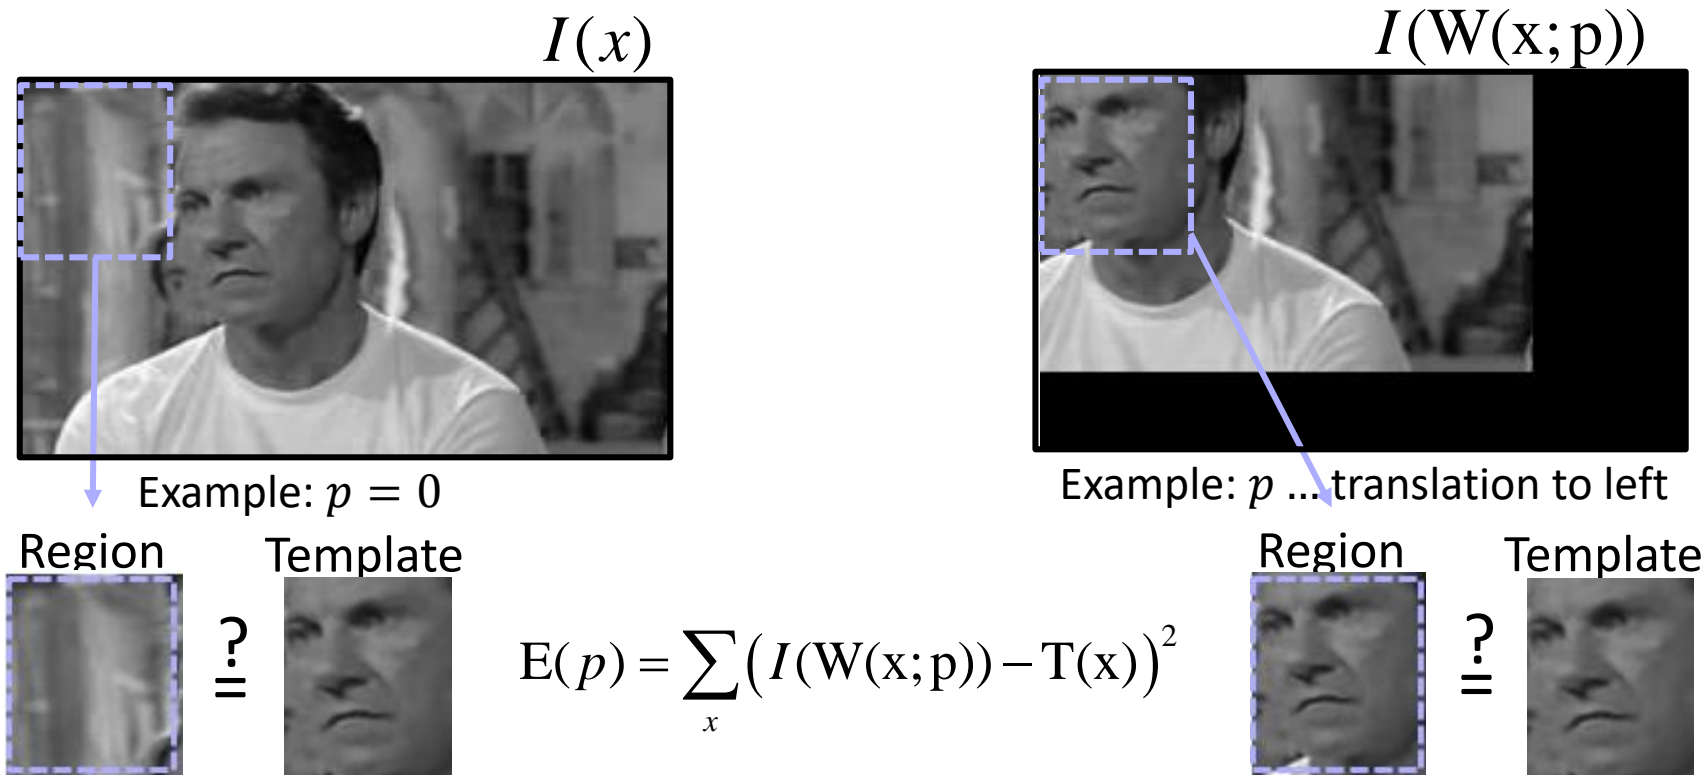
- We know how to **minimize a cost function** $E(p_1, p_2, \dots, p_N)$, w.r.t. \mathbf{p} , where $\mathbf{p} = [p_1, p_2, \dots, p_N]^T$ are parameters of our model.
- We know how to **compute** $E(p_1, p_2)$.



Next: Require *deformation model* that depends on \mathbf{p} .

Displacement models

- Introduce a **warp function** $W(\mathbf{x})$ that warps image onto the coordinate frame of the template – *we can think about the warp as a transformation model* $W(\mathbf{x}; \mathbf{p})$ that takes coordinate \mathbf{x} and transforms it according to parameters \mathbf{p} .



Displacement models

- Simple example:
Translation to left-up in x by 5 and y by 10.




Warped image



Original image


Displacement models

- Popular parametric 2D transformations



Displacement models

- Rigid body motion
 - Rotate, translate



$$\begin{aligned}x' &= x \cos p_1 - y \sin p_1 + p_2 \\y' &= x \sin p_1 + y \cos p_1 + p_3\end{aligned}\quad p = [p_1, p_2, p_3]^T$$

- Compact matrix notation for $W(\mathbf{x}; \mathbf{p})$:

$$W(\mathbf{x}; \mathbf{p}) = \begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} x \cos p_1 - y \sin p_1 + p_2 \\ x \sin p_1 + y \cos p_1 + p_3 \end{bmatrix} = \begin{bmatrix} \cos(p_1) & -\sin(p_1) & p_2 \\ \sin(p_1) & \cos(p_1) & p_3 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$


Displacement models

- Affine motion
 - Rotation, translation, scale, shear

$$x' = p_1x + p_2y + p_3$$

$$y' = p_4x + p_5y + p_6$$

$$p = [p_1, p_2, p_3, p_4, p_5, p_6]^T$$


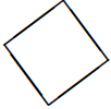


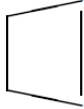


- Compact matrix notation for $W(\mathbf{x}; \mathbf{p})$:

$$W(\mathbf{x}; \mathbf{p}) = \begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} p_1x + p_2y + p_3 \\ p_4x + p_5y + p_6 \end{bmatrix} = \begin{bmatrix} p_1 & p_2 & p_3 \\ p_4 & p_5 & p_6 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

How many free parameters?

- Degrees of freedom DoF (dim. of \mathbf{p})

Transformation	Matrix	# DoF	Preserves	Icon
translation	$\begin{bmatrix} \mathbf{I} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	2	orientation	
rigid (Euclidean)	$\begin{bmatrix} \mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	3	lengths	
similarity	$\begin{bmatrix} s\mathbf{R} & \mathbf{t} \end{bmatrix}_{2 \times 3}$	4	angles	
affine	$\begin{bmatrix} \mathbf{A} \end{bmatrix}_{2 \times 3}$	6	parallelism	
projective	$\begin{bmatrix} \tilde{\mathbf{H}} \end{bmatrix}_{3 \times 3}$	8	straight lines	

Richard Szeliski: [Computer Vision – algorithms and applications](#) (Section 2.1.2)

Tracking as gradient ascent/descent

- Lucas-Kanade tracker
- Initially published in 1981 as an image registration method¹.
- Improved many times, most importantly by Carlo Tomasi².
- Also part of the OpenCV library.
- Single algorithm and results published in a premium journal³.
- Our derivations will follow³
 - See Section 2 in that paper.
 - If you're interested: *See other Sections for improvements of LK and the results obtained by these.*

¹ Lucas and Kanade. An iterative image registration technique with an application to stereo vision. ICAI, 1981.

² Shi and Tomasi. Good features to track. CVPR, 1994.

³ Baker and Matthews. Lucas-Kanade 20 years on: A unifying framework. IJCV, 2004.

Lucas-Kanade algorithm

- Task: *Find the warp* $W(\mathbf{x}; \mathbf{p})$ parameterized by \mathbf{p} , that aligns the image $I(\mathbf{x})$ with a template $T(\mathbf{x})$.
- For example, the warp could be a translation, i.e.,

$$W(\mathbf{x}; \mathbf{p}) = \begin{bmatrix} x + p_1 \\ y + p_2 \end{bmatrix},$$

but in general $W(\mathbf{x}; \mathbf{p})$ can be arbitrary.

- Problem formulation – Find the parameter values of \mathbf{p} that minimize the image differences:

$$E(\mathbf{p}) = \sum_{\mathbf{x}} \left(I(W(\mathbf{x}; \mathbf{p})) - T(\mathbf{x}) \right)^2$$

Lucas-Kanade algorithm

$$E(p) = \sum_x \left(I(W(x; p)) - T(x) \right)^2$$

Finding minimum of $E(p)$ w.r.t. p is a **nonlinear optimization problem**.

We therefore **assume we have initial guess** of p and search for the best increment Δp .

$$E(p, \Delta p) = \sum_x \left(I(W(x; p + \Delta p)) - T(x) \right)^2$$

Iterative solution (think of gradient descent):

$$p \leftarrow p + \Delta p$$



Lucas-Kanade algorithm

- Task: Find the best Δp : $\Delta p = \arg \min_{\Delta p} E(p, \Delta p)$

$$E(p, \Delta p) = \sum_x \left(I(W(x; p + \Delta p)) - T(x) \right)^2$$

Would have been easy if $E(p)$ was quadratic in Δp ...

- To simplify, linearize $I(W(x; p + \Delta p))$ at p :

$$I(W(x; p + \Delta p)) \approx I(W(x; p)) + \nabla I^T \frac{dW}{dp} \Delta p$$

$$\nabla I = \begin{bmatrix} I_x \\ I_y \end{bmatrix} \Big|_{W(x, p)}$$

Jacobian

Note: This is a gradient of image I evaluated at $W(x; p)$, i.e., ∇I is computed in the coordinate frame of I and then *warped* back into the coordinate frame of T by the current estimate of the warp $W(x, p)$.

Jacobians of displacement models

- Translation $W(\mathbf{x}; \mathbf{p}) = \begin{bmatrix} \tilde{x} \\ \tilde{y} \end{bmatrix} = \begin{bmatrix} x + p_1 \\ y + p_2 \end{bmatrix}$

$$\frac{dW(\mathbf{x}; \mathbf{p})}{d\mathbf{p}} = \begin{bmatrix} \frac{\partial \tilde{x}}{\partial p_1} & \frac{\partial \tilde{x}}{\partial p_2} \\ \frac{\partial \tilde{y}}{\partial p_1} & \frac{\partial \tilde{y}}{\partial p_2} \end{bmatrix} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

- Affine $W(\mathbf{x}; \mathbf{p}) = \begin{bmatrix} p_1 x + p_3 y + p_5 \\ p_2 x + p_4 y + p_6 \end{bmatrix}$

$$\frac{dW(\mathbf{x}; \mathbf{p})}{d\mathbf{p}} = \quad ???$$

$$J(W) = \begin{bmatrix} \frac{\partial f_1}{\partial p_1} & \dots & \frac{\partial f_1}{\partial p_n} \\ \frac{\partial f_2}{\partial p_1} & \dots & \frac{\partial f_2}{\partial p_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial p_1} & \dots & \frac{\partial f_m}{\partial p_n} \end{bmatrix}$$

Some pre-computed Jacobians

Transform	Matrix	Parameters p	Jacobian J
translation	$\begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \end{bmatrix}$	(t_x, t_y)	$\begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$
Euclidean	$\begin{bmatrix} c_\theta & -s_\theta & t_x \\ s_\theta & c_\theta & t_y \end{bmatrix}$	(t_x, t_y, θ)	$\begin{bmatrix} 1 & 0 & -s_\theta x - c_\theta y \\ 0 & 1 & c_\theta x - s_\theta y \end{bmatrix}$
similarity	$\begin{bmatrix} 1 + a & -b & t_x \\ b & 1 + a & t_y \end{bmatrix}$	(t_x, t_y, a, b)	$\begin{bmatrix} 1 & 0 & x & -y \\ 0 & 1 & y & x \end{bmatrix}$
affine	$\begin{bmatrix} 1 + a_{00} & a_{01} & t_x \\ a_{10} & 1 + a_{11} & t_y \end{bmatrix}$	$(t_x, t_y, a_{00}, a_{01}, a_{10}, a_{11})$	$\begin{bmatrix} 1 & 0 & x & y & 0 & 0 \\ 0 & 1 & 0 & 0 & x & y \end{bmatrix}$
projective	$\begin{bmatrix} 1 + h_{00} & h_{01} & h_{02} \\ h_{10} & 1 + h_{11} & h_{12} \\ h_{20} & h_{21} & 1 \end{bmatrix}$	$(h_{00}, h_{01}, \dots, h_{21})$	(see Section 6.1.3)

Richard Szeliski: [Computer Vision – algorithms and applications](#) (6.1.1.)

Lucas-Kanade algorithm

- Recall the **original cost function**, i.e.,

$$E(\mathbf{p}, \Delta \mathbf{p}) = \sum_x \left(I(W(\mathbf{x}; \mathbf{p} + \Delta \mathbf{p})) - T(\mathbf{x}) \right)^2$$

- Plugging the linearized term into the above eq. gives

$$E(\mathbf{p}, \Delta \mathbf{p}) \approx \sum_x \left(I(W(\mathbf{x}; \mathbf{p})) + \nabla I^T \frac{dW}{d\mathbf{p}} \Delta \mathbf{p} - T(\mathbf{x}) \right)^2$$

- Observe that $E(\mathbf{p}, \Delta \mathbf{p})$ is quadratic in $\Delta \mathbf{p}$ which means that **$E(\mathbf{p}, \Delta \mathbf{p})$ can be directly minimized** w.r.t. $\Delta \mathbf{p}$:

$$\frac{\partial E(\mathbf{p}, \Delta \mathbf{p})}{\partial \Delta \mathbf{p}} \equiv 0 \quad \Delta \mathbf{p} = ?$$

Lucas-Kanade algorithm

$$\frac{\partial E(\mathbf{p}, \Delta \mathbf{p})}{\partial \Delta \mathbf{p}} \equiv 0$$

$$\Delta \mathbf{p} = \mathbf{H}^{-1} \sum_x \left[\nabla I^T \frac{dW}{d\mathbf{p}} \right]^T [T(\mathbf{x}) - I(W(\mathbf{x}; \mathbf{p}))]$$

Show this at home!


- Where **\mathbf{H}** can be interpreted as a Gauss-Newton approximation of the Hessian

$$\mathbf{H} = \sum_x \left[\nabla I^T \frac{dW}{d\mathbf{p}} \right]^T \left[\nabla I^T \frac{dW}{d\mathbf{p}} \right]$$

Lucas-Kanade algorithm

Iterative solution (think of gradient descent):

- Guess initial parameters \mathbf{p} .
- Construct a linearized cost function $E(\mathbf{p}, \Delta\mathbf{p})$ evaluated at \mathbf{p} .
- Minimize $E(\mathbf{p}, \Delta\mathbf{p})$ w.r.t. $\Delta\mathbf{p}$.
- Update parameters:



LK Implementation

Start with initial \mathbf{p} and iterate:

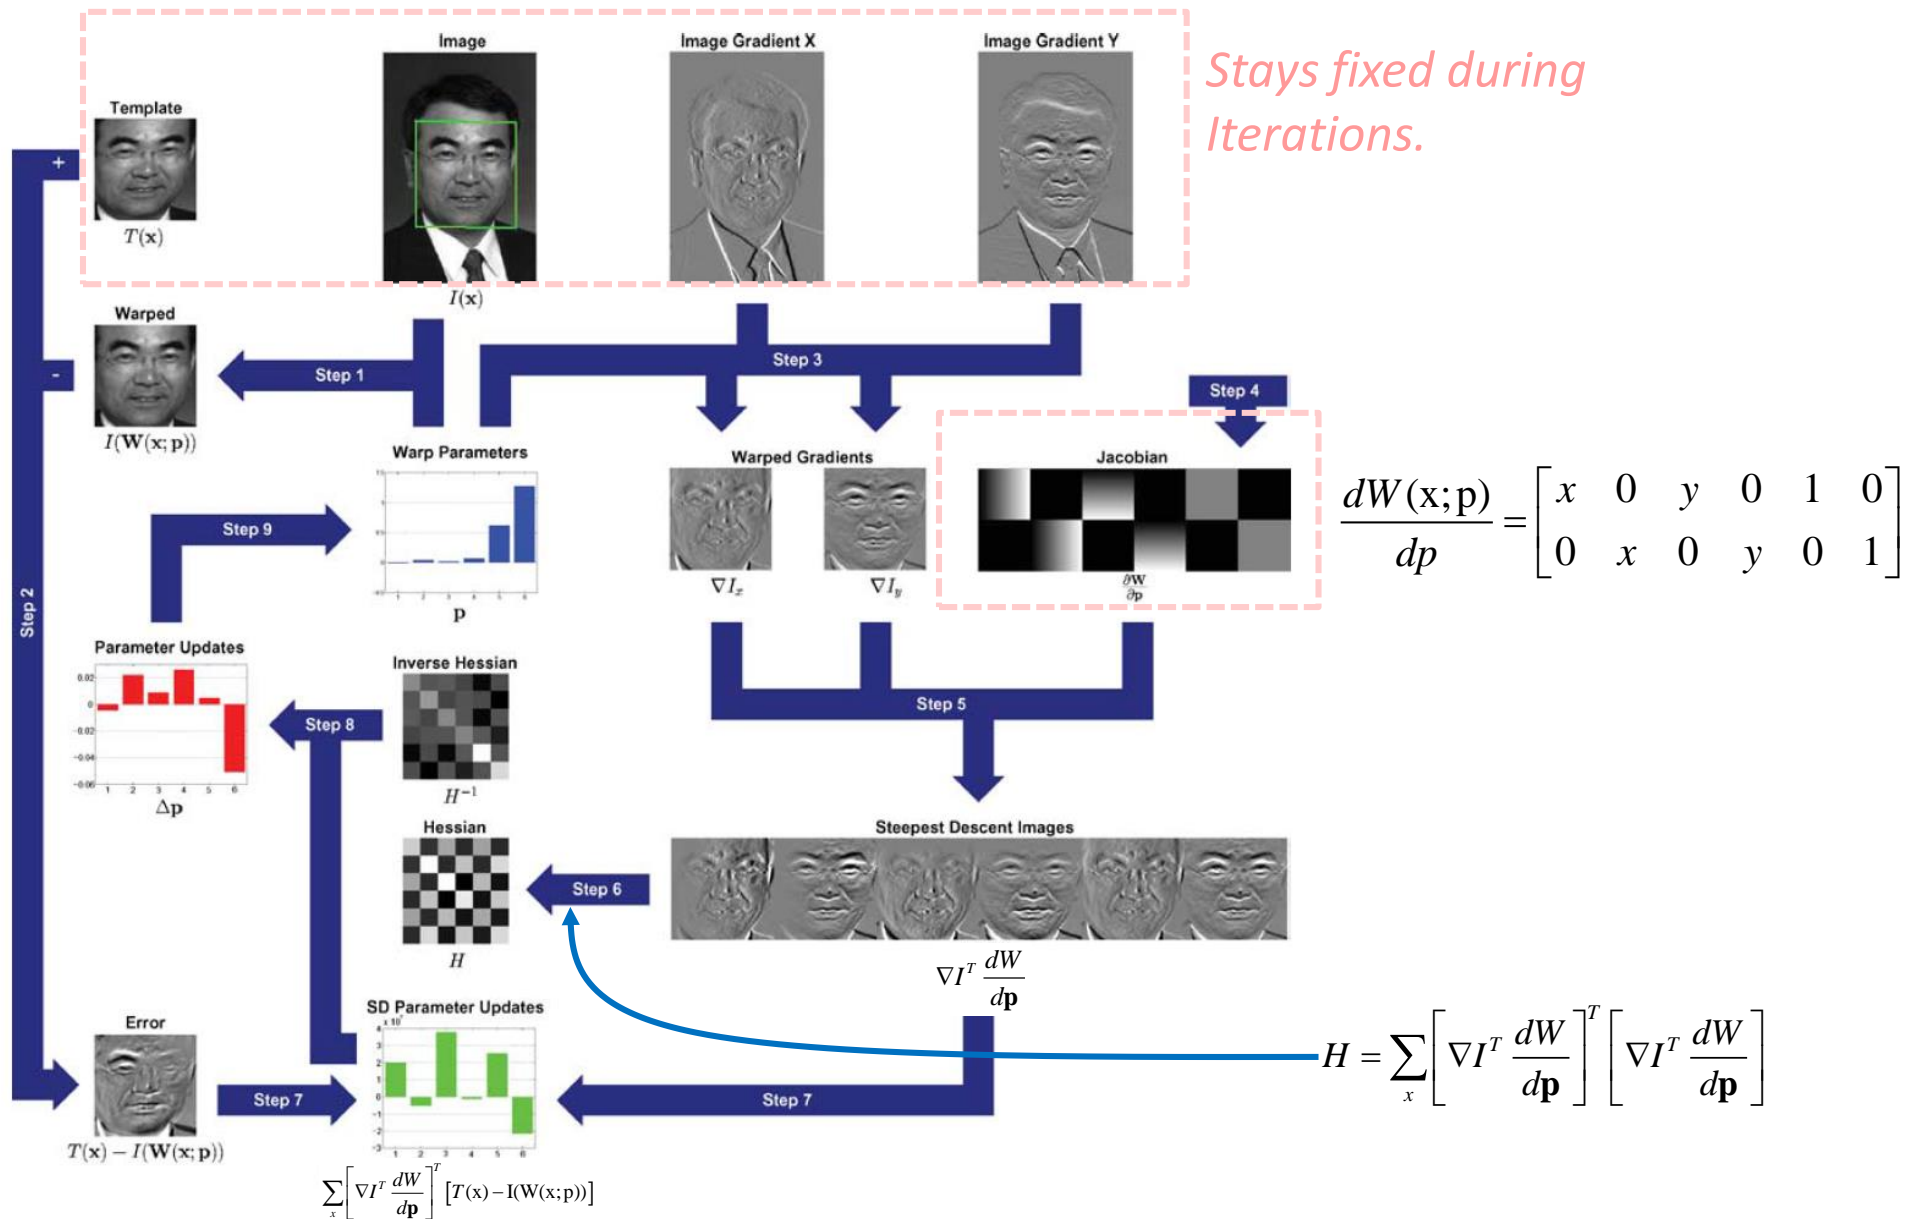
1. Warp image $I(\mathbf{x})$ with $W(\mathbf{x}; \mathbf{p})$.
2. Warp the gradient image $\nabla I(\mathbf{x})$ with $W(\mathbf{x}; \mathbf{p})$.
3. Evaluate the Jacobian $\frac{\partial W}{\partial \mathbf{p}}$ at $(\mathbf{x}; \mathbf{p})$ and compute the steepest descent image $\nabla I^T \frac{dW}{d\mathbf{p}}$.
4. Compute the Hessian $H = \sum_x \left[\nabla I^T \frac{dW}{d\mathbf{p}} \right]^T \left[\nabla I^T \frac{dW}{d\mathbf{p}} \right]$
5. Compute increment $\Delta \mathbf{p} = H^{-1} \sum_x \left[\nabla I^T \frac{dW}{d\mathbf{p}} \right]^T [T(\mathbf{x}) - I(W(\mathbf{x}; \mathbf{p}))]$
6. Update parameters: $\mathbf{p} \leftarrow \mathbf{p} + \Delta \mathbf{p}$

Until $\Delta \mathbf{p} < \epsilon$

(For the sake of completeness – no need to learn by heart)


LK Implementation

$$\Delta p = H^{-1} \sum_x \left[\nabla I^T \frac{dW}{dp} \right]^T [T(x) - I(W(x; p))]$$



Gradient descent visualization

- Assume that warp is translation only $W(x;p) = \begin{bmatrix} \tilde{x} \\ \tilde{y} \end{bmatrix} = \begin{bmatrix} x + p_1 \\ y + p_2 \end{bmatrix}$



Speeded up Lucas Kanade


- The original LK, spends a lot of computation on warping the image and its derivatives.
- The paper¹ suggests a simplification.

Original:

$$E(\Delta p) = \sum_x \left(I(W(x; p + \Delta p)) - T(x) \right)^2$$

New:

$$E(\Delta p) = \sum_x \left(I(W(x; p)) - T(W(x; \Delta p)) \right)^2$$



“The Inverse Compositional Algorithm” (see paper¹, Section 3.2 for details of derivation)

¹Baker and Matthews. Lucas-Kanade 20 years on: A unifying framework. IJCV, 2004.

Lucas-Kanade Inverse Compositional Algorithm

Pre-compute (!!):

- Evaluate gradient ∇T of template $T(x)$.
- Evaluate Jacobian $dW/d\mathbf{p}$.
- Compute steepest descent images $\nabla T^T \frac{dW}{d\mathbf{p}}$.
- Compute hessian $H = \sum_x \left[\nabla T^T \frac{dW}{d\mathbf{p}} \right]^T \left[\nabla T^T \frac{dW}{d\mathbf{p}} \right]$

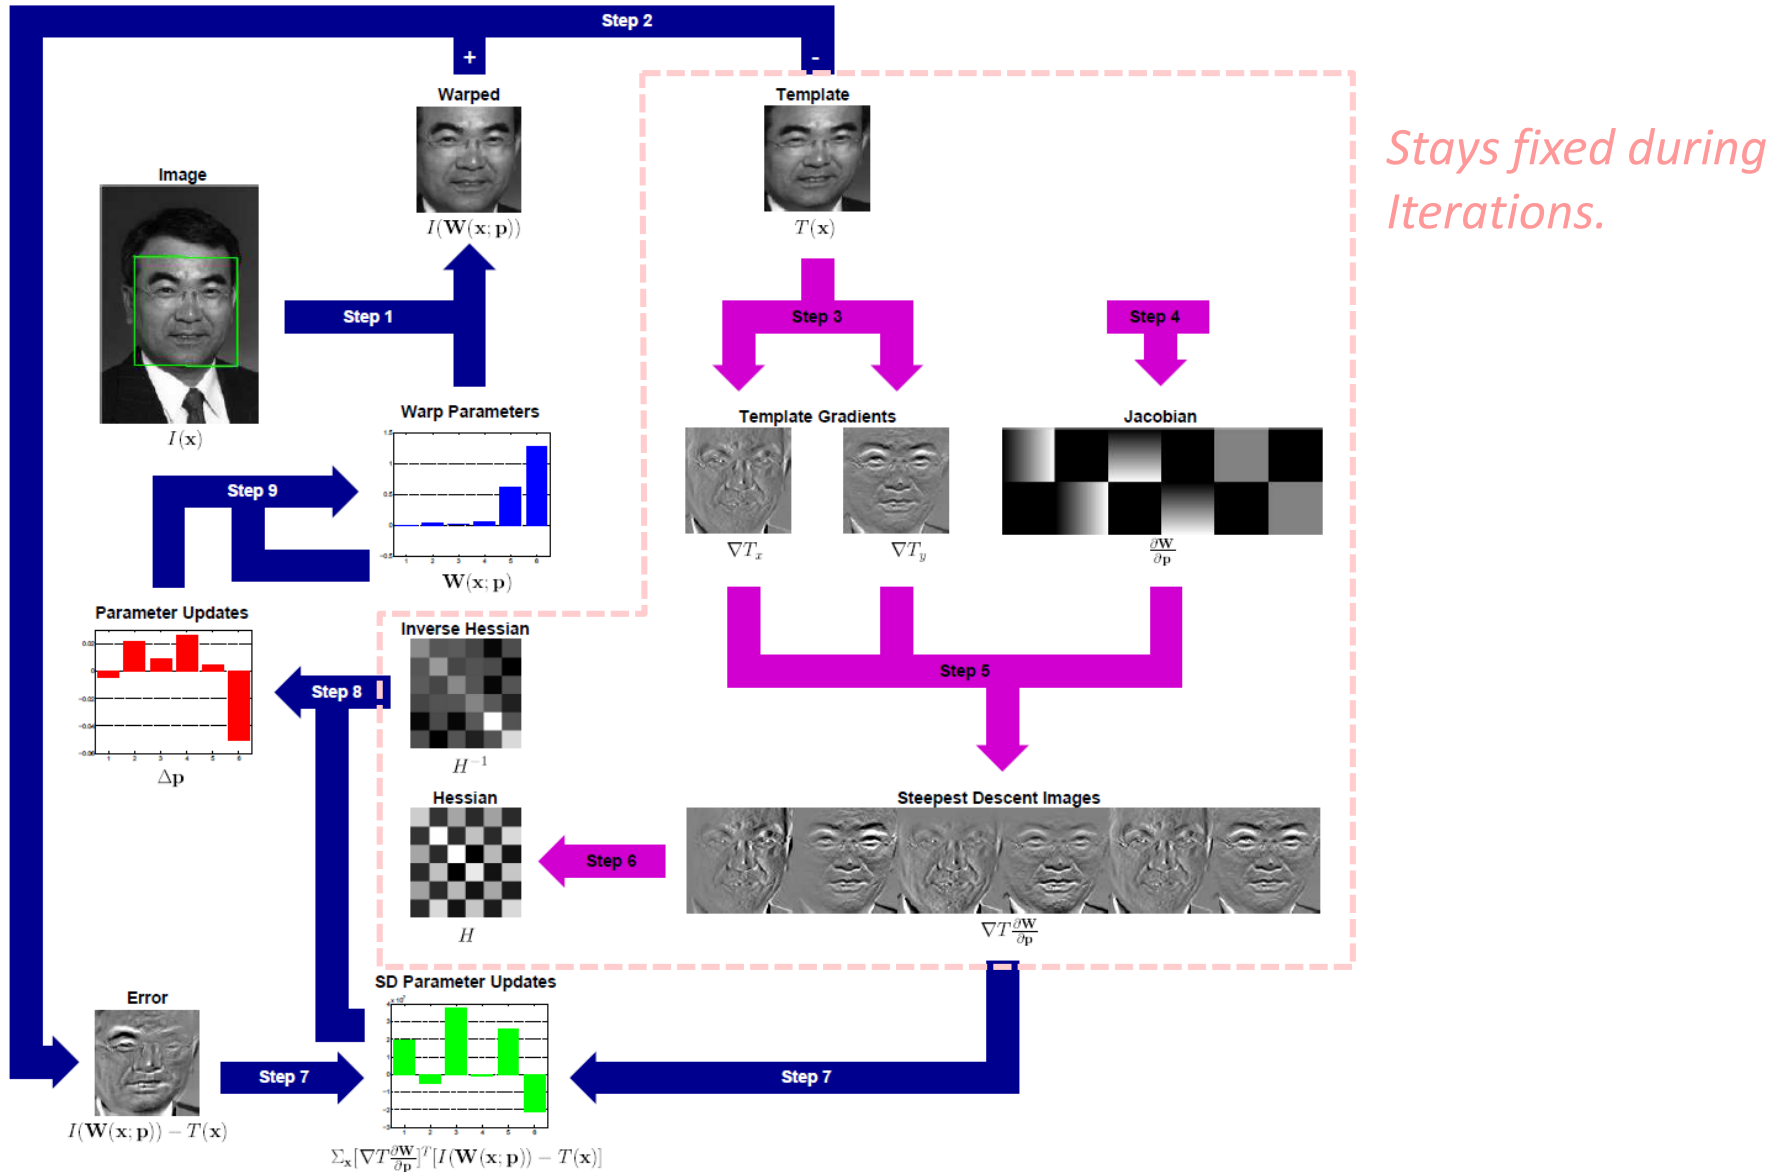
Iterate:

1. Warp image $I(\mathbf{x})$ with $W(\mathbf{x}; \mathbf{p})$
2. Compute steepest descent $\sum_x \left[\nabla T^T \frac{dW}{d\mathbf{p}} \right]^T [I(W(\mathbf{x}; \mathbf{p})) - T(\mathbf{x})]$
3. Compute increment $\Delta \mathbf{p} = H^{-1} \sum_x \left[\nabla T^T \frac{dW}{d\mathbf{p}} \right]^T [I(W(\mathbf{x}; \mathbf{p})) - T(\mathbf{x})]$
4. Update parameters $W(\mathbf{x}; \mathbf{p}) \leftarrow W(\mathbf{x}; \mathbf{p}) \circ W(\mathbf{x}; \Delta \mathbf{p})^{-1}$

(Just for the sake of completeness – no need to learn by heart)


Lucas Kanade ICA

$$\Delta p = H^{-1} \sum_x \left[\nabla T^T \frac{\partial W}{\partial p} \right]^T [I(W(x; p)) - T(x)]$$



What are good features to track?

- Which patches (templates) $T(x)$ should we consider?
- Remember this discussion at LK flow estimation?



Let's look at the maths...

- Which patches (templates) $T(x)$ should we consider?
- The ones for which we can solve the updates

$$\Delta p = H^{-1} \sum_x \left[\nabla I^T \frac{dW}{d\mathbf{p}} \right]^T [T(x) - I(W(x; p))]$$

- Stability depends on whether the Hessian is invertible

$$H = \sum_x \left[\nabla I^T \frac{dW}{d\mathbf{p}} \right]^T \left[\nabla I^T \frac{dW}{d\mathbf{p}} \right]$$

What are good features to track?

- Assume that the warp function is pure translation

$$W(\mathbf{x}; \mathbf{p}) = (x + p_1, y + p_2)$$

$$\frac{dW(\mathbf{x}; \mathbf{p})}{d\mathbf{p}} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

$$H = \sum_x \left[\nabla I^T \frac{dW}{d\mathbf{p}} \right]^T \left[\nabla I^T \frac{dW}{d\mathbf{p}} \right]$$

Note that the Jacobian
is not necessarily constant in general,
but for the translational motion
it is constant!

- Then we can show that the \mathbf{H} is in fact

$$H = \begin{bmatrix} \sum_x I_x^2 & \sum_x I_x I_y \\ \sum_x I_x I_y & \sum_x I_y^2 \end{bmatrix}$$

This is used in the Harris corner detector!

- Means that corners make good features to track.

Verify this by
yourself.

Tracking patches

Without checking similarity
with the initial patch




With checking similarity
with the initial patch



Approach: remove a patch if similarity to
initial template drops below a threshold.

People counting by clustering KLT



Vincent Rabaud and Serge Belongie, Counting Crowded Moving Objects [[pdf](#)] [[poster](#)] [CVPR 2006](#), New York, NY.

Tracking facial points by LK ICA



(a)



(d)



>200 frames per second



(c)




(f)

- [1] Iain Matthews and Simon Baker, "[Active Appearance Models Revisited](#)," International Journal of Computer Vision, Vol. 60, No. 2, 2004
- [2] Simon Baker, Iain Matthews, Jing Xiao, Ralph Gross, Takeo Kanade, and Takahiro Ishikawa, "[Real-Time Non-Rigid Driver Head Tracking for Driver Mental State Estimation](#)," 11th World Congress on Intelligent Transportation Systems, October, 2004.

Motion stabilization and stitching



- LK can be used for **motion compensation**
- We can consider the entire image as template



- Choose a **pseudo-perspective** transform for $W(x;p)$
(pseudo-perspective is approximation for perspective)

Motion stabilization and stitching

- LK can be used for **motion compensation**
- We can consider the entire image as template



Tracking by sparse flow

- Apply Lucas Kanade (pyramidal) to **estimate sparse flow**.
- **Fit a parametric model** to the flows, e.g., affine, by least squares or RANSAC, (or by a simple median if the model is simple enough).



t



t+1

For least squares and RANSAC, see Richard Zseliski:

[Computer Vision – algorithms and applications](#) (6.1.1-6.1.4)

Tracking by a grid of flow vectors




- Apply a **grid of LK flows** and estimate **reliability** of each computed **flow vector**.



Tomas Vojir and Jiri Matas, “[The Enhanced Flock of Trackers](#)”. *Registration and Recognition in Images and Videos - Studies in Computational Intelligence*, Springer 2014. ([bib](#))

LK combination with deep features

- “Recent” work proposed learning deep features such that the cost function optimized in LK tracker becomes smooth with a large attraction perimeter



Wang et al., [Deep-LK for Efficient Adaptive Object Tracking](#), ICRA2018 [[Youtube link](#)]

References on LK

Recommended read:

- Baker and Matthews. Lucas-Kanade 20 years on: A unifying framework. IJCV, 2004.
 - [At least the section on basic Lucas&Kanade optimization](#)

If you are interested in some milestone papers:

- Lucas and Kanade. An iterative image registration technique with an application to stereo vision. ICAI, 1981.
- Shi and Tomasi. Good features to track. CVPR, 1994.