Univerza v Ljubljani

Advanced CV methods
Long-Term tracking

Matej Kristan

Laboratorij za Umetne Vizualne Spoznavne Sisteme,
Fakulteta za računalništvo in informatiko,
Univerza v Ljubljani

Long-term tracking (LTT)



- Regardless of how well the visual model is designed, any short-term tracker will eventually fail
- Disappears from the field of view, gets fully occluded, etc.

Long-term tracking (LTT)

- The general LT tracking properties:
  - Determine when the target has been lost (or disappeared)
  - Re-detect the target after losing the target
  - Update the visual model very carefully to minimize drifting

Taxonomy: Short-term/long-term spectrum[1]

	Position reported	Tracking failure detection	Target re-detection
ST_0: Basic ST	each frame	no	no
ST_1: Basic ST with conservative updating	each frame	not explicitly, selective update of visual model	no
LT_0: Pseudo LT	only when visible	yes	no
LT_1: Re-detecting LT	only when visible	yes	yes

- ST_0 (e.g., vanilla DCF, MS); ST_1 (e.g., MDNet) -> easily converted to LT_0
- LT_1 most sophisticated, typical composition:
  - Short-term tracker (ST) for frame-to-frame localization
  - Detector for target re-detection
  - Algorithm for interaction between ST and detector

[1] Lukežič, et al., *Now you see me: evaluating performance in long-term visual tracking*, TCyb 2020

LT1 trackers origin

- Most of the LT_1 originate from two main paradigms introduced by *TLD*[1] (aka Predator) and *Alien*[2]
- In the following we will overview the TLD

[1]Kalal, Mikolajczyk, Matas, Tracking-Learning-Detection, TPAMI2010

[2]Pernici, F. and Del Bimbo, A., Object Tracking by Oversampling Local Features, TPAMI2013

Advanced computer vision methods

TRACKING BY TRACKING, LEARNING, DETECTION (PREDATOR)

Tracking learning detection: TLD aka Predator[1]

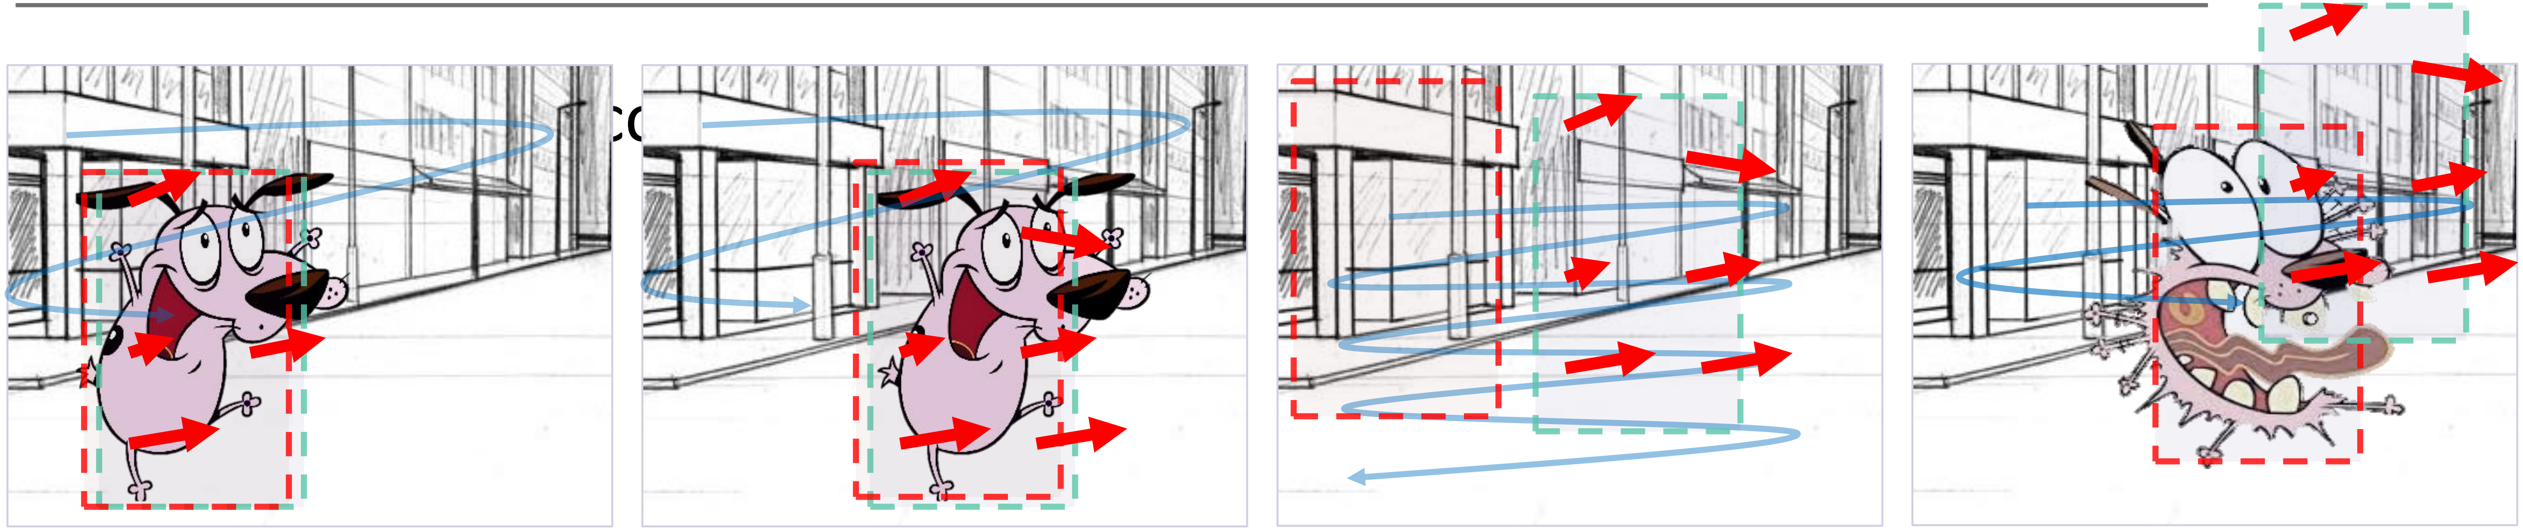

- Detector is the main component
- It's all about robust detector updating
- Run Detector and ST tracker in parallel
- Use the ST and Detector output to construct training samples for Detector

Short-term:

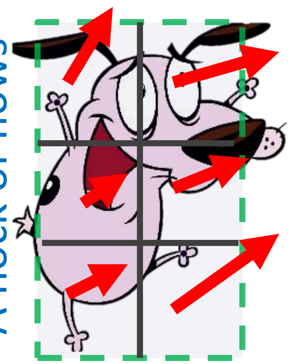

Detector:

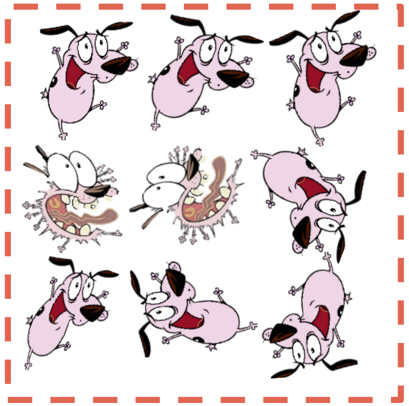

[1]Kalal, Mikolajczyk, Matas, Tracking-Learning-Detection, TPAMI2010

Fast-forward… "TLD in action"

Kalal, Mikolajczyk, Matas, Tracking-Learning-Detection, TPAMI2010

The short-term tracker

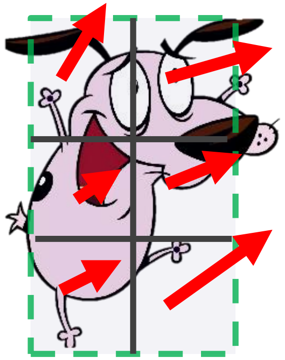

- A “cell” grid of Lucas-Kanade trackers
- Each LK tracker has a reliability estimate
- Robustly estimates motion from 50% of most reliable displacements
  (could also use a robust estimator, e.g., RANSAC)
- 2 layers of Pyramidal LK tracker
  with 10×10 pixels patches.
- Fairly robust frame-to-frame localization
  in absence of severe occlusion

Z. Kalal, K. Mikolajczyk, and J. Matas. Forward-Backward Error: Automatic Detection of Tracking Failures. ICPR, 2010
Improved version:
T. Vojir and J. Matas. Robustifying the flock of trackers. CVWW2011

The detector visual model

- Appearance model: a grayscale patch
- Bounding box with fixed aspect
  (only scale changes, proportions constant)
- Patch resampled into 15x15 size

- Object model is a collection of multiple positive and negative patches!
- Forget patches (randomly) to keep the number of patches low enough
  (memory and speed efficiency)

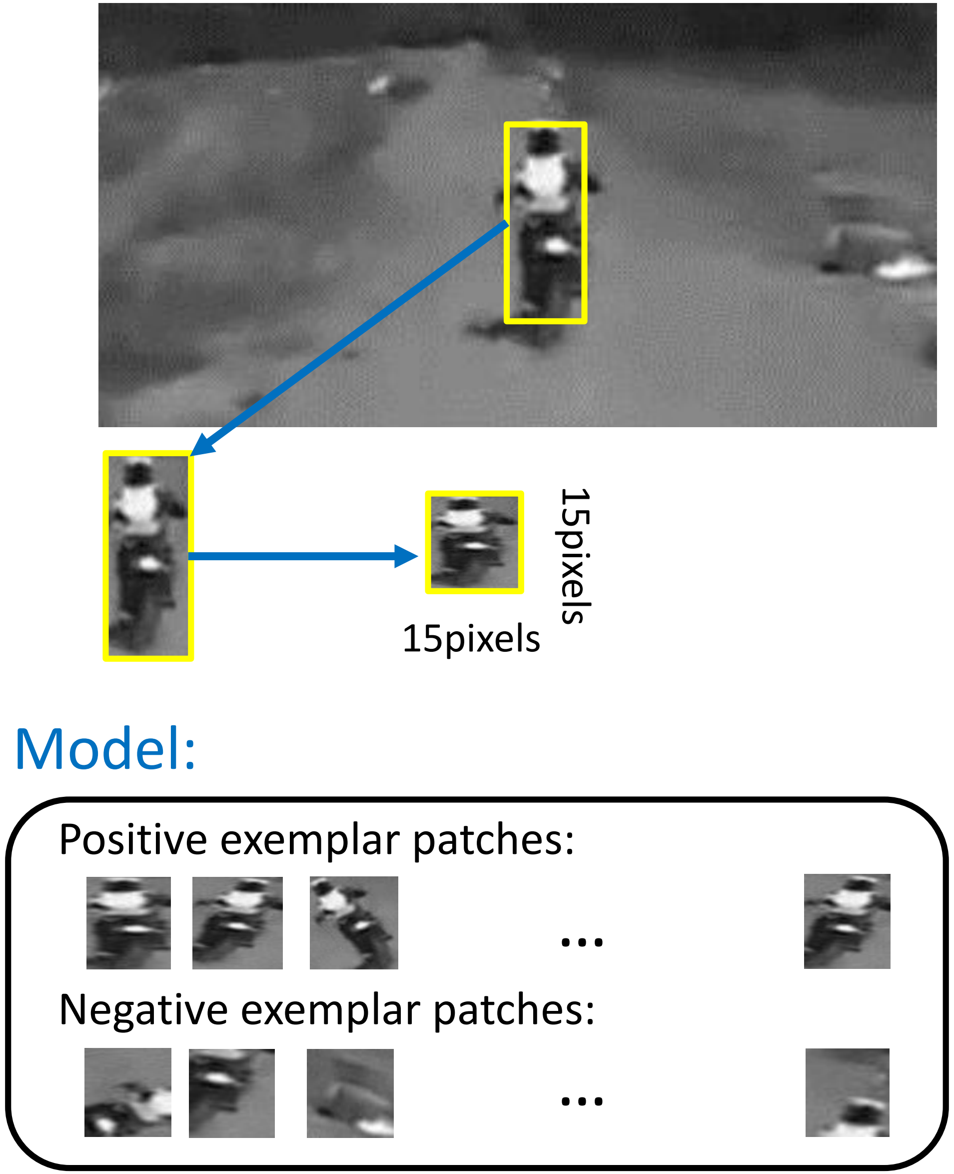

The detector application

- A scanning window
- Compare patches using a normalized cross correlation (NCC)
- A nearest-neighbor classifier using the NCC score
- Problem: A brute force would require comparing all locations with all patches in the model!
- Solution: Apply cascaded approach that quickly rejects many potential image locations by using simple and fast features.

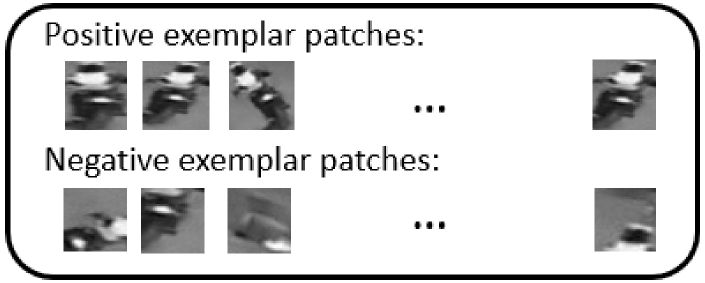

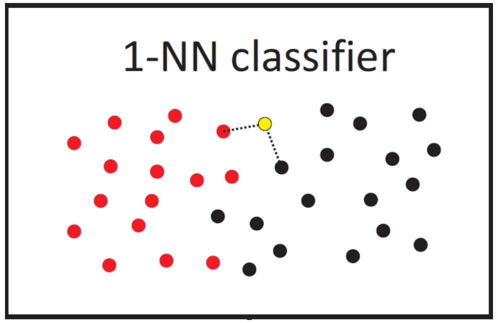

Fast classifiers with low FP/FN, high TP

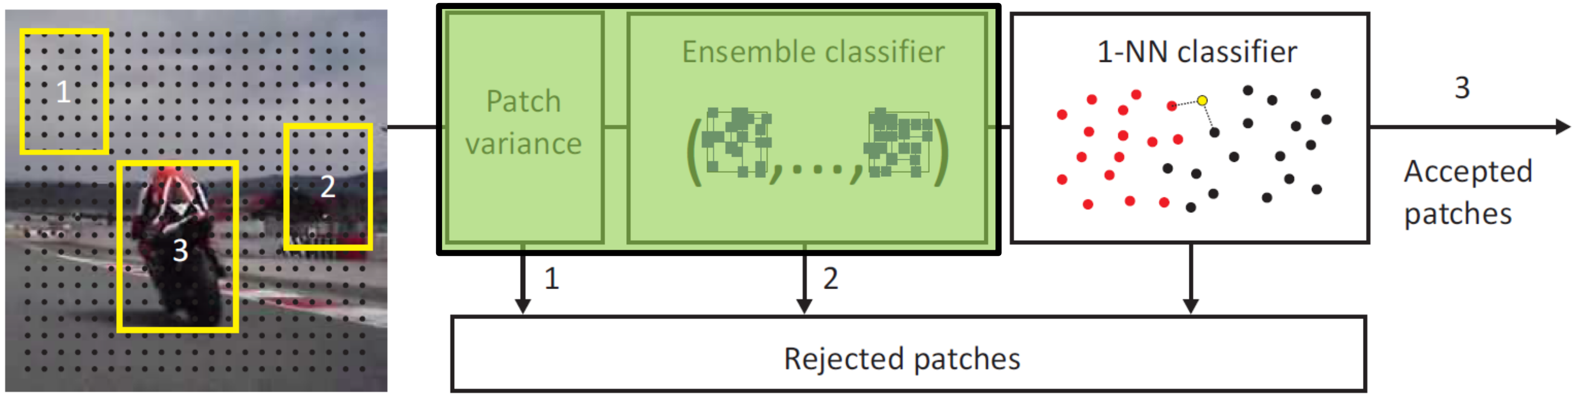

The ST-Detector interaction algorithm

- PN learning: Responsible for training the Detector

- PN (semi-supervised) learning assumptions:
  - Two classes of labelling processes are available: P and N
  - "P" proposes positive, the "N" proposes negative examples only.
  - Both processes are noisy and can make mistakes
  - By carefully addressing the conflicts between the two labelling processes, a long-term stability is achieved.

Interaction algorithm P-event: “Loop”

- Guideline: *Do not trust the learning examples until you are absolutely sure about their labels!*
- Exploits temporal structure
- Assumption: If an adaptive tracker fails, it is unlikely to recover.
- Rule: Patches from a track starting and ending in the current model (red), i.e. are validated by the detector, are added to the model.

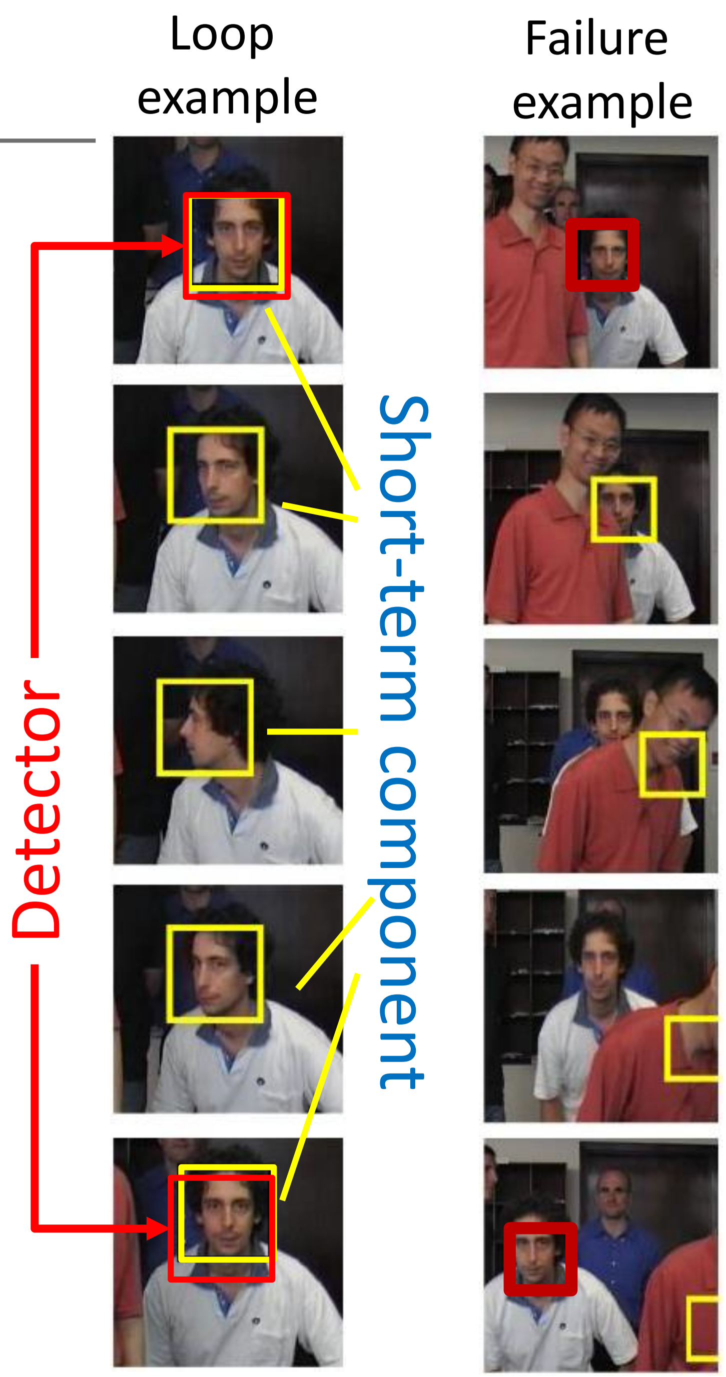

Interaction algorithm N-event: "Uniqueness"

- Exploits spatial structure
- Assumption:
  Object is unique in a single frame
  (no other object looks alike)
- Rule: If the tracked patch is in the model, all other detections within the current frame (red) are assumed wrong
  → *are pruned from the model*

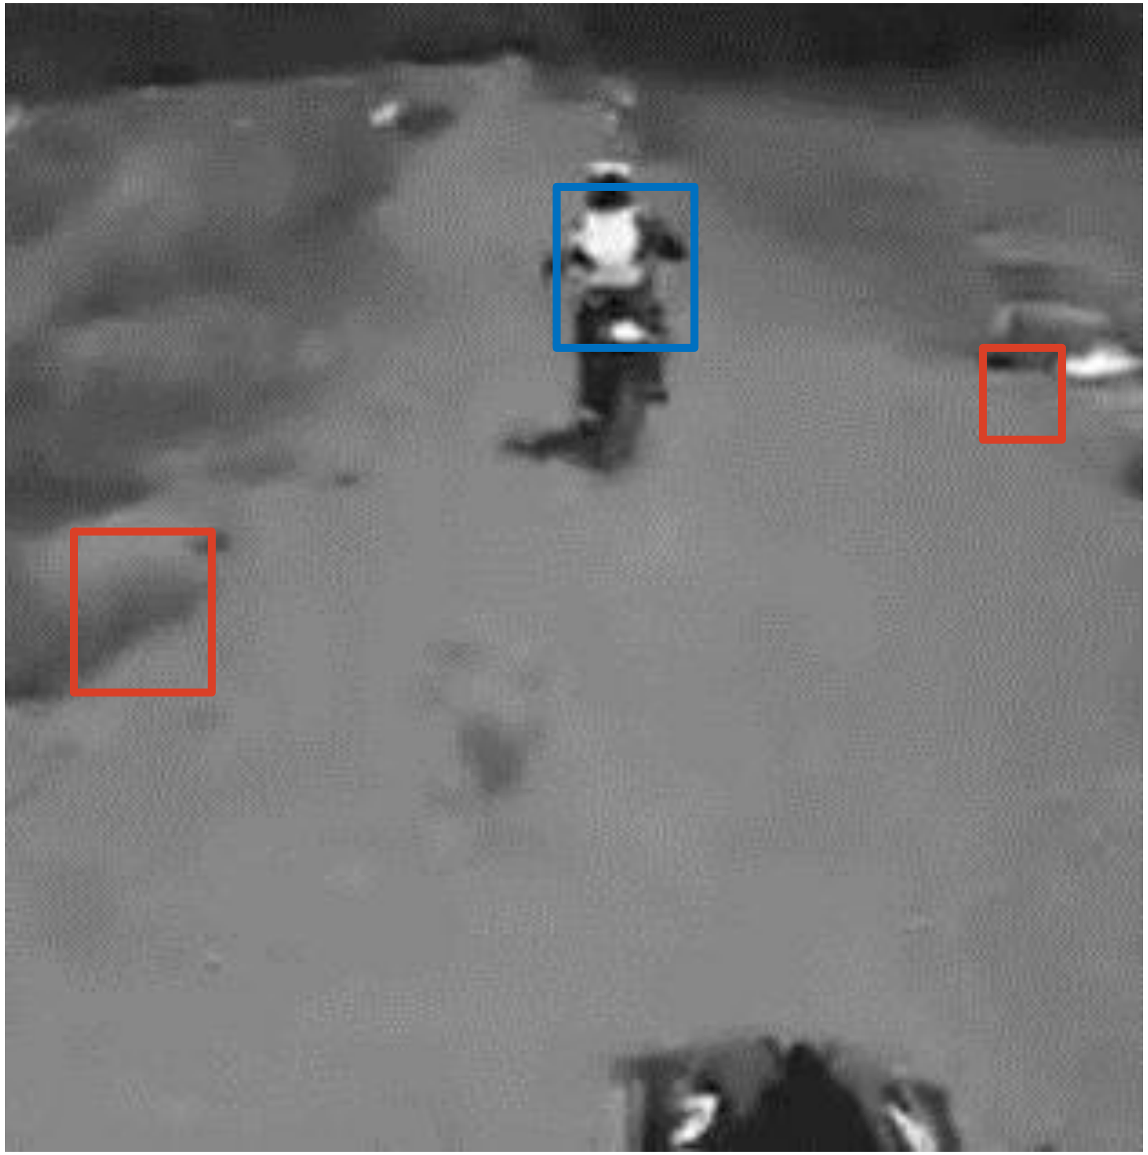

Interaction algorithm: Model learning

Defined by:

- P-events, N-events, detector learning method
- P and N events are defined in terms of tracker and detector outputs

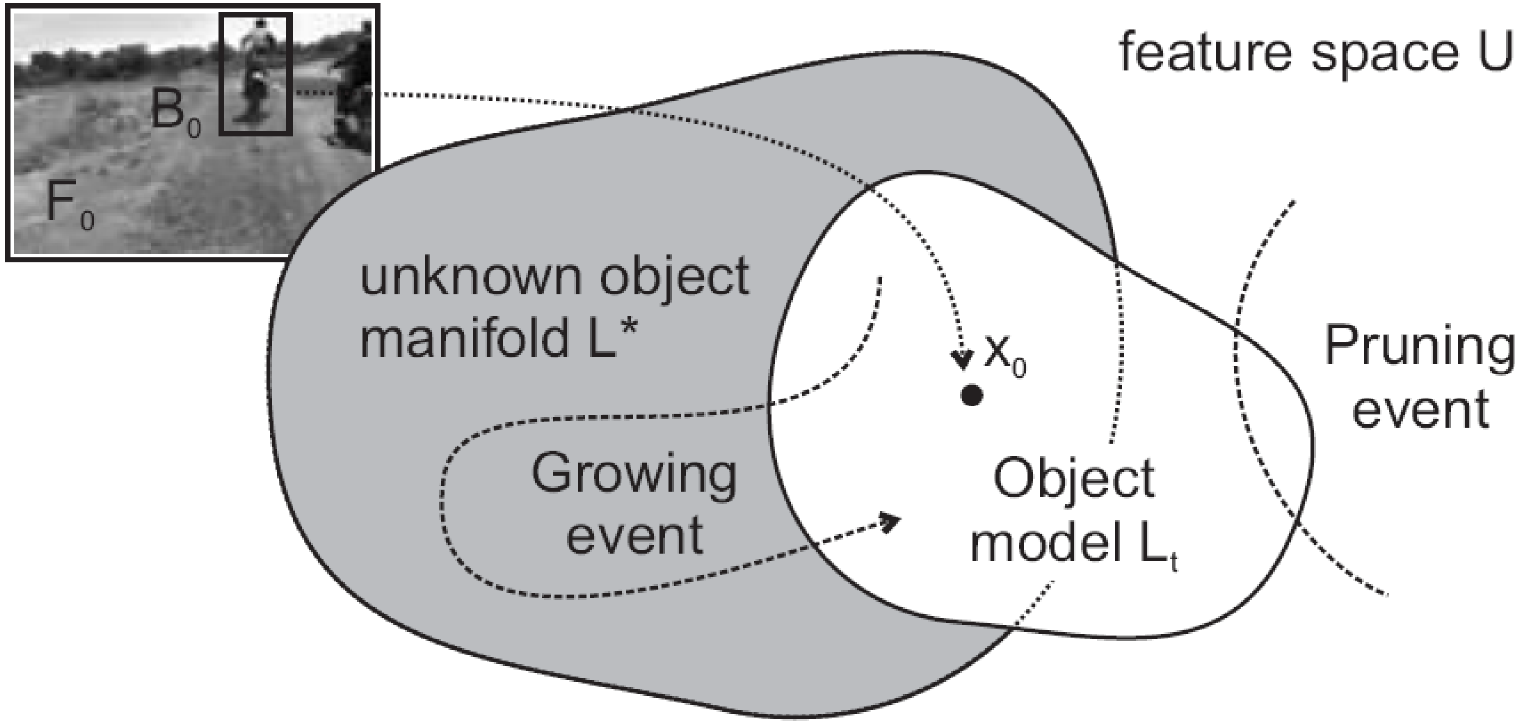

TLD tracking-learning example

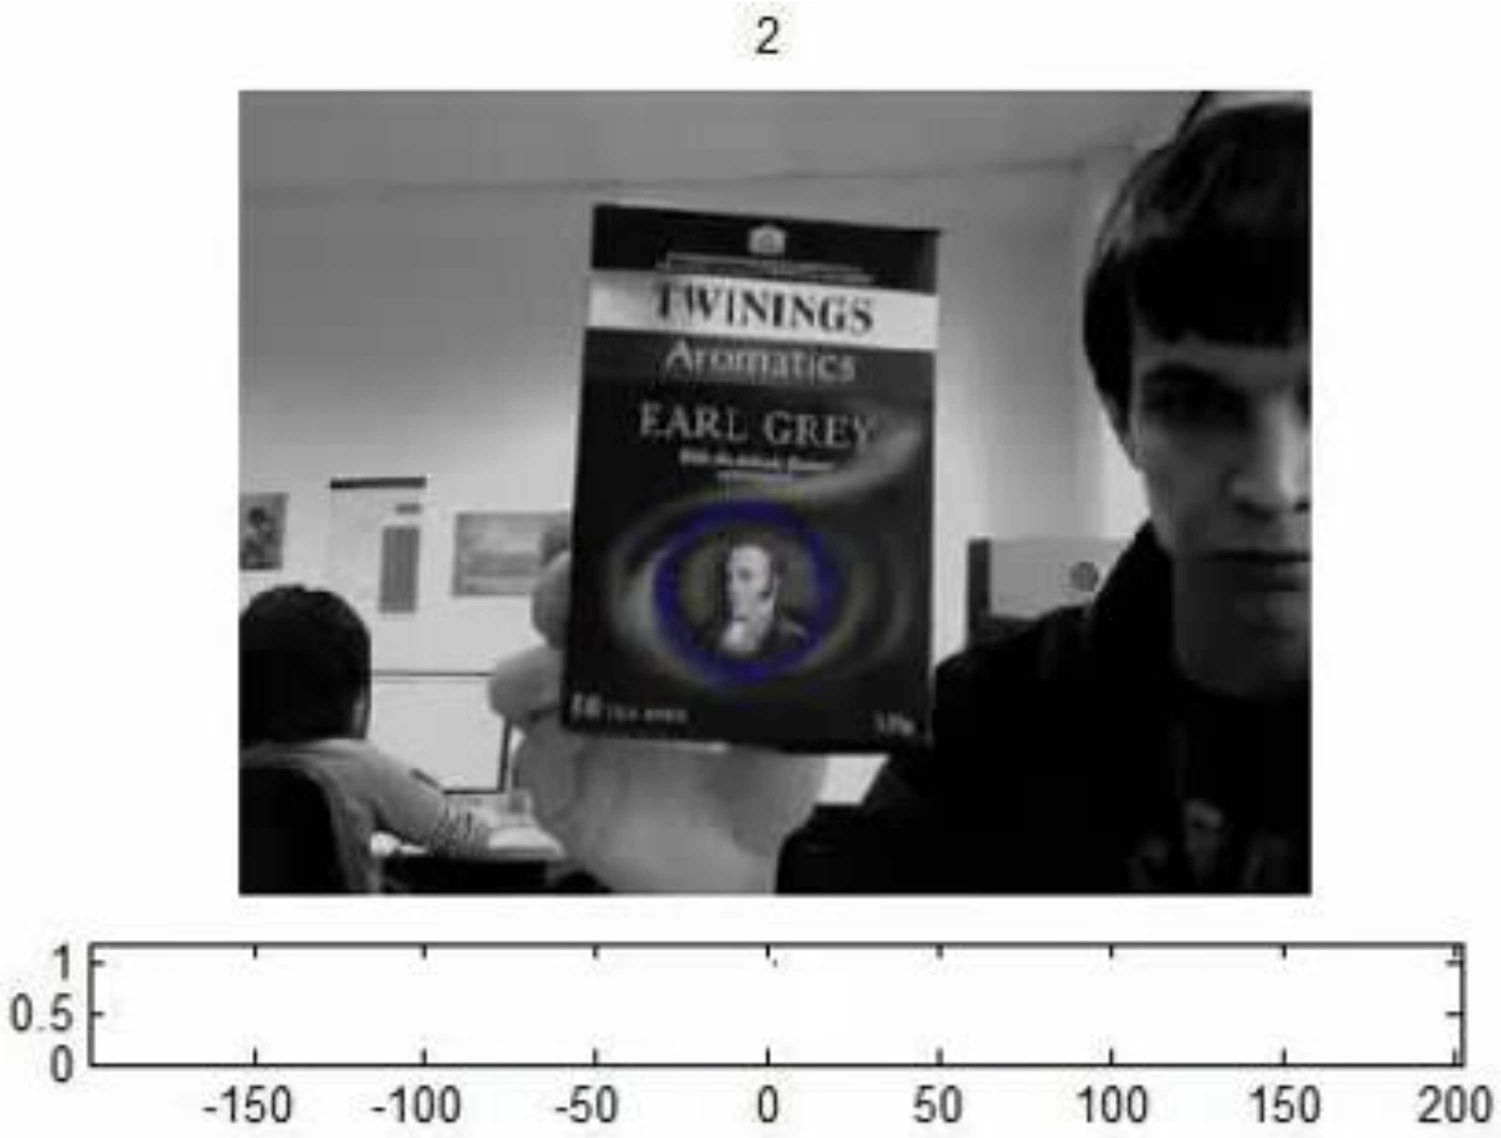

Detector templates (positives)

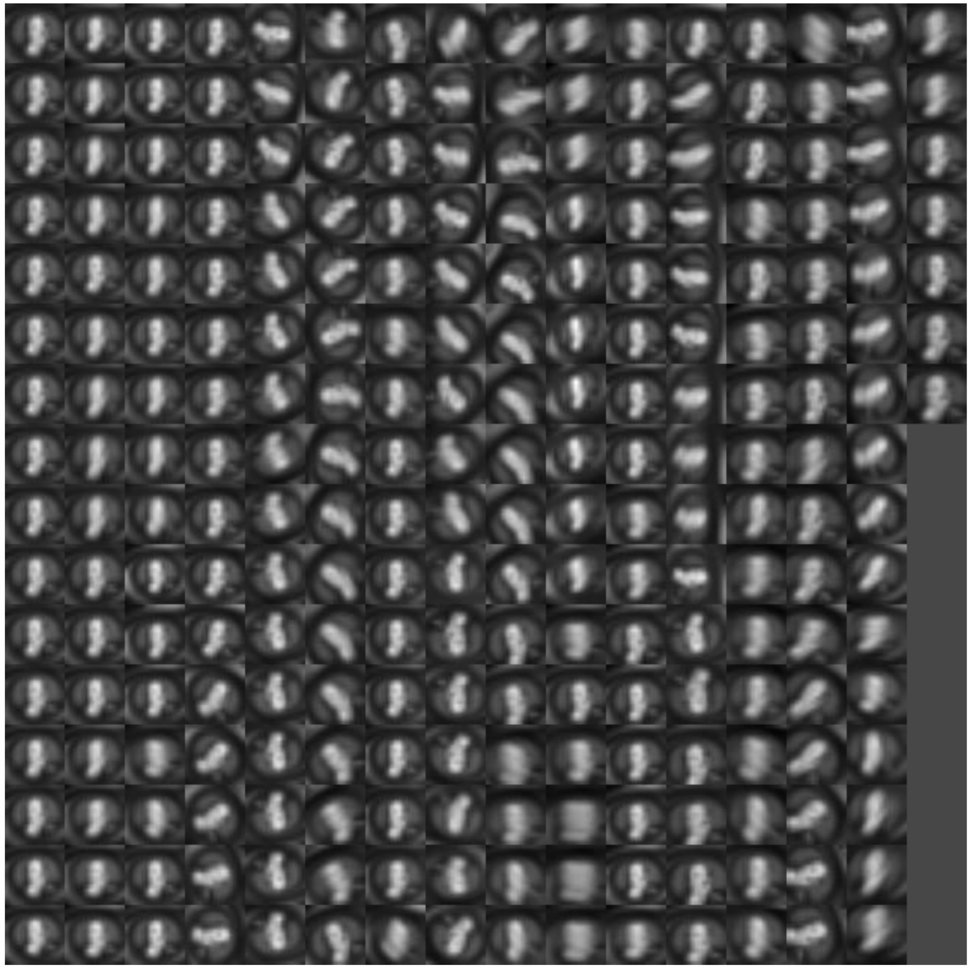

TLD tracking example

TLD summary

- PN Learning trains a robust detector by observing the object of interest (no a priori labelled training data, no constraints on the video)
- Detector improves over time (experimentally validated)
- A stable semi-supervised learning algorithm
- Matlab/C++ implementation runs at > 20 fps (back in 2010)

- Code is available online: http://personal.ee.surrey.ac.uk/Personal/Z.Kalal/

Kalal, Mikolajczyk, Matas, Tracking-Learning-Detection, TPAMI2010

Long-Term Architecture Implementation Issues

Tracker	Short-term tracker	Detector	Interaction
Alien [6]	Keypoints (SIFT)	Keypoints (SIFT)	F-B, Ransac
TLD [1]	Optical flow	Random forest	P-N learning
MUSTER [2]	Correlation filter	Keypoints (SIFT)	F-B, Ransac
LCT [3]	Correlation filter	Random fern	K-NN, response thresh.
CMT [4]	Keypoints (flow)	Keypoints (static)	F-B, clustering
PTAV [5]	Correlation filter	CNN (Siam. Net.)	CNN confidence score

Approaches from different methodologies

- Prohibits tight interaction e.g., feature/model sharing
- Leads to complicated implementation

[1] Kalal et al., Tracking-Learning-detection, TPAMI 2010
[2] Ma et al., Long-Term Correlation Tracking, CVPR 2015
[3] Hong et al., MUlti-Store Tracker (MUSTer): a Cognitive Psychology Inspired Approach to Object Tracking, CVPR 2015
[4] Nebehay et al., Clustering of Static-Adaptive Correspondences for Deformable Object Tracking, CVPR 2015
[5] Fan et al., Parallel Tracking and Verifying: A Framework for Real-Time and High Accuracy Visual Tracking, ICCV 2017
[6] Pernici, F. and Del Bimbo, A., Object Tracking by Oversampling Local Features, TPAMI2013

Long-Term Architecture Implementation Issues

Tracker	Short-term tracker	Detector	Interaction
Alien [6]	Keypoints (SIFT)	Keypoints (SIFT)	F-B, Ransac
TLD [1]	Optical flow	Random forest	P-N learning
MUSTER [2]	Correlation filter	Keypoints (SIFT)	F-B, Ransac
LCT [3]	Correlation filter	Random fern	K-NN, response thresh.
CMT [4]	Keypoints (flow)	Keypoints (static)	F-B, clustering
PTAV [5]	Correlation filter	CNN (Siam. Net.)	CNN confidence score
FCLT [7]	Correlation filter	Correlation filter	Correlation uncertainty

Shared target representation: tight interaction, efficient implementation

- Short-term tracker and a detector within a single methodology
- A single DCF learner, two interacting models

[7] Lukežič, Čehovin, Vojir, Matas, Kristan, *FuCoLoT -- A Fully-Correlational Long-Term Tracker*, ACCV 2018

Fully Correlational Long-term Tracker (FCLT)

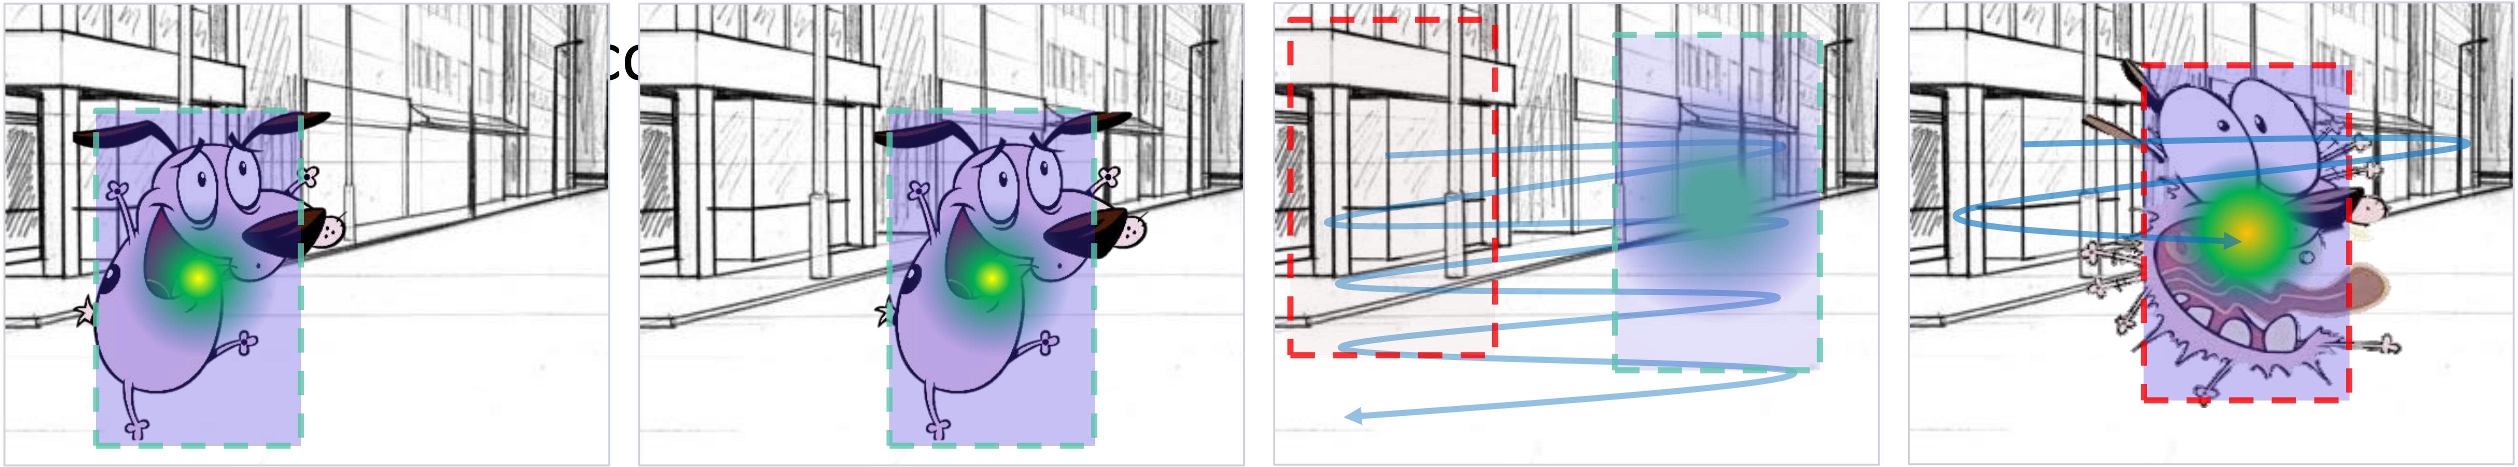

- Discriminative correlation filter in two separate components.
- Detector activated when ST not confident.
- Motion model used with detector.

Short-term:

Detector:

[1]Lukežič et al., *Discriminative Correlation Filter Tracker with Channel and Spatial Reliability*, IJCV 2018

FCLT: ST and Detector learning

- Short-term (ST) model is a CSRDCF[1] with standard update
- Detector:
  - Standard DCF cannot be used for image-wide detection
  - Utilize constrained learning from CSRDCF[1] from a wider region
  - Several object models updated at various time scales

Detector 1:

Never update

Detector 2:

Update every 250th

Detector 3:

Update every 50th

…

Detector N:

Update every frame

[1]Lukežič, Vojir, Čehovin Zajc, Matas and Kristan, *Discriminative Correlation Filter Tracker with Channel and Spatial Reliability*, IJCV 2018

FCLT: Detector application

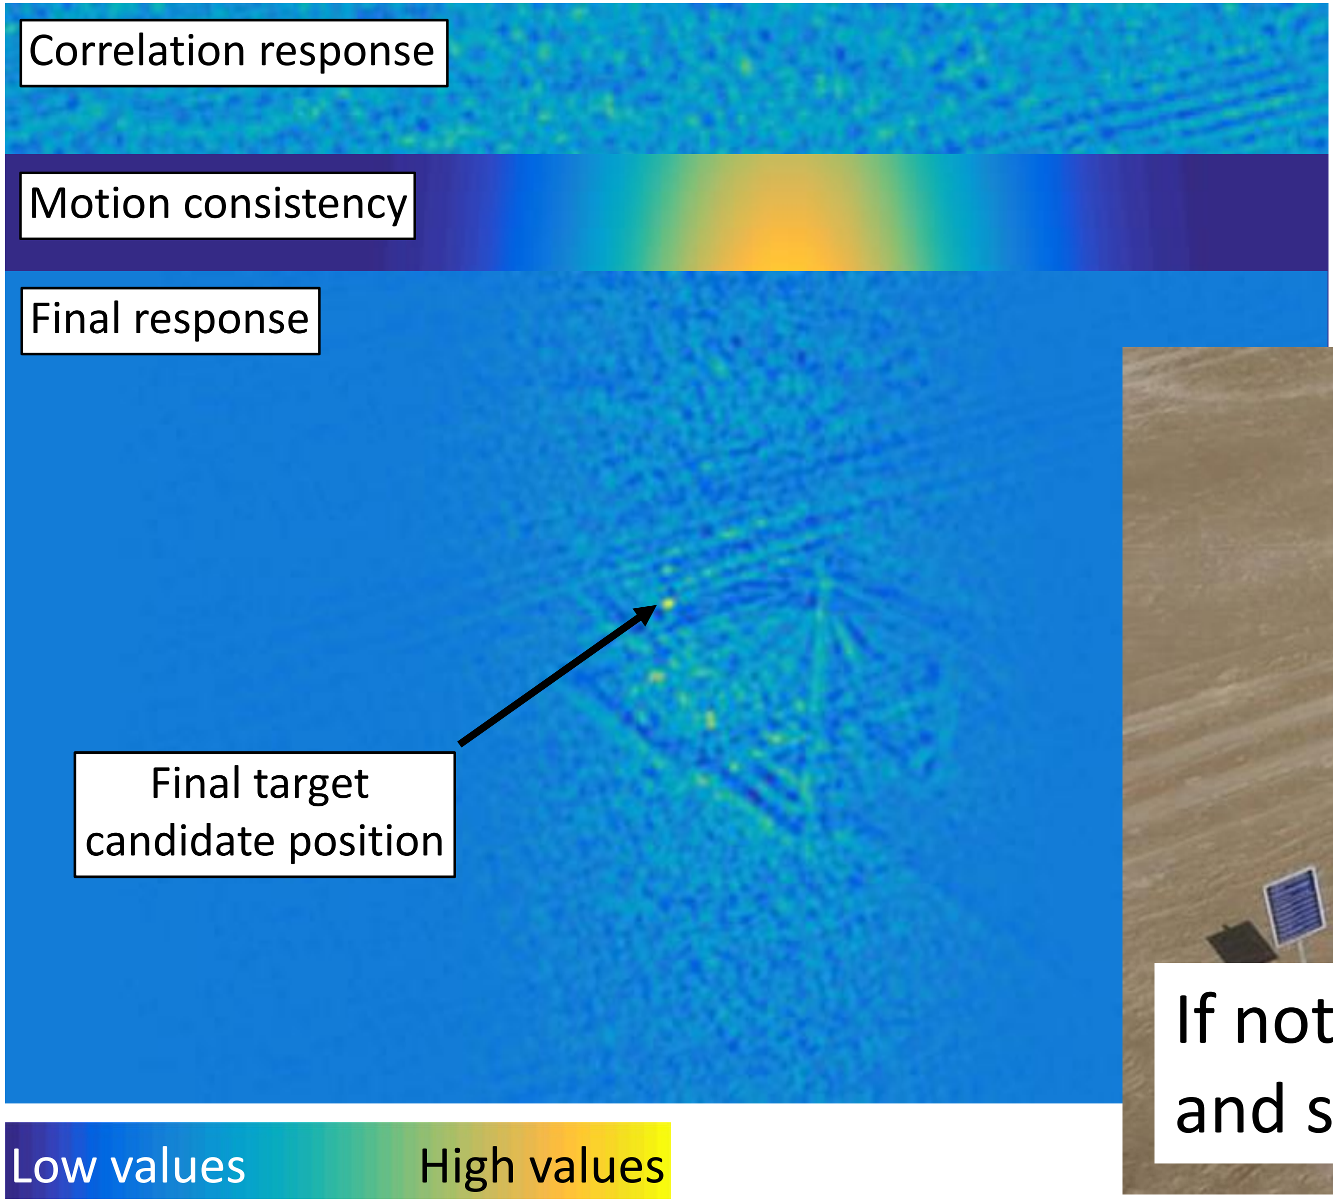

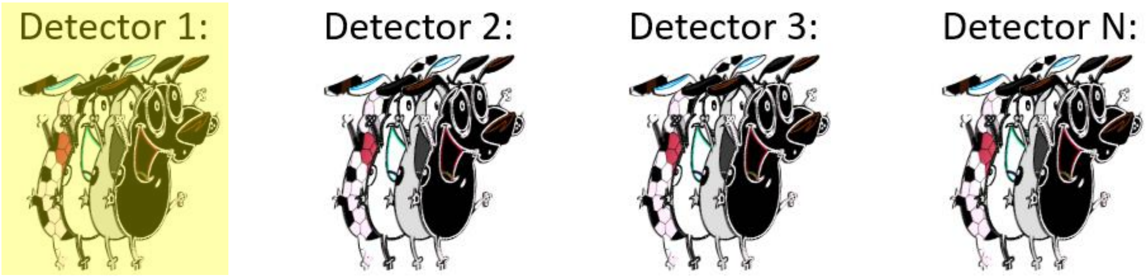

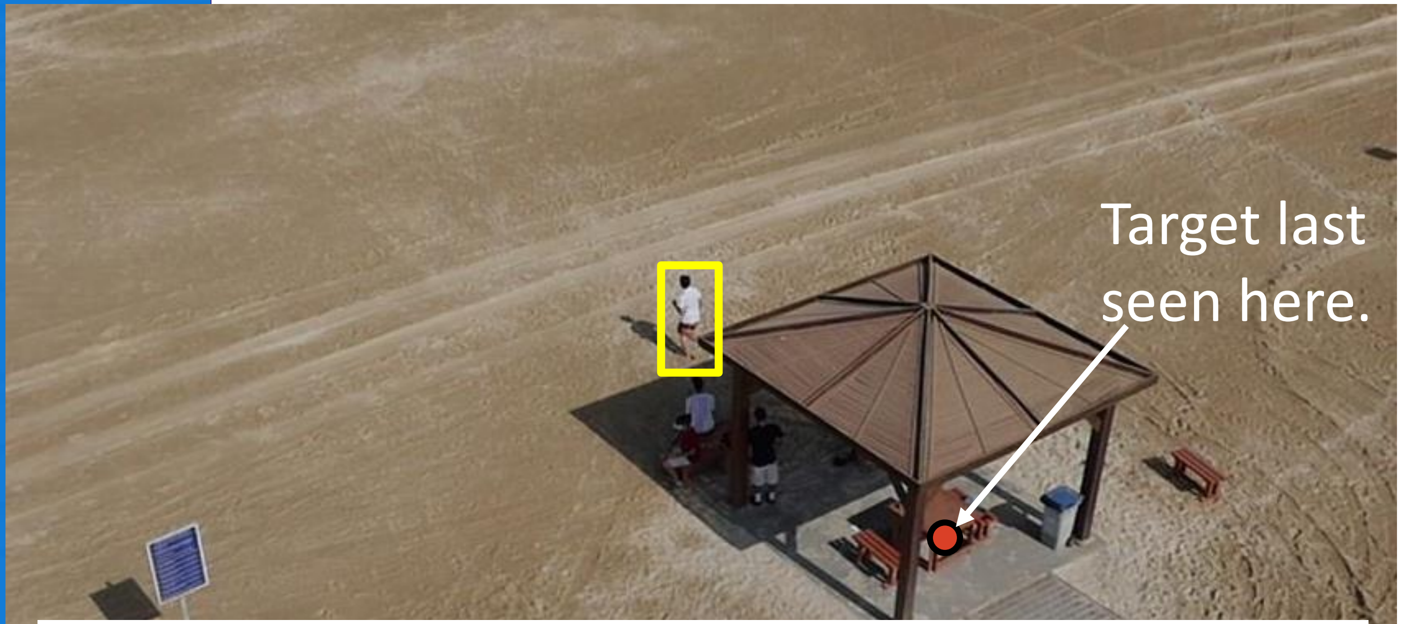

If not detected: Cycle through N detectors and scales in subsequent frames.

FCLT : ST tracking failure detecton

- Reliability score q_t
  on correlation response R_t^{ST}

$$q_t = \underline{MAX(R_t^{ST})} \times \underline{PSR(R_t^{ST})}$$

 $* \; H_t^{ST} = R_t^{ST}$ 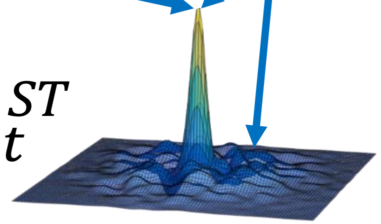

- Threshold on the ratio: $\dfrac{\overline{q_t}}{q_t}$
  $\overline{q_t}$ is mean over past frames
- When failure detected:
  - Activate detector
  - Stop updating visual model

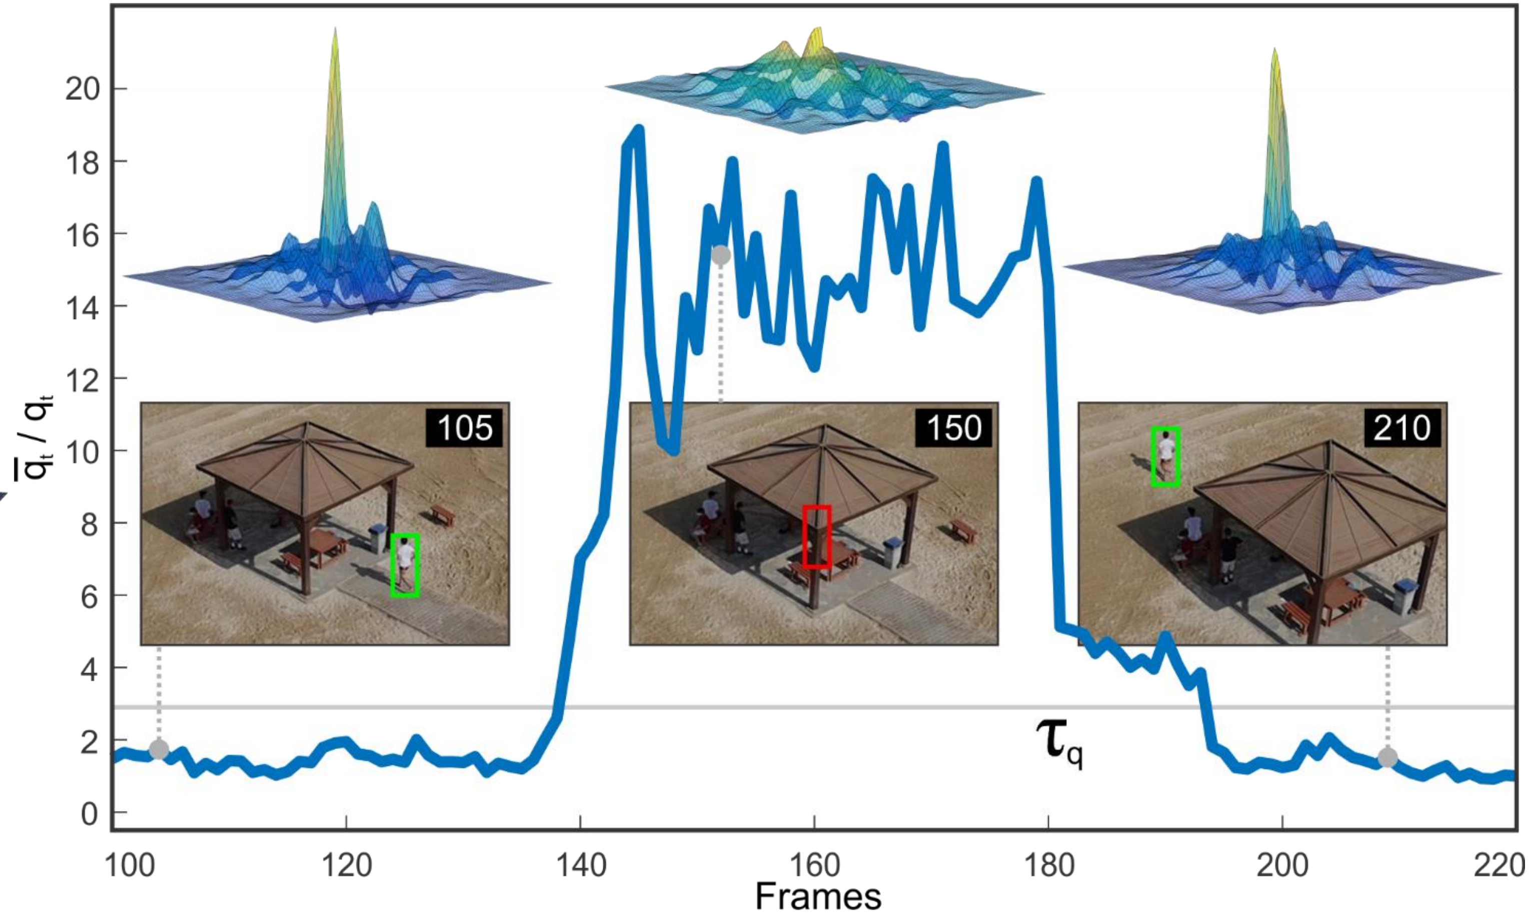

Example: Tracking with FCLT

Short-term tracker

Detector

Tracking uncertainty

Lukežič, Čehovin, Vojir, Matas, Kristan, *FuCoLoT -- A Fully-Correlational Long-Term Tracker*, ACCV 2018

Redetection capability (LT_0 vs LT_1)

FCLT[1]

Re-detects after target re-appears

MDNet[2]

Never recovers after drift

[1] Lukežič, Čehovin, Vojir, Matas, Kristan, *FuCoLoT -- A Fully-Correlational Long-Term Tracker*, ACCV2018
[2] Nam, Han, Learning, Multi-Domain Convolutional Neural Networks for Visual Tracking, CVPR2016

Extension of D3S to LT setup

- Similar to FCLT, only using DCF from GEM for global re-detection (and few additional upgrades, such as MDNet verifier)

Džubur et al., A Long-Term Discriminative Single Shot Segmentation Tracker, ERK2022

A 2D Object Assumption in Standard Trackers

- Existing tracking methods treat a tracked object as a 2D structure
- Problem: Cannot distinguish between pose change and (self)occlusion

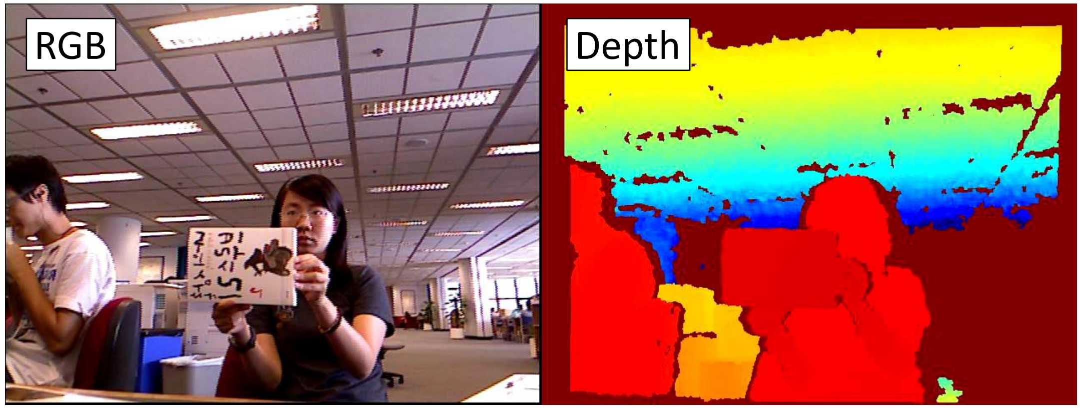

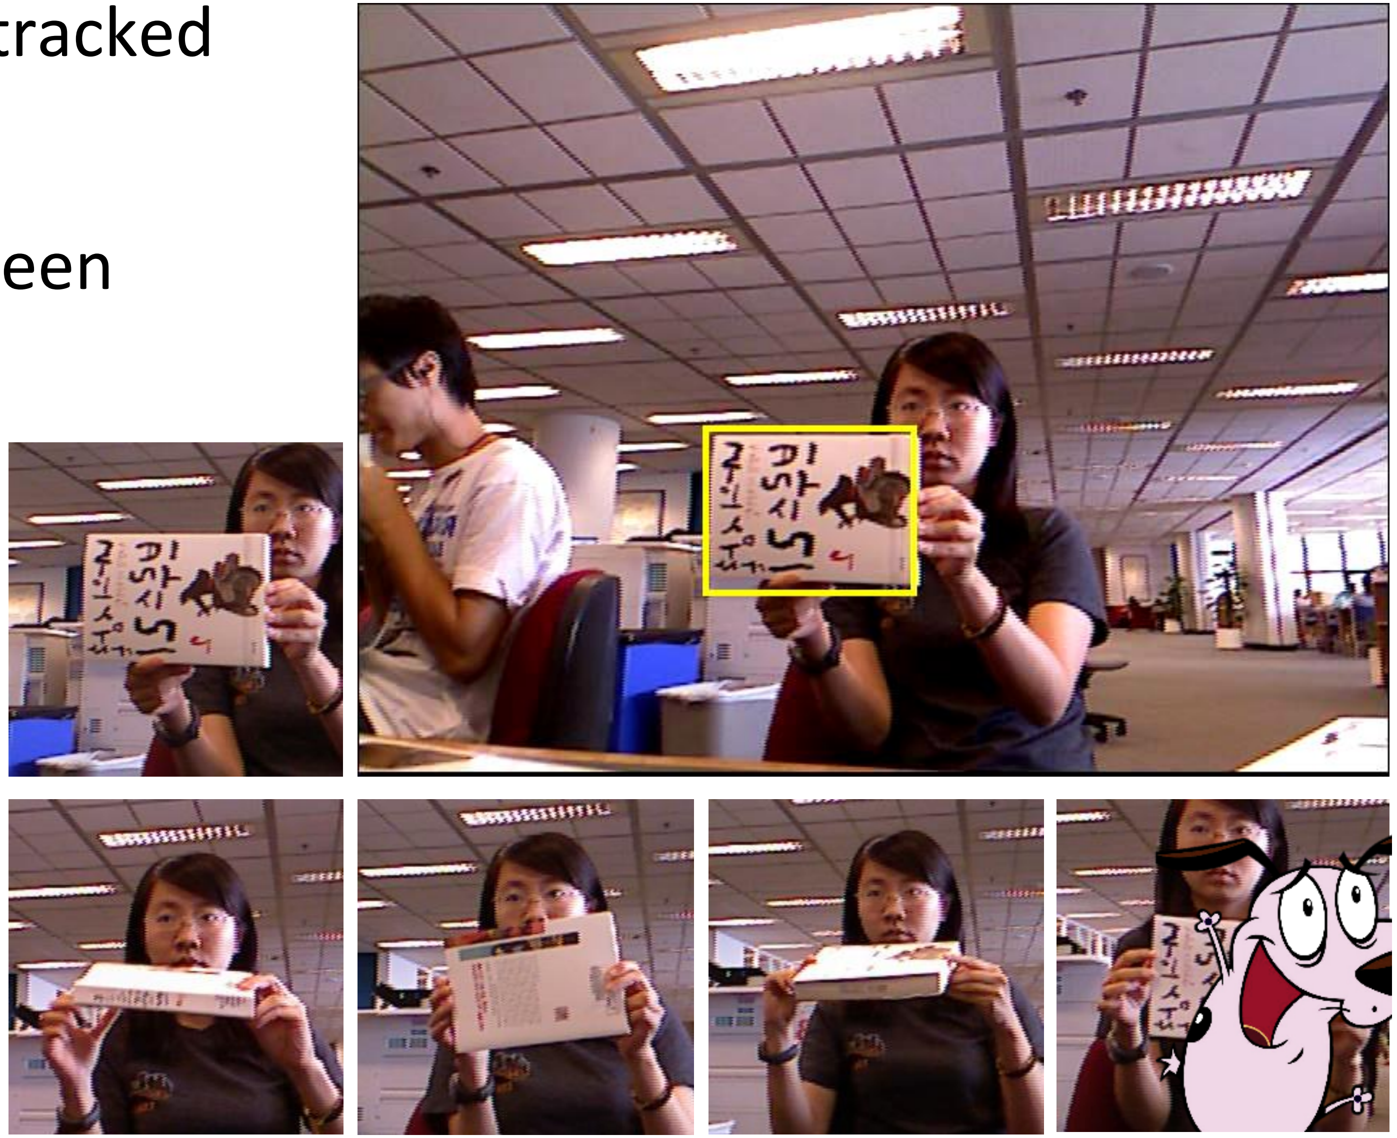

Extension to RGBD tracking

- Extend FCLT by 3D reconstruction to improve occlusion detection

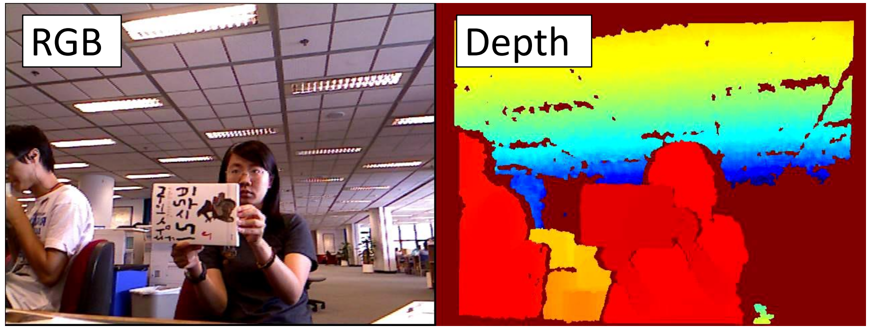

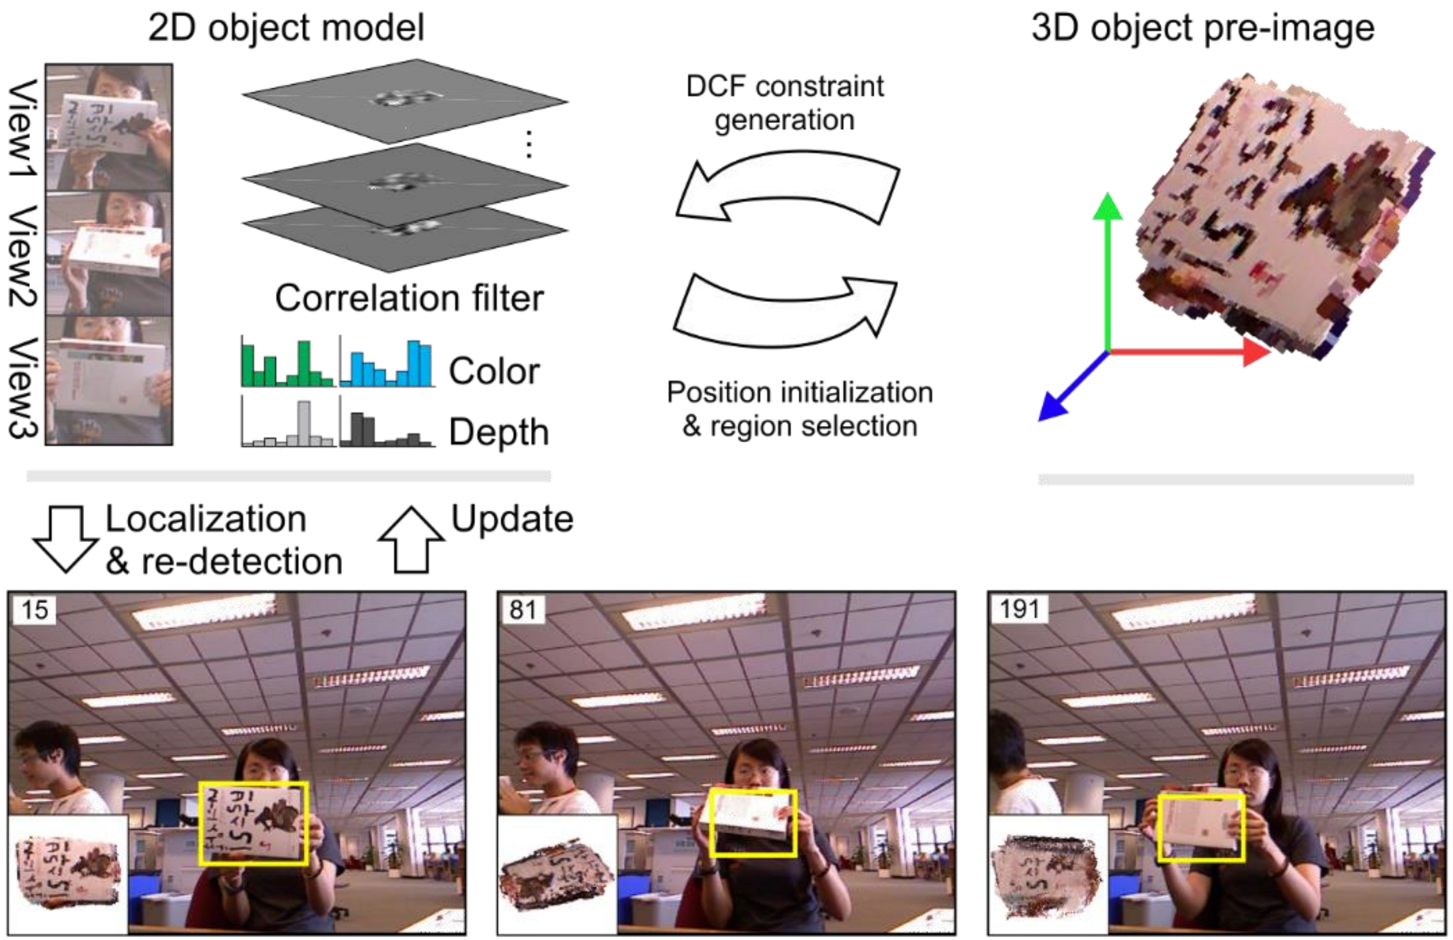

Kart, Lukezic, Kristan, Kämäräinen, Matas, Object Tracking by Reconstruction with View-Specific Discriminative Correlation Filters CVPR2019

Object tracking by reconstruction (OTR)

- Top performance among all RGBD trackers on PTB [Song et al., ICCV2013] and STC [Xiao et al.] benchmarks.

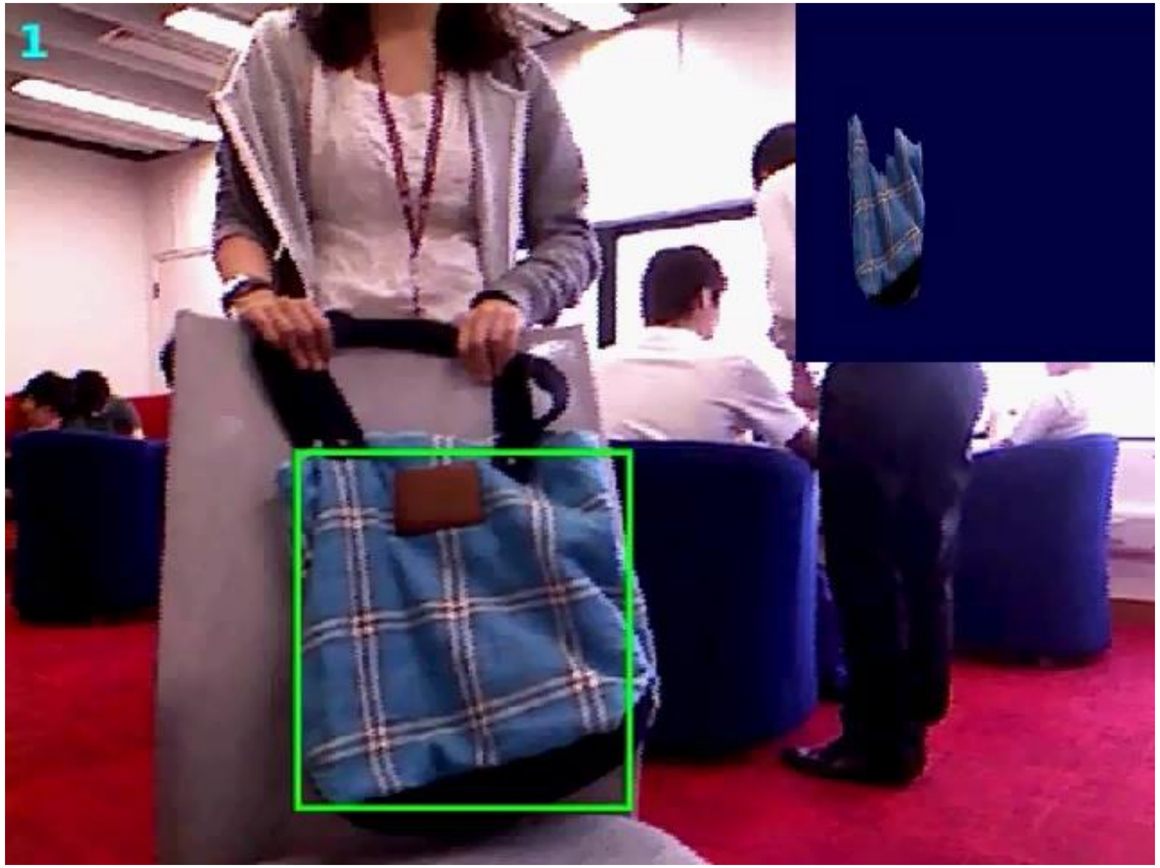

Kart, Lukezic, Kristan, Kämäräinen, Matas, Object Tracking by Reconstruction with View-Specific Discriminative Correlation Filters CVPR2019

Recent deep LT developments (2018)

https://github.com/xiaobai1217/MBMD

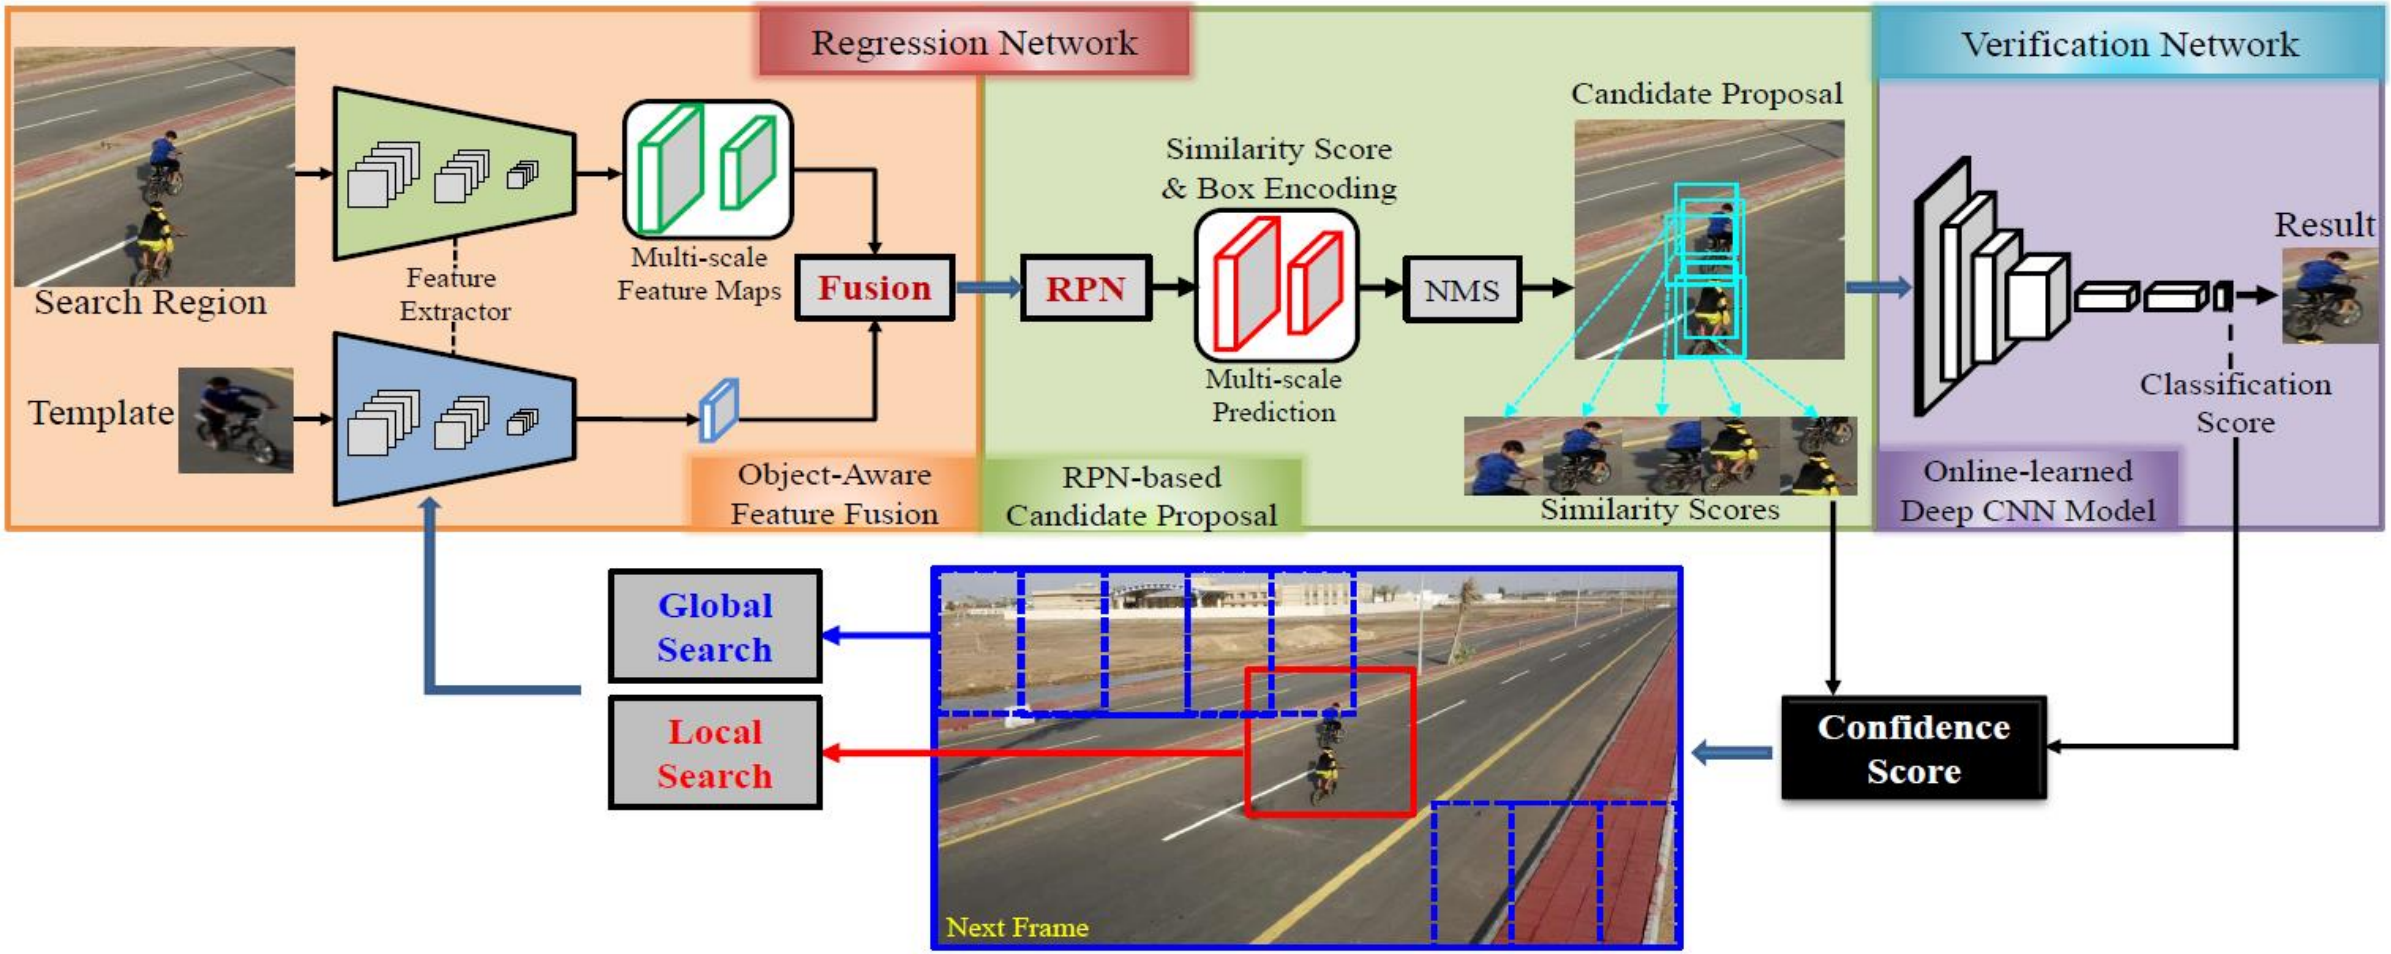

- Region proposal network akin to SSD[1] and SiamRPN[2]
- Verification network, essentially MDNet[3]
- Interaction akin to FCLT

[1]Liu et al., SSD: Single shot multibox detector, ECCV2016
[2]Li et al., High Performance Visual Tracking with Siamese Region Proposal Network, CVPR2018
[3]Nam et al., Learning multi–domain convolutional neural networks for visual tracking, CVPR2016

Zhang et al., Learning regression and verification networks for long-term visual tracking, ArXiv 2018

MBMD deep long-term tracker

- Modern state-of-the-art trackers are based on transformers (e.g., STARK-like) with a large localization range + a discriminator like Dimp

Zhang et al., Learning regression and verification networks for long-term visual tracking, ArXiv 2018

https://github.com/xiaobai1217/MBMD

References

- TLD:
  - Kalal, Z., Mikolajczyk, K. and Matas, J., Tracking-Learning-Detection, IEEE TPAMI2010
  - Page + code: http://personal.ee.surrey.ac.uk/Personal/Z.Kalal/
- FCLT:
  - Lukežič, Čehovin, Vojir, Matas, Kristan, *FuCoLoT -- A Fully-Correlational Long-Term Tracker*, ACCV 2018

Acknowledgment

- Thanks to Jiri Matas for kindly sharing some of their slides that I used in preparation of this lecture.