

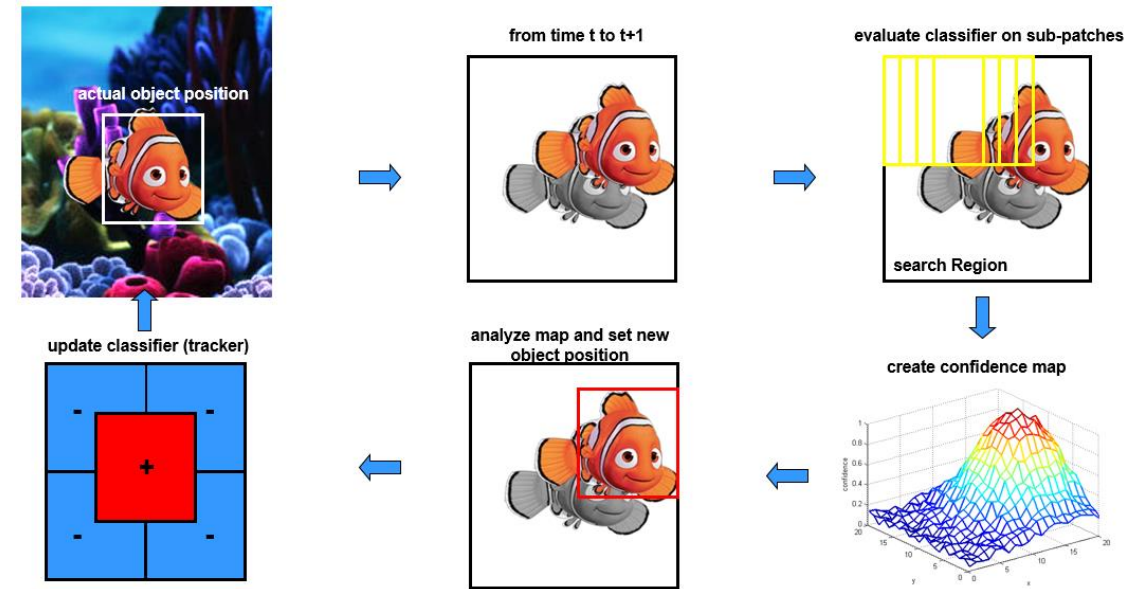
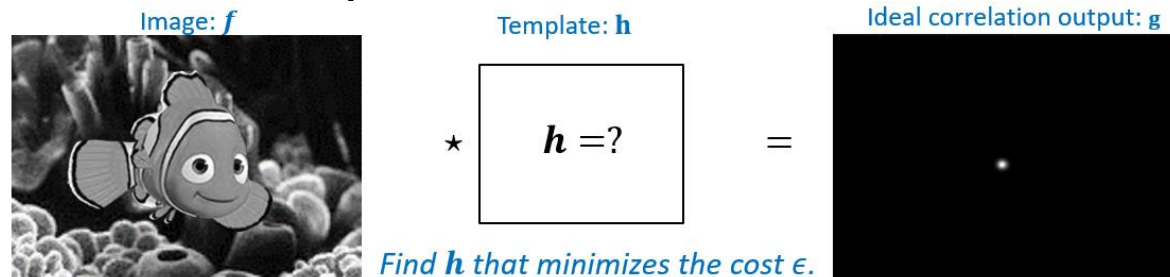
Previously at ACVM...

- Discriminative trackers

- Adaboost
- TMIL
- Structured SVM

- Discriminative correlation filters

- Linear classifiers (ridge regression learning)
- Efficient computation via FFT





Advanced computer vision methods

Tracking by Recursive Bayes Filters

Part I: Introduction

Matej Kristan

Laboratorij za Umetne Vizualne Spoznavne Sisteme,
Fakulteta za računalništvo in informatiko,
Univerza v Ljubljani

Classes of trackers

- A tracker can be roughly classified by considering the following two properties:

Property 1: Batch tracking vs. Online tracking

- How many images are considered to estimate the state at time-step t ?

Property 2: Non Bayesian vs. Bayesian tracking

- How is the notion of the target state encoded?

Online vs Batch tracking

- Batch tracking: Can consider all frames before t and after t to infer the target position at time-step t .



Potentially robust, appropriate for offline systems

- Online tracking: Can consider only frames before t to infer position at t .



Potentially fast, appropriate for real-time systems

Non Bayesian vs Bayesian tracking

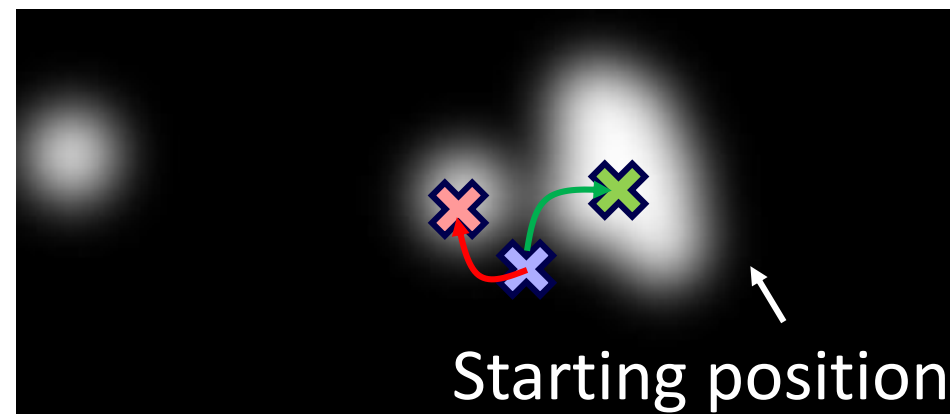
Question: “How is the information of the target state encoded?”

NON-BAYESIAN

- Local optimum
 - Gradient descent
 - Mean Shift
 - Greedy local search, etc.
- Fast convergence
- Single solution for the state value
 - But is it correct?



input image



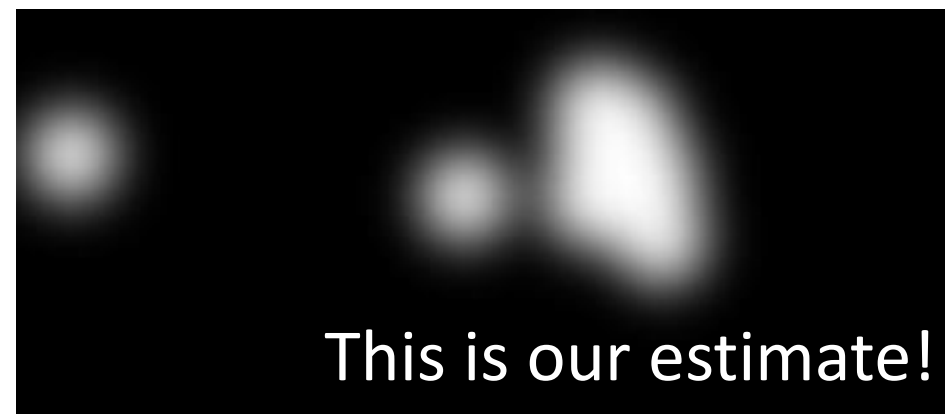
similarity/probability

Non Bayesian vs Bayesian tracking

Question: “How is the notion of the target state encoded?”

BAYESIAN

- Assign a **probability** to each position of the target
- Relevant info is encoded in the **pdf over the target state!**
- Implicitly remembers **multiple hypotheses** of “location”.
- Interpret the pdf when required
- Typically slower



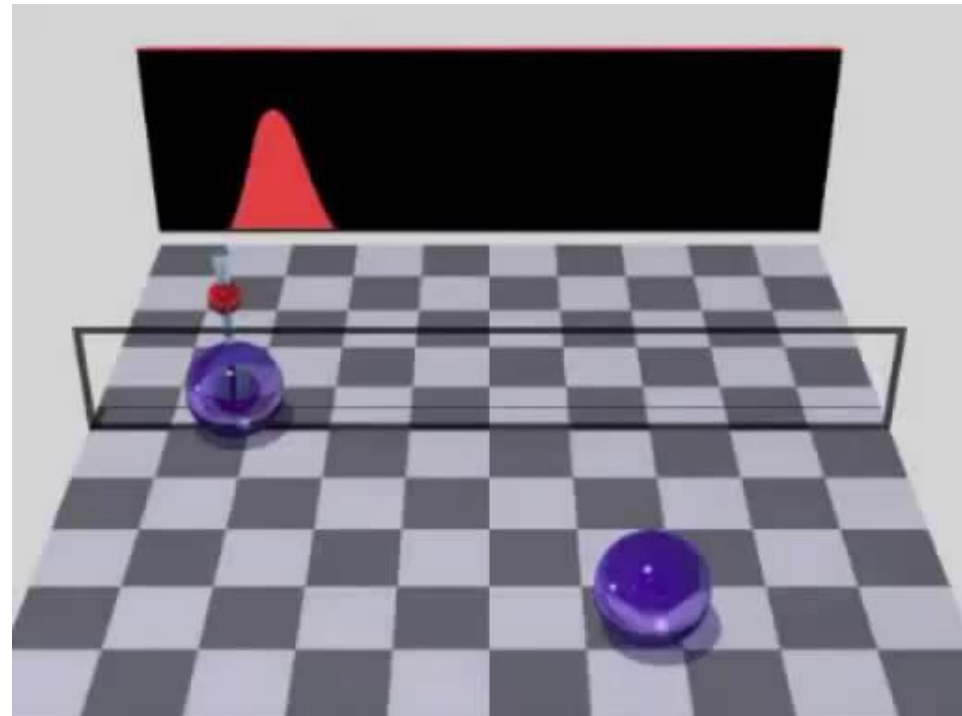
similarity/probability

Non Bayesian vs Bayesian tracking

Question: “How is the notion of the target state encoded?”

BAYESIAN

- The pdf changes with time – an **entire pdf is tracked**
- Example of a pdf:
 - $p(Ball|x_k)$,
 - Expected value:
 - $\hat{x} = \langle x_k \rangle_{p(Ball|x_k)}$



Examples: Online tracking

- Non Bayesian
 - Mean Shift



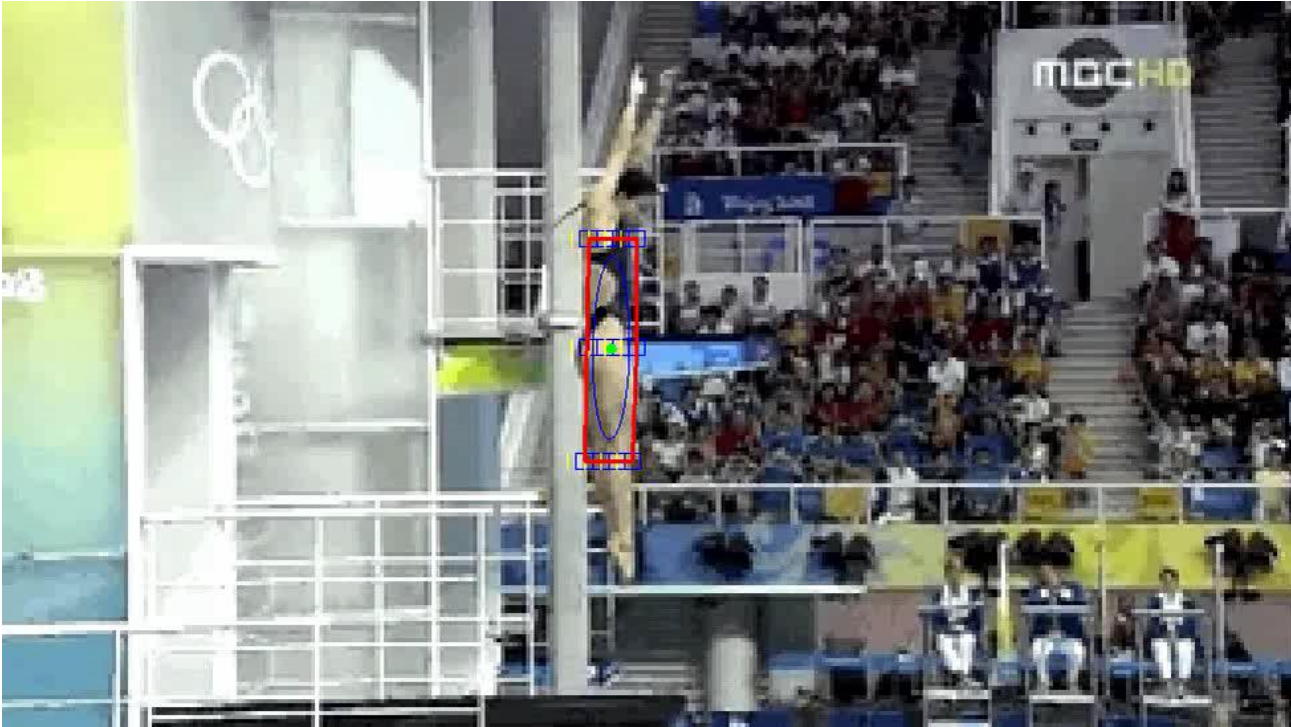
Comaniciu et al. "Kernel-Based Object Tracking",
IEEE TPAMI., 2003

- Fully Bayesian
 - Bayes recursive filters



Isard et al., "CONDENSATION -- conditional density
propagation for visual tracking" IJCV, 1998

Practical challenges



➔ Change in appearance

- Level of detail
- Occlusion by visually similar objects
- Clutter
- Target motion
- Interacting targets

Practical challenges



- Change in appearance
- ➔ Level of detail
- Occlusion by visually similar objects
- Clutter
- Target motion
- Interacting targets

Kristan et. al, "Closed-world tracking of multiple interacting targets for indoor-sports applications", CVIU2009

Practical challenges



- Change in appearance
- Level of detail
- ➔ Occlusion by visually similar objects
- Clutter
- Target motion
- Interacting targets

M. Kristan et al., "A Local-motion-based probabilistic model for visual tracking", *Pattern Recognition*, 2009

Practical challenges



- Change in appearance
- Level of detail
- Occlusion by visually similar objects
- ➔ Clutter
- Target motion
- Interacting targets

Z. Khan, T. Balch, and F. Dellaert, "MCMC-Based Particle Filtering for Tracking a Variable Number of Interacting Targets" IEEE TPAMI, 2005.

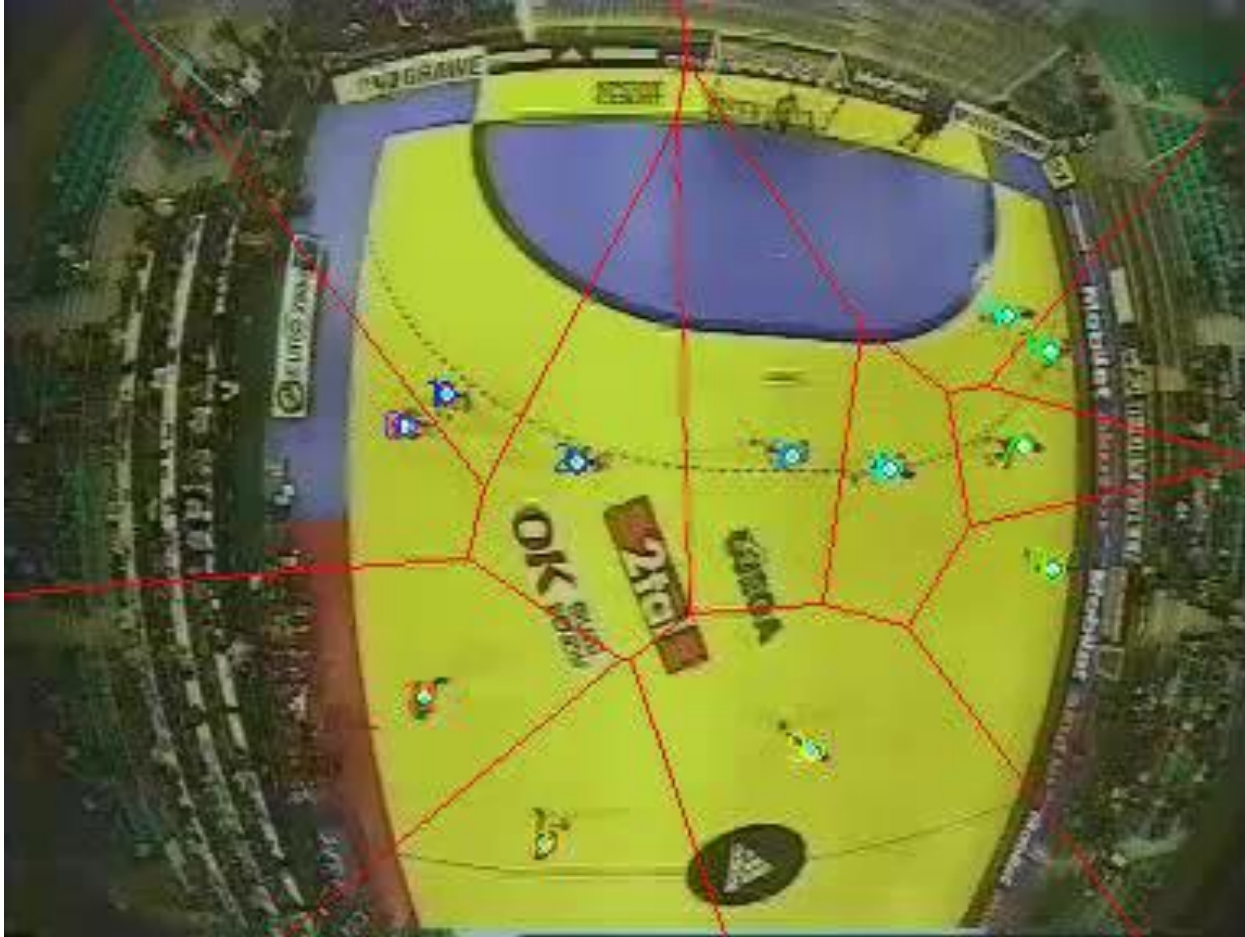
Practical challenges



- Change in appearance
- Level of detail
- Occlusion by visually similar objects
- Clutter
- ➔ Target motion
- Interacting targets

Isard, M. and Blake, A. "CONDENSATION --conditional density propagation for visual tracking", IJCV1998

Practical challenges

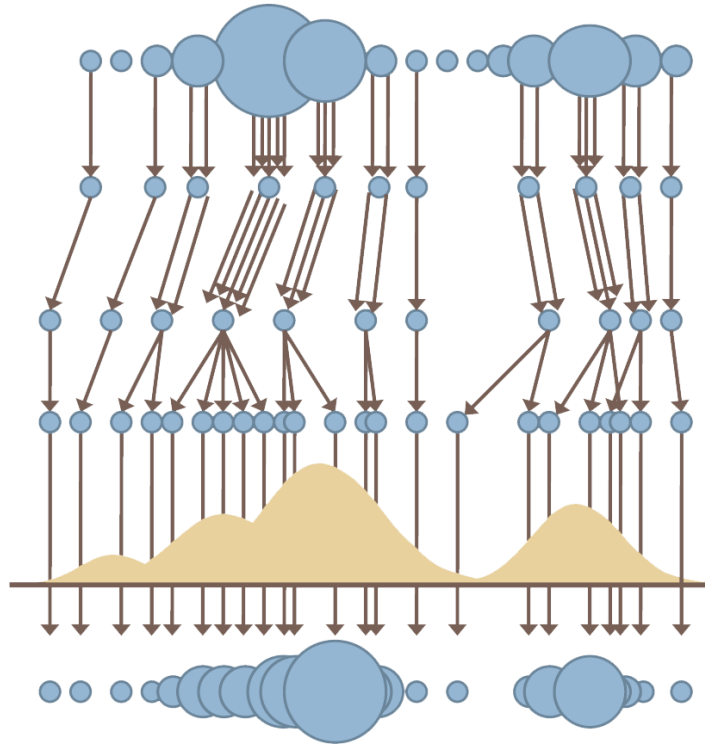


- Change in appearance
- Level of detail
- Occlusion by visually similar objects
- Clutter
- Target motion
- ➔ Interacting targets

Kristan et al., "Closed-world tracking of multiple interacting targets for indoor-sports applications" *CVIU* 2009.

Practical challenges

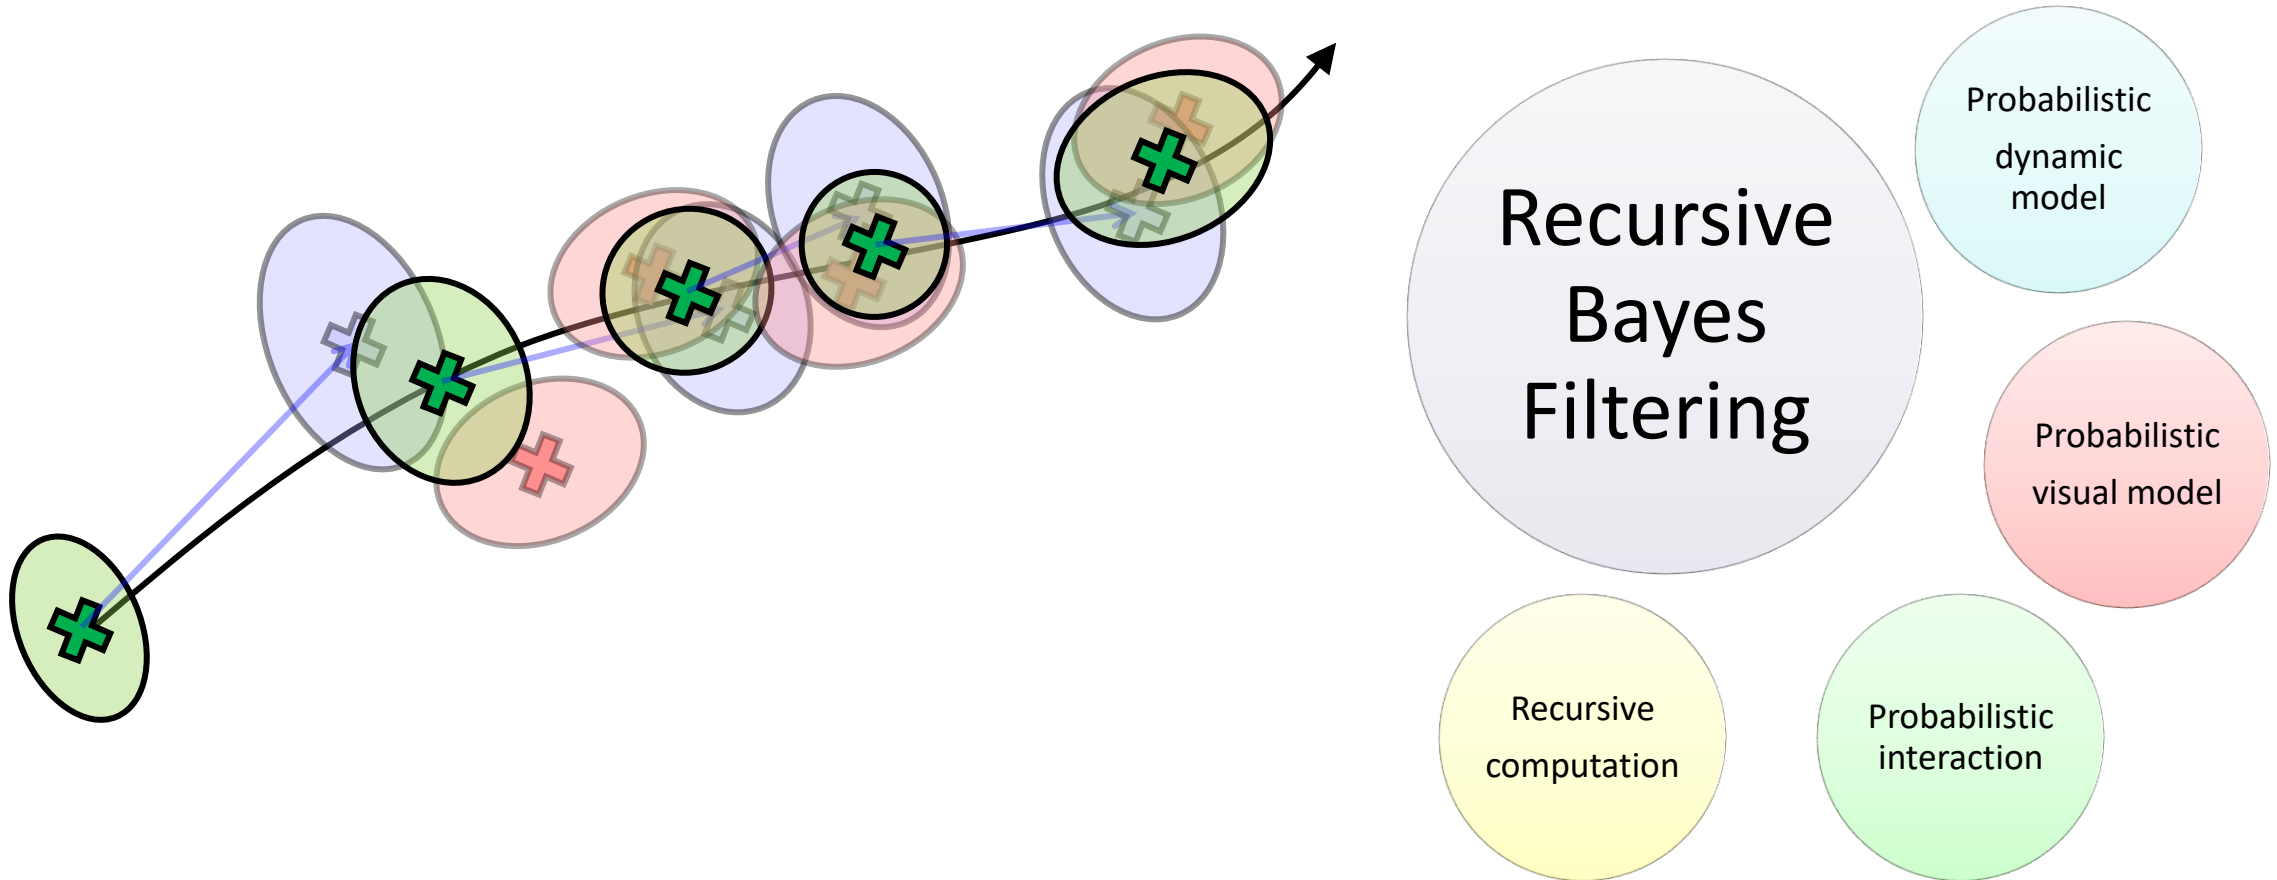
These issues can be addressed efficiently using probabilistic approaches!



- ➔ Change in appearance
- ➔ Level of detail
- ➔ Occlusion by visually similar objects
- ➔ Clutter
- ➔ Target motion
- ➔ Interacting targets

Recursive Bayes Filters

- A principled way to address uncertainty in visual tracking



Consider tracking an airplane as an example

Observe a scene at $t-1$



Observe a scene at t



Observed scene at t

Plane at $t-1$



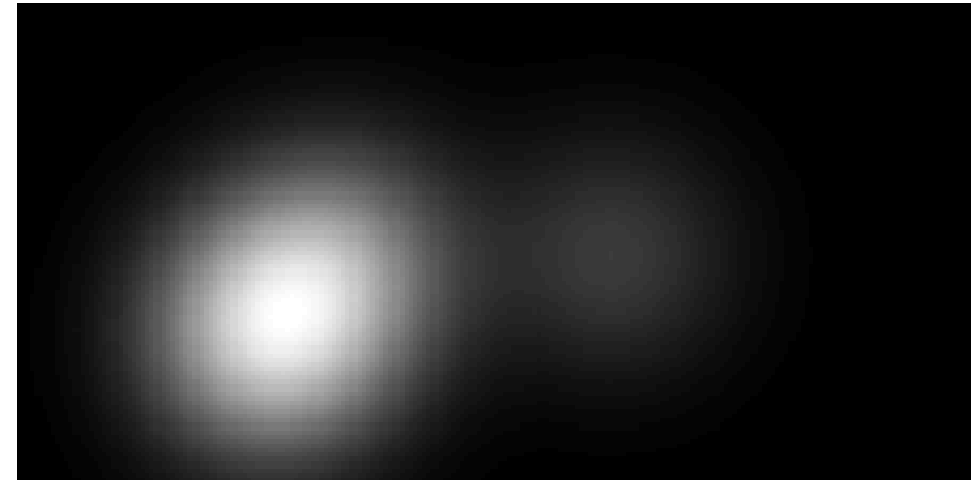
Bayesian tracking as a state estimation problem

- State at time t : x_t (e.g., position)
- Measurement at t : y_t (e.g., location obtained by detector)
- **Approach:** Given all we know about the target and the measurements we take, what is the probability that a target is at state x_t ?

From this:
Observed scene at t



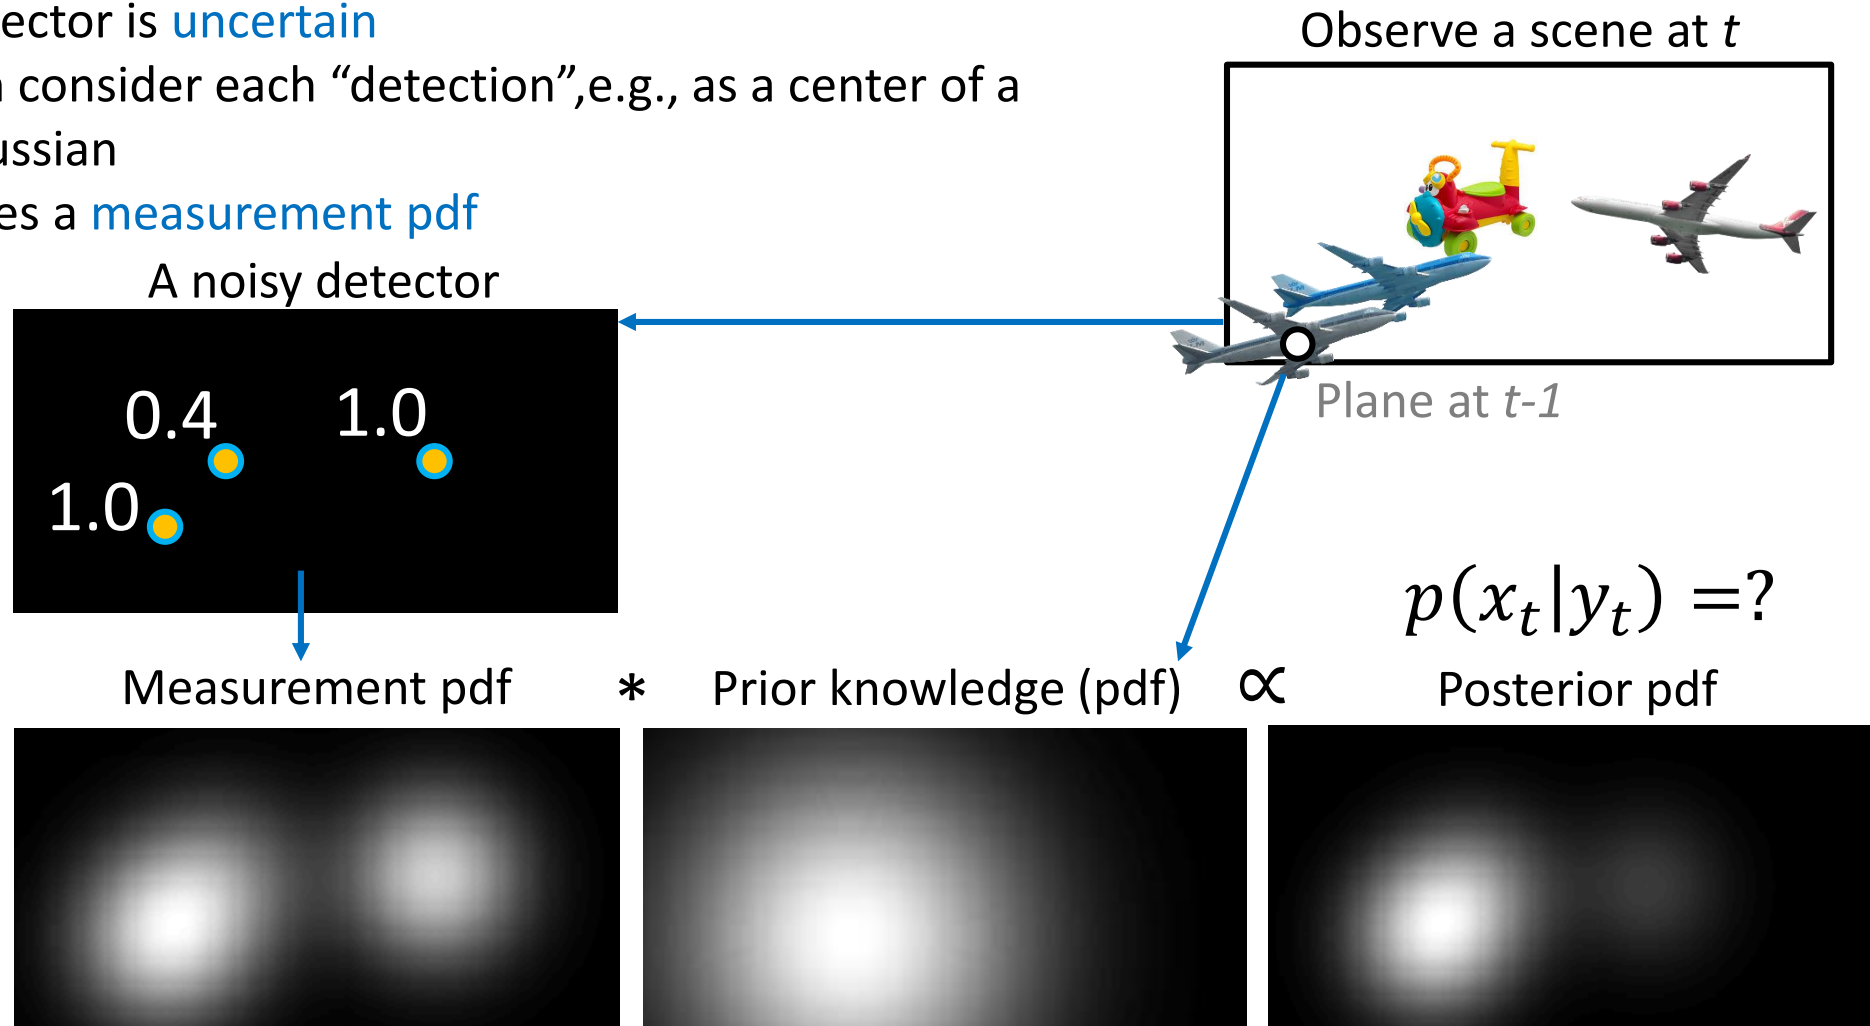
Infer this: $p(x_t|y_t) = ?$



What is a Bayes Filter?

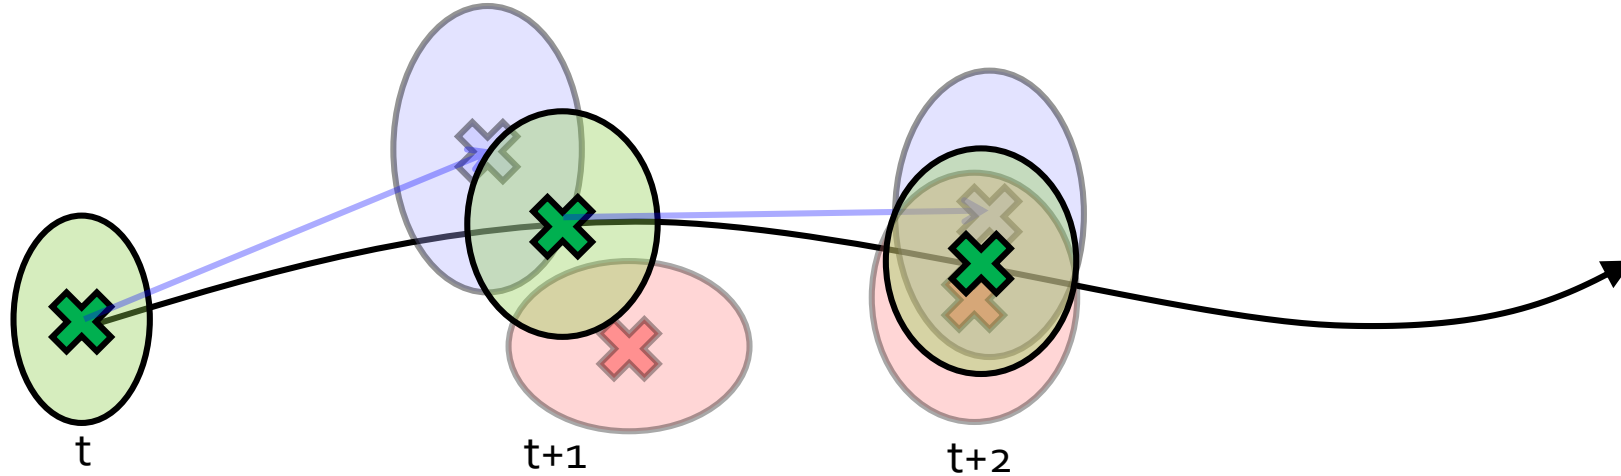
- **Key idea 1:** Reason about the target states in terms of pdfs

- Detector is **uncertain**
- Can consider each “detection”, e.g., as a center of a Gaussian
- Gives a **measurement pdf**

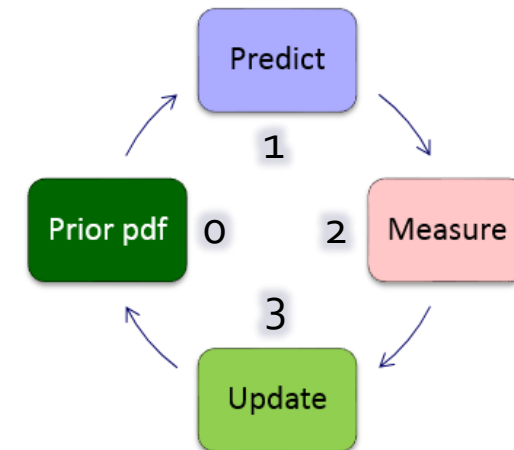


What is a *Recursive Bayes Filter*?

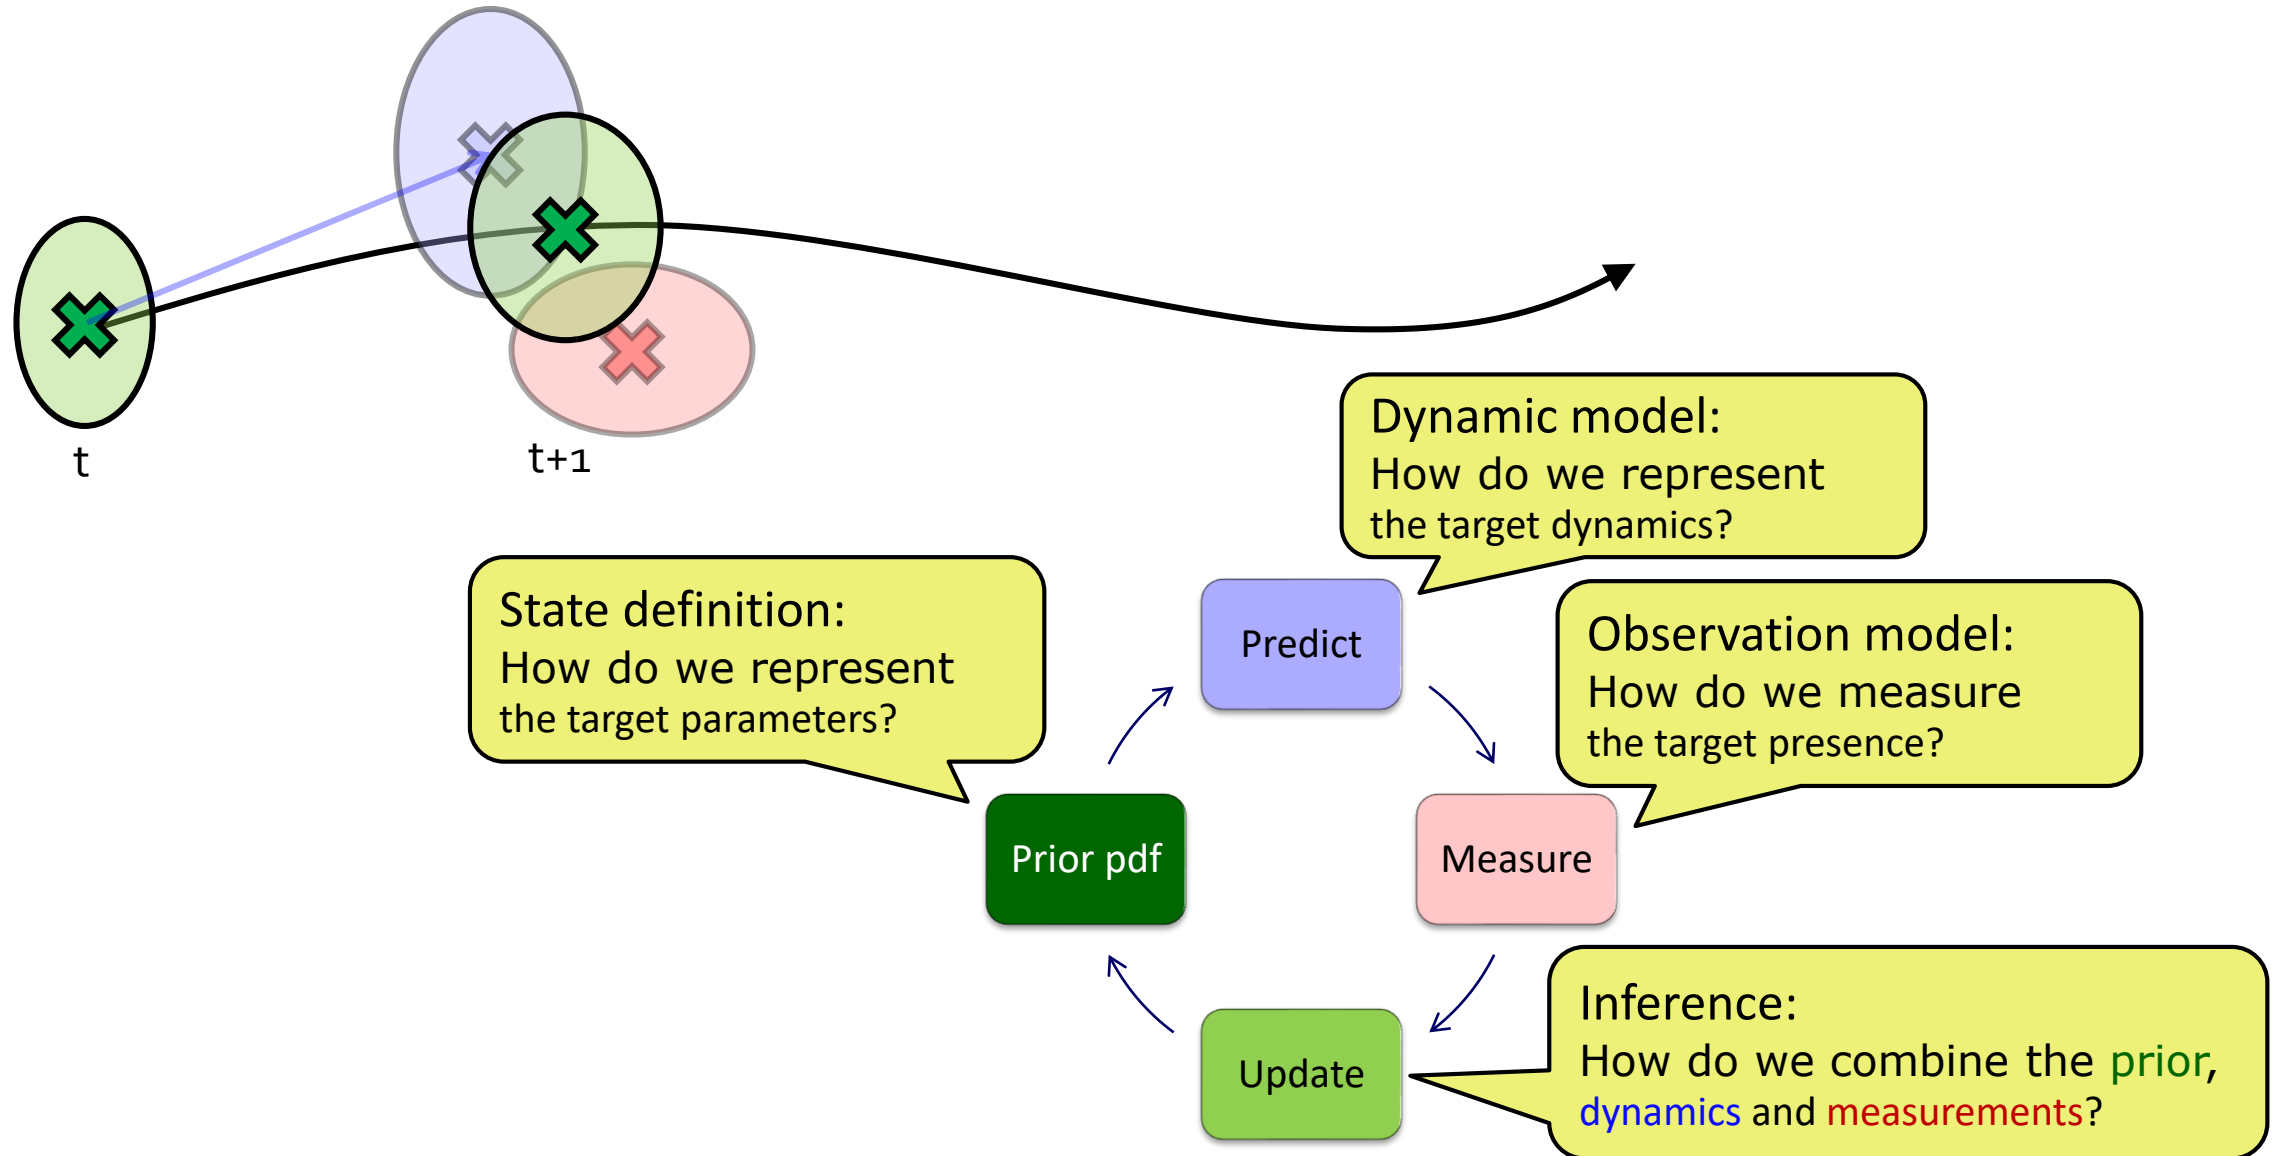
- **Key idea 1:** Encode beliefs about states in a pdf



- **Key idea 2:** Recursively estimate the posterior
 - **Predict** from *uncertain* motion model
 - **Measure** from *uncertain* sensor
 - **Update** distribution



Recursive Bayes Filter: Key ingredients

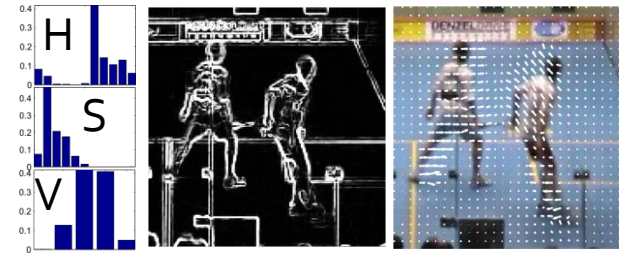


Key Ingredients

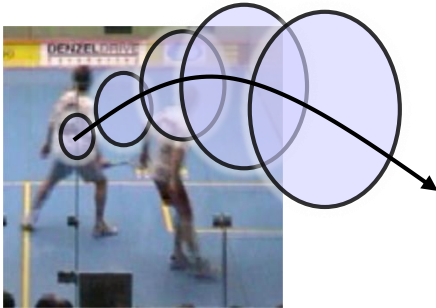
State definition



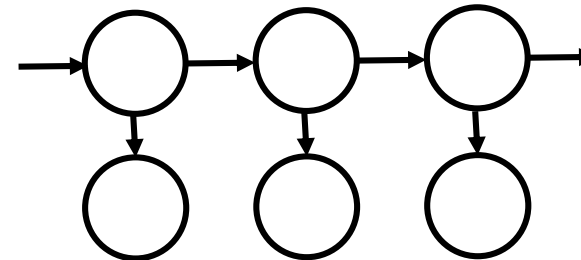
Observation model



Dynamic model



Inference

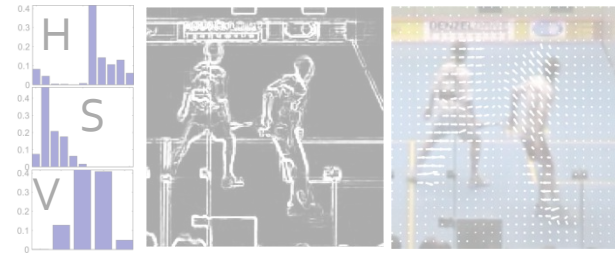


Key Ingredients

State definition



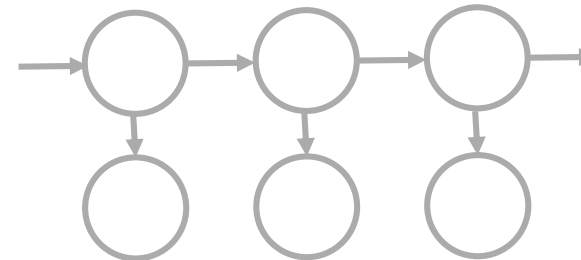
Observation model



Dynamic model



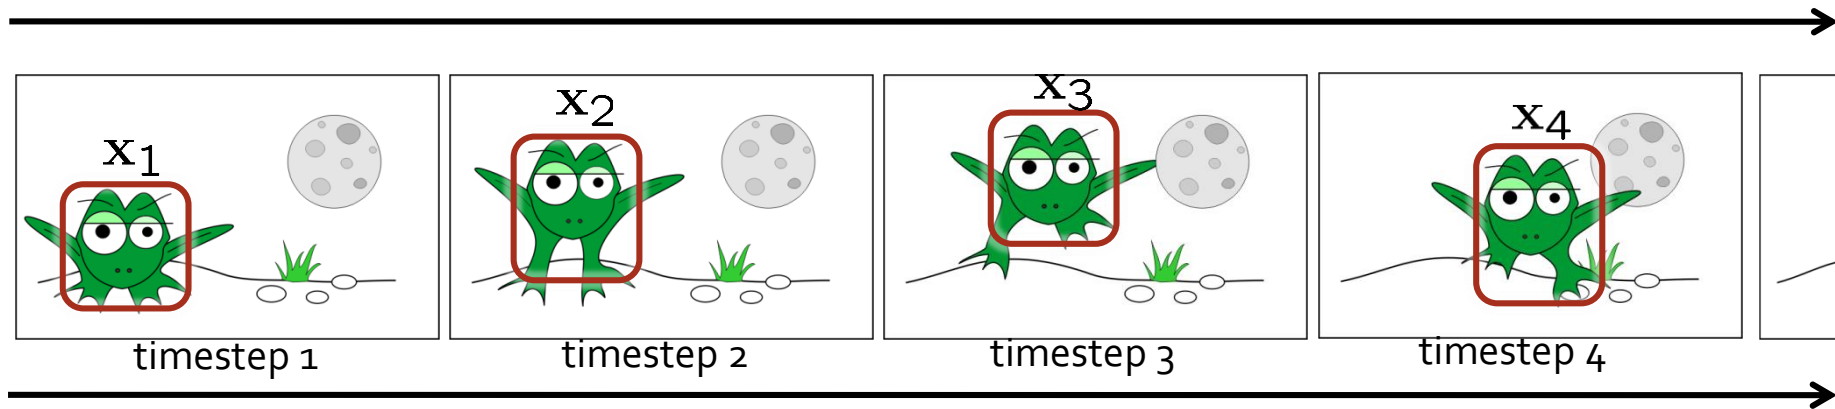
Inference



What is a state of the target?

- Target properties at a time-step
- Encodes parameters (which we want to estimate)

$$\mathbf{x}_{1:N} = \{\mathbf{x}_1, \dots, \mathbf{x}_{k-1}, \mathbf{x}_k, \mathbf{x}_{k+1}, \dots, \mathbf{x}_N\} = \{\mathbf{x}_k\}_{1:N}$$



- **Parametric form** of the state \mathbf{x}_k depends on the model by which we describe the target.

How do we define a state?

- Define the state by the target “free” parameters.

- Examples:

- Location

$$\mathbf{x}_k = [x, y]$$

- Location + size

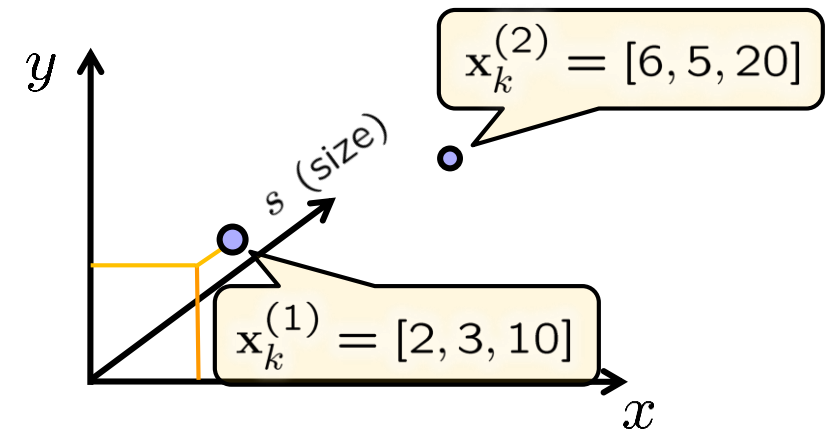
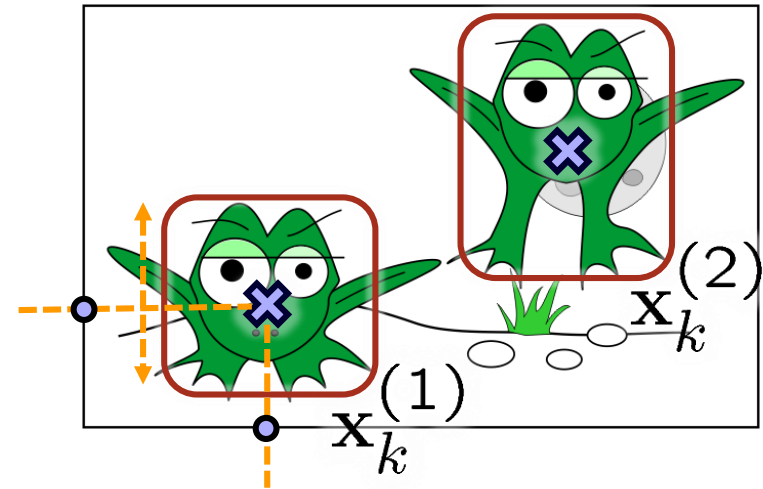
$$\mathbf{x}_k = [x, y, s]$$

- Location + velocity

$$\mathbf{x}_k = [x, y, \dot{x}, \dot{y}]$$

- Multiple objects (joint state)

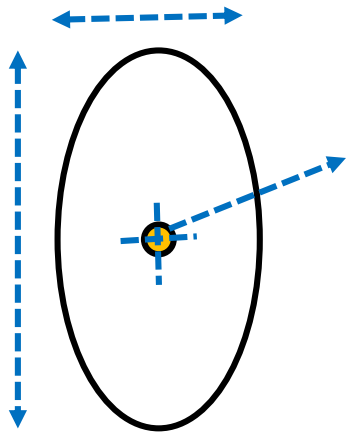
$$\mathbf{x}_k = \{\mathbf{x}_k^{(1)}, \mathbf{x}_k^{(2)}\}$$



State definition: Example 1

- Axis-aligned blobs (bounding box, ellipse)
 - center
 - width + height
 - velocity

$$6D \quad \mathbf{x}_k = [x, y, \dot{x}, \dot{y}, H_x, H_y]$$

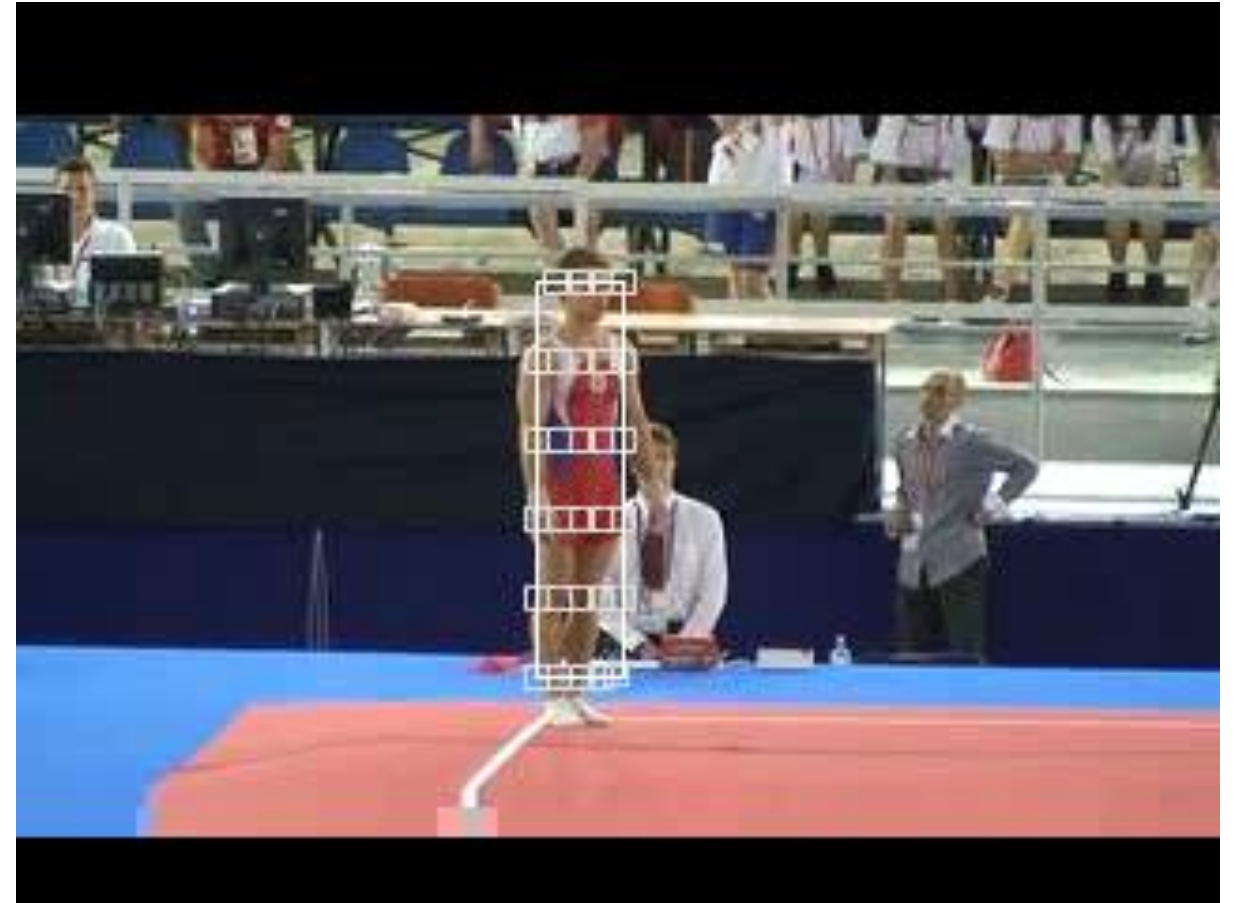
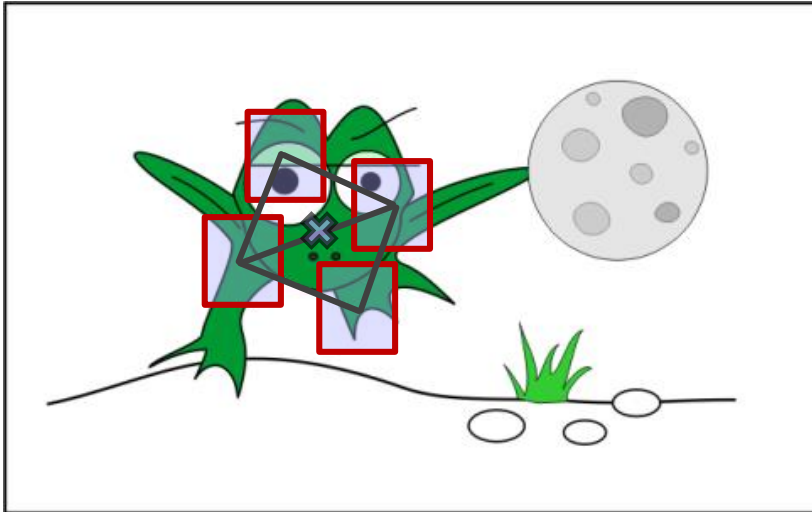


Kristan et al., "A Local-motion-based probabilistic model for visual tracking". *Pattern Recognition*, 2009.

State definition: Example 2

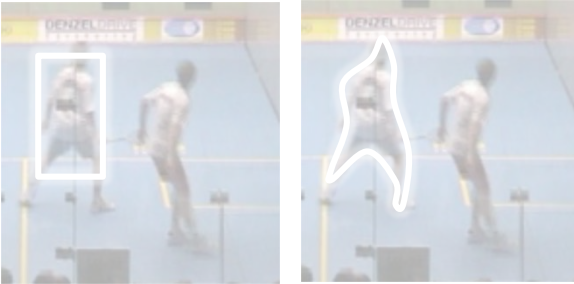
- Part-based models, Constellation models
 - Center, velocity
 - Relative part locations
 - Varying number of parts

$$\mathbf{x}_k = [x_c, y_c, \dot{x}_c, \dot{y}_c, \mathbf{x}_k^{(1)}, \mathbf{v}_k^{(1)}, \dots, \mathbf{x}_k^{(N_k)}, \mathbf{v}_k^{(N_k)}]$$

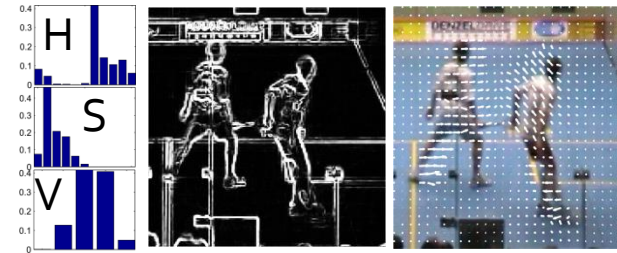


Key Ingredients

State definition



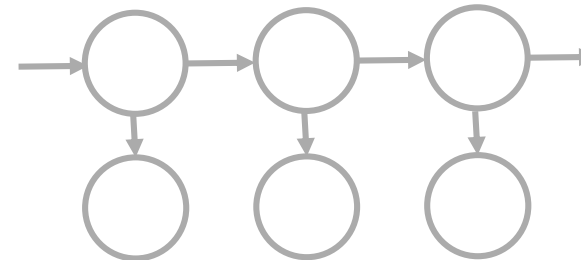
Observation model



Dynamic model

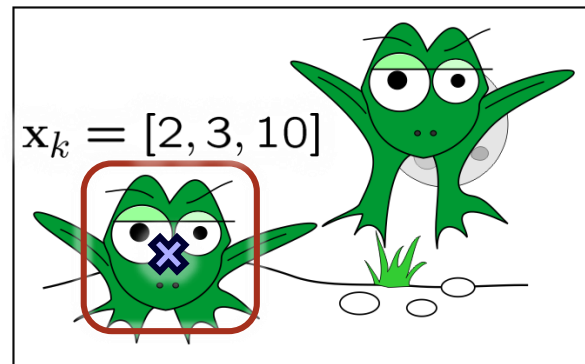


Inference



What is the observation model?

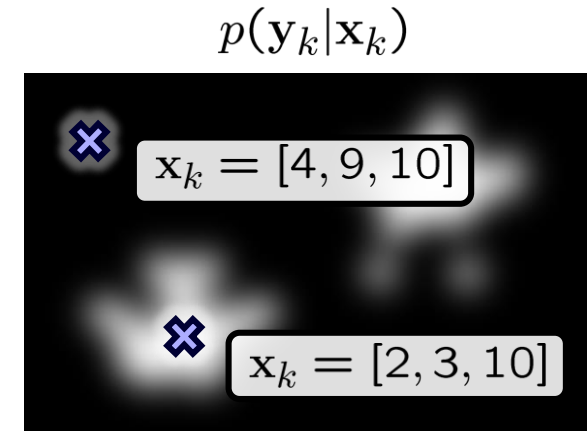
- Transforms measurement into a probability
- The likelihood of observing y_k assuming the target is located at state x_k : $p(y_k|x_k)$



observed image

Low likelihood.

High likelihood.



likelihood map

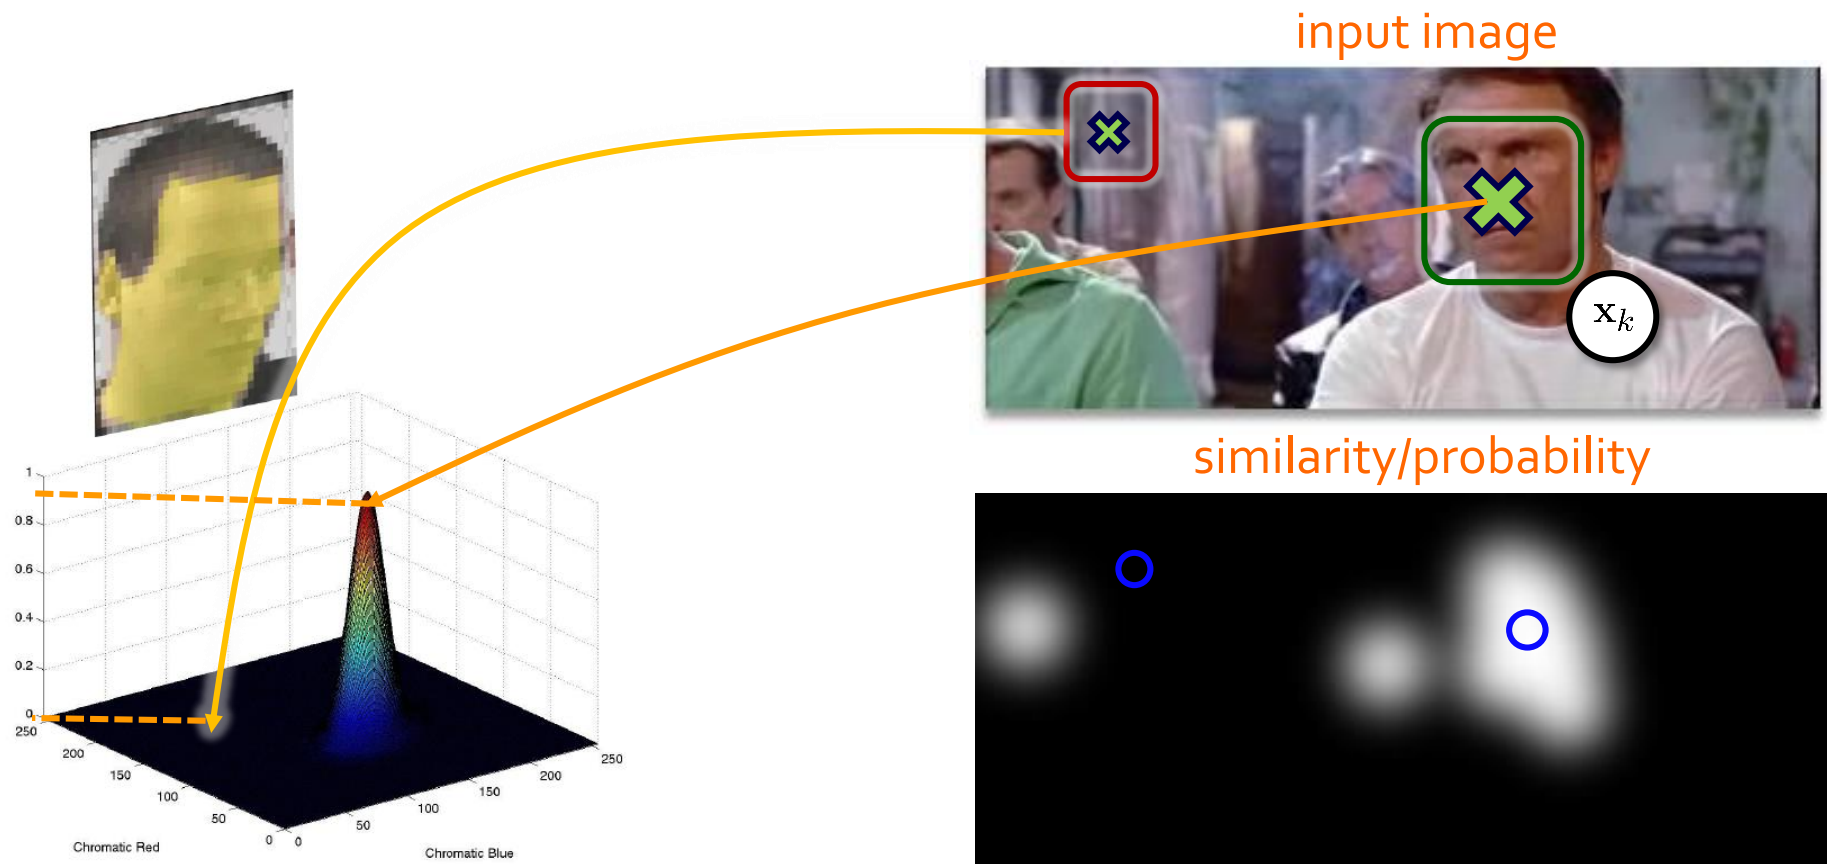
If the observation model was evaluated at each possible state x_k ...

Observation model

1. Choose a visual model
(e.g., histograms, HOG, template, ...)
2. Define similarity function with the visual model
3. Define a function that maps similarity to probability
(i.e., zero similarity \rightarrow zero probability and vice versa)

Observation model: Example 1

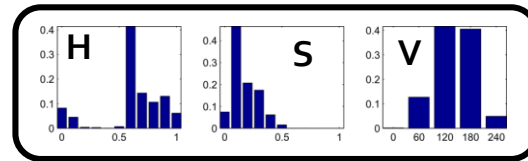
- Skin color sampled from a region
 - clusters in chromatic space – model by a Gaussian $p(y_k | x_k) \propto \exp(-\frac{1}{2}(y_k - \mu)^T \Sigma^{-1}(y_k - \mu))$



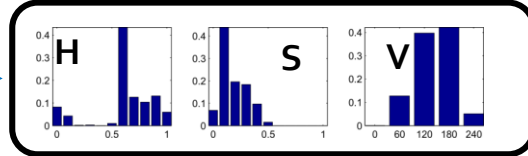
Observation model: Example 2

- Histograms
 - Color histograms
 - Hellinger distance between reference h_r and sampled histogram h_s : $d_{Hell}(h_r, h_s)$

reference model h_r



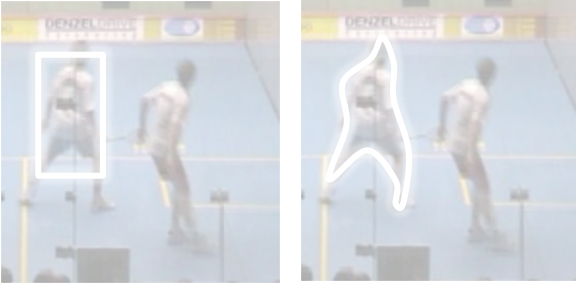
sampled model h_s



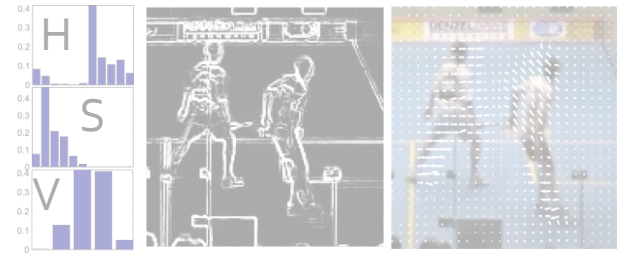
$$p(y_k | x_k) \propto \exp\left(-\frac{1}{2} d_{Hell}^2 / \sigma^2\right)$$

Key Ingredients

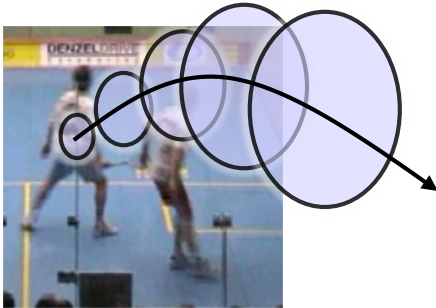
State definition



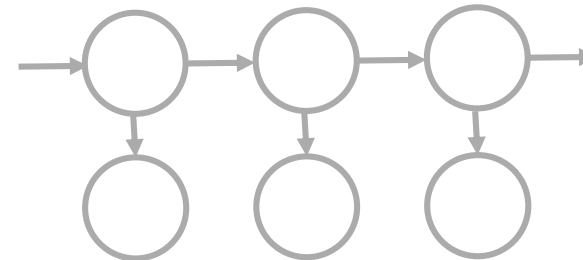
Observation model



Dynamic model

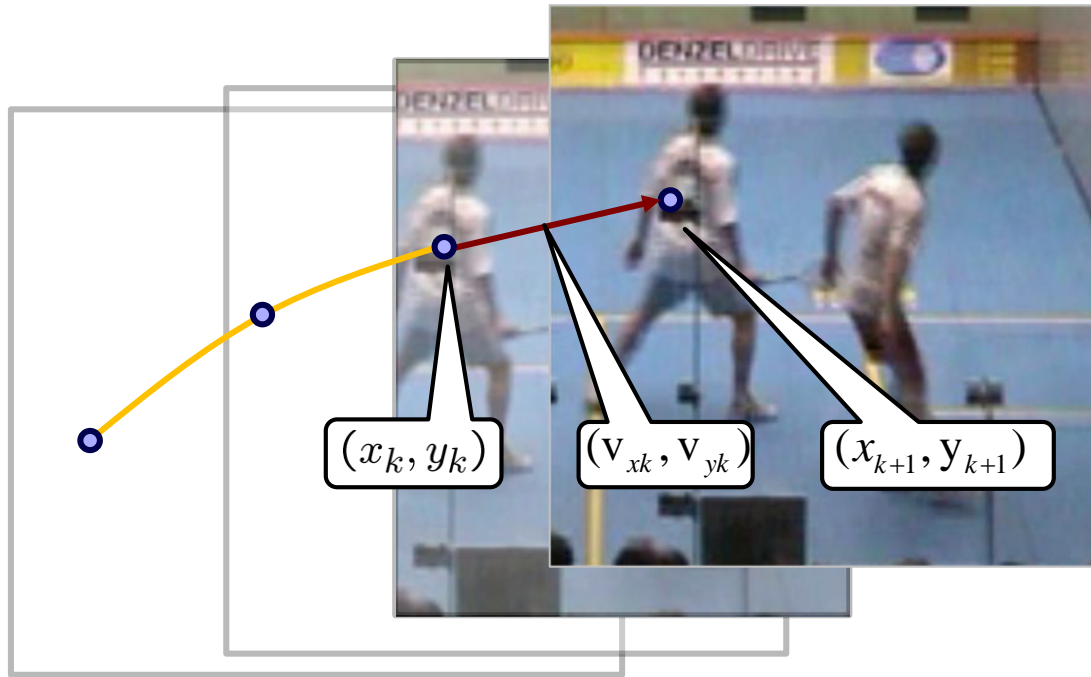


Inference



What is a dynamic model?

- Predicts the target state from its previous estimate.



$$x_{k+1} = x_k + v_{xk} \Delta t$$

$$y_{k+1} = y_k + v_{yk} \Delta t$$

- This is an example of a **constant-velocity model**
- Assumption: velocity at $k+1$ is equal to velocity at k .

A constant velocity model

- 1D problem, but **2D state space** with position and velocity

$$\mathbf{x} = \begin{bmatrix} x \\ \dot{x} \end{bmatrix} \longrightarrow \dot{\mathbf{x}} = F\mathbf{x} \qquad F = ?$$

- Velocity does not change:

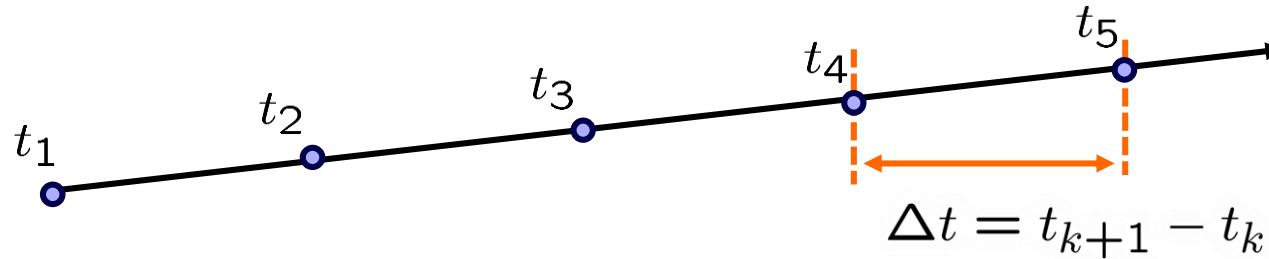
$$\dot{\mathbf{x}} = F\mathbf{x}$$

$$\dot{\mathbf{x}} = \begin{bmatrix} \dot{x} \\ \ddot{x} \end{bmatrix} = \begin{bmatrix} ? \\ ? \end{bmatrix} \begin{bmatrix} x \\ \dot{x} \end{bmatrix}$$

From continuous to discrete

- A continuous motion is sampled at equally spaced time-steps (spacing Δt):

$$\dot{\mathbf{x}} = F\mathbf{x}$$



- Solution according to Stengel (p.84)

$$\mathbf{x}(t_k) = \Phi(\Delta t)\mathbf{x}(t_{k-1})$$

$$\Phi(\Delta t) = e^{F\Delta t} = I + F\Delta t + \frac{1}{2!}F^2\Delta t^2 + \frac{1}{3!}F^3\Delta t^3 + \dots$$

From continuous to discrete

- For the constant-velocity model:

$$\dot{\mathbf{x}} = F\mathbf{x} \quad F = \begin{bmatrix} 0 & 1 \\ 0 & 0 \end{bmatrix}$$

Matlab symbolic toolbox:

```
>> syms T
>> F = [0 1; 0 0]
>> Fi = expm(F*T)
```

Compute using your favorite symbolic toolbox:

$$\Phi(\Delta t) = e^{F\Delta t}$$
$$\Phi(\Delta t) = \begin{bmatrix} 1 & \Delta t \\ 0 & 1 \end{bmatrix}$$

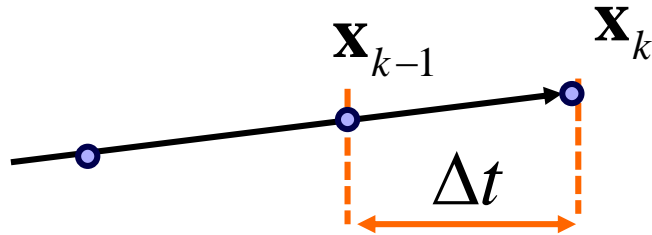
Python symbolic toolbox:

```
>> import sympy as sp
>> from sympy.interactive.printing import init_printing
>> init_printing(use_unicode=False, wrap_line=False)
T = sp.symbols('T')
>> F = sp.Matrix([[0, 1],[0, 0]])
>> Fi = sp.exp(F*T)
```

Discrete constant velocity model

- See if the derived CV model makes any sense:

$$\mathbf{x}_k = \Phi(\Delta t) \mathbf{x}_{k-1} \quad , \quad \Phi(\Delta t) = \begin{bmatrix} 1 & \Delta t \\ 0 & 1 \end{bmatrix} \quad , \quad \mathbf{x}_k = \begin{bmatrix} x_k \\ \dot{x}_k \end{bmatrix}$$



$$\begin{bmatrix} x_k \\ \dot{x}_k \end{bmatrix} = \begin{bmatrix} 1 & \Delta t \\ 0 & 1 \end{bmatrix} \begin{bmatrix} x_{k-1} \\ \dot{x}_{k-1} \end{bmatrix} \longrightarrow \begin{aligned} x_k &= x_{k-1} + \dot{x}_{k-1} \Delta t \\ \dot{x}_k &= \dot{x}_{k-1} \end{aligned}$$

But constant velocity is *not a very realistic assumption...*

A nearly-constant-velocity model

- Assume that **acceleration** is not zero, but **is noisy**:

$$\ddot{x} = w$$

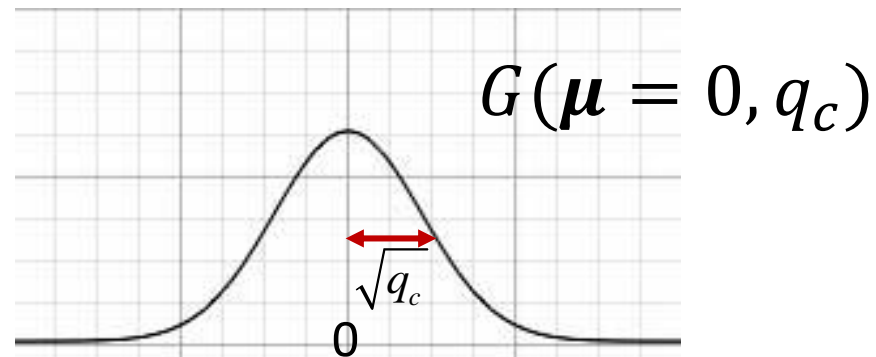
$$\mathbf{x} = \begin{bmatrix} x \\ \dot{x} \end{bmatrix}$$

$$\dot{\mathbf{x}} = \begin{bmatrix} \dot{x} \\ \ddot{x} \end{bmatrix} = \begin{bmatrix} 0 & 1 \\ 0 & 0 \end{bmatrix} \begin{bmatrix} x \\ \dot{x} \end{bmatrix} + \begin{bmatrix} ? \\ ? \end{bmatrix} w$$

$$\dot{\mathbf{x}} = F\mathbf{x} + Lw$$

Deterministic (points to $F\mathbf{x}$)

Stochastic (points to Lw) — a white noise sequence specified by its covariance (spectral density) q_c !



A discrete counterpart

- Solution of $\dot{\mathbf{x}} = F\mathbf{x} + Lw$ according to Stengel (p.84)

$$\mathbf{x}_k = \Phi(\Delta t)\mathbf{x}_{k-1} + W_{k-1}, \quad \Phi(\Delta t) = e^{F\Delta t} \text{ (deterministic)}$$

W_{k-1} is a random variable: $W_{k-1} = \int_{t_{k-1}}^{t_k} \Phi(\tau) L w(\tau) d\tau$

Governed by a pdf and specified by the covariance matrix:

$$Q_{k-1} = \int_0^{\Delta t} (\Phi(\xi) L) q_c (\Phi(\xi) L)^T d\xi$$

Might want to apply Matlab/Python/Mathematica to solve for $Q_{k-1} \dots$

The covariance Q of a NCV

- Recall:

$$\Phi(\Delta t) = \begin{bmatrix} 1 & \Delta t \\ 0 & 1 \end{bmatrix} \quad \dot{\mathbf{x}} = \begin{bmatrix} \dot{x} \\ \ddot{x} \end{bmatrix} = \begin{bmatrix} 0 & 1 \\ 0 & 0 \end{bmatrix} \begin{bmatrix} x \\ \dot{x} \end{bmatrix} + \begin{bmatrix} 0 \\ 1 \end{bmatrix} w$$

$$Q_{k-1} = \int_0^{\Delta t} (\Phi(\xi) L) q_c (\Phi(\xi) L)^T d\xi \quad Q = q \begin{bmatrix} \frac{1}{3} \Delta t^3 & \frac{1}{2} \Delta t^2 \\ \frac{1}{2} \Delta t^2 & \Delta t \end{bmatrix}$$

Matlab symbolic toolbox:

```
>>syms T q
>> Fi = [1 T;0 1]
>> L=[0 ;1 ]
>> Q=int((Fi*L)*q*(Fi*L)',T,0,T)
```

Python symbolic toolbox:

```
>> import sympy as sp
```

```
>> T, q = sp.symbols('T q')
>> Fi = sp.Matrix([[1, T],[0, 1]])
>> L = sp.Matrix([[0], [1]])
>> Q = sp.integrate( (Fi*L)*q*(Fi*L).T, (T, 0, T) )
```

The nearly-constant-velocity model

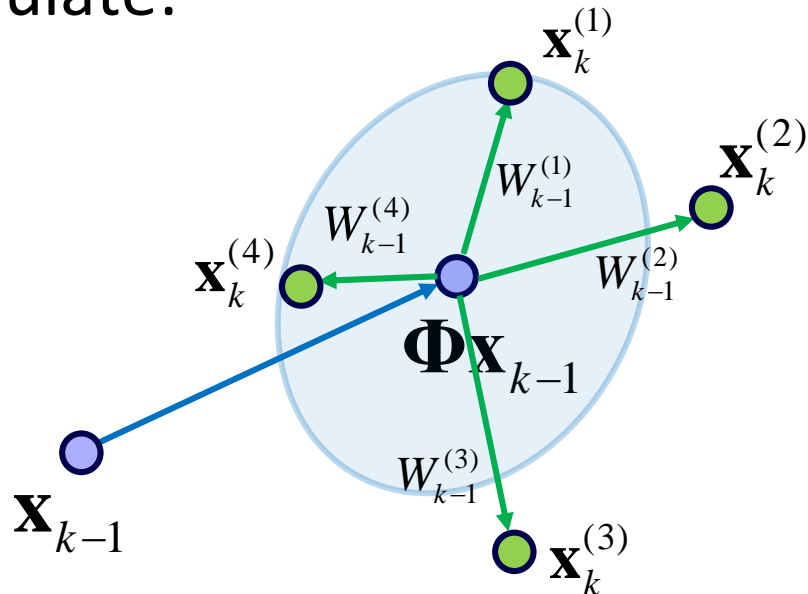
- We are done: $\mathbf{x}_k = \Phi \mathbf{x}_{k-1} + W_{k-1}$

$$W_{k-1} \sim G(\mu = 0, Q)$$

$$\Phi = \begin{bmatrix} 1 & \Delta t \\ 0 & 1 \end{bmatrix}$$

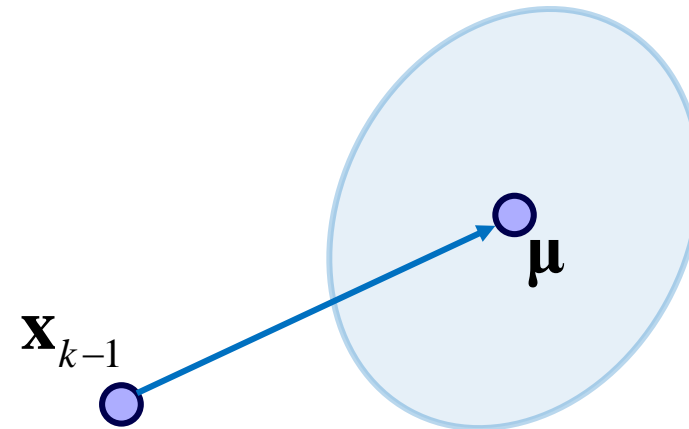
$$Q = q \begin{bmatrix} \frac{1}{3} \Delta t^3 & \frac{1}{2} \Delta t^2 \\ \frac{1}{2} \Delta t^2 & \Delta t \end{bmatrix}$$

Let's simulate:



Probabilistic model:

$$p(\mathbf{x}_k | \mathbf{x}_{k-1}) = G(\mu = \Phi \mathbf{x}_{k-1}, Q)$$



It is easy to extend to 2D or higher

- A 2D NCV example:
 - If you like compact derivation...

$$\mathbf{x} = \begin{bmatrix} x \\ \dot{x} \\ y \\ \dot{y} \end{bmatrix} \quad \dot{\mathbf{x}} = \begin{bmatrix} 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 \end{bmatrix} \begin{bmatrix} x \\ \dot{x} \\ y \\ \dot{y} \end{bmatrix} + \begin{bmatrix} 0 & 0 \\ 1 & 0 \\ 0 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} w_1 \\ w_2 \end{bmatrix}$$

Will be part of the assignment...

- Or, more common and simpler (but equivalent):
 - Two separate instances of the NCV model that we derived in previous slides (one for x and one for y).

Random walk dynamic model

- **Brownian motion** – velocity is noise!

$$\mathbf{x} = \begin{bmatrix} x \\ \dot{x} \end{bmatrix} \quad \dot{x} = w$$

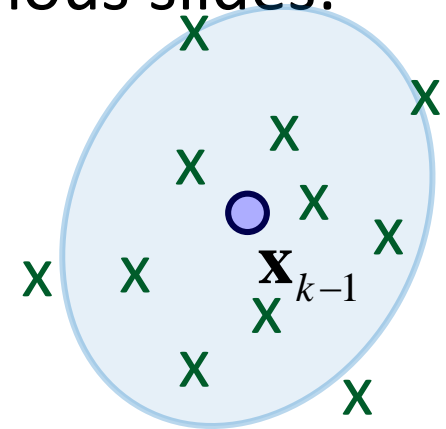
Step 1: Write the equation in a form of
 $\dot{x} = Fx + Lw$

$$\dot{\mathbf{x}} = \begin{bmatrix} \dot{x} \\ \ddot{x} \end{bmatrix} = \begin{bmatrix} 0 & 0 \\ 0 & 0 \end{bmatrix} \begin{bmatrix} x \\ \dot{x} \end{bmatrix} + \begin{bmatrix} 1 \\ 0 \end{bmatrix} w$$

- **Apply the same** (Matlab/Python) **derivation** as in previous slides:

$$\mathbf{x}_k = \mathbf{x}_{k-1} + W_{k-1}$$

$$p(\mathbf{x}_k | \mathbf{x}_{k-1}) = G(\boldsymbol{\mu} = \mathbf{x}_{k-1}, Q)$$

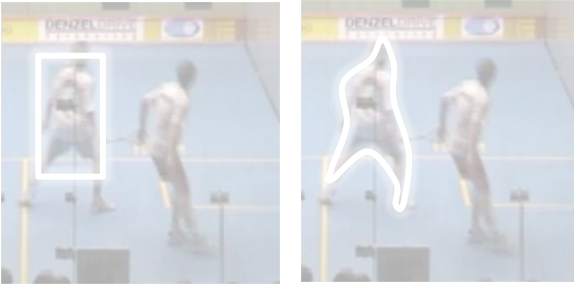


Widely used dynamic models

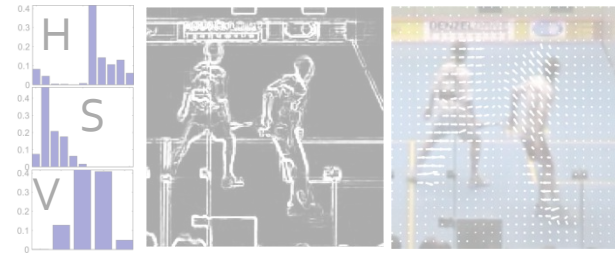
- **Velocity** is non-correlated:
 - **Velocity** modelled by a **white noise** sequence
 - Random Walk model (**RW**), Brownian motion
- **Acceleration** non-correlated:
 - **Acceleration** modelled by a **white noise** sequence
 - Nearly constant velocity (**NCV**)
- Derivative of acceleration (**jerk**) non-correlated:
 - **Jerk** modelled by a **white noise** sequence
 - Nearly constant acceleration (**NCA**)

Key Ingredients

State definition



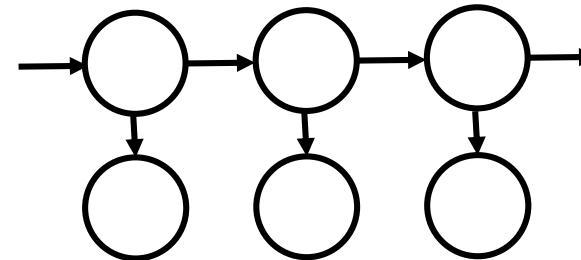
Observation model



Dynamic model



Inference



Probabilistic view

- Given a sequence of observations $\mathbf{y}_{1:k} = \{\mathbf{y}_i\}_{i=1:k}$
(think about the observation in most abstract way – an image)

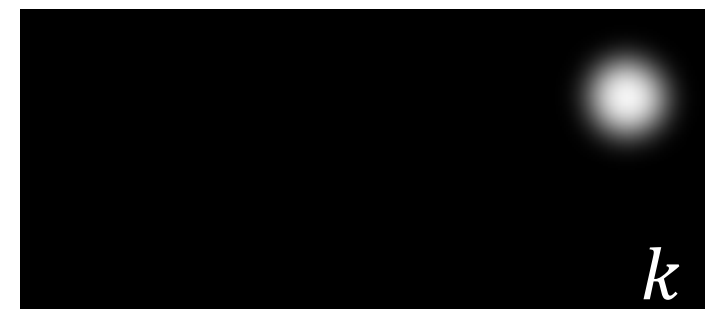


- ...want to find the density over the current state \mathbf{x}_k

$$p(\mathbf{x}_k | \mathbf{y}_{1:k})$$

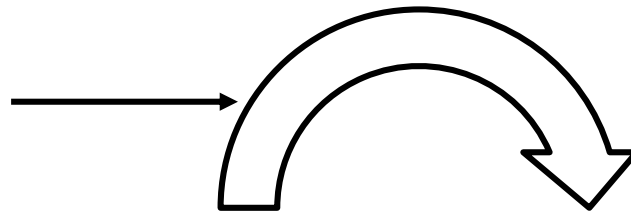
Using the Bayesian terminology:

The **posterior** over the \mathbf{x}_k

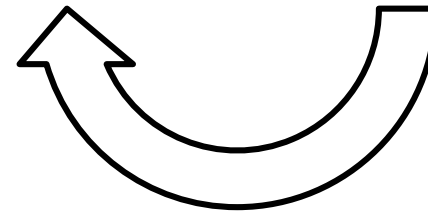


Towards Recursive Bayes Filter

- The goal is to **rewrite the posterior** in the current time-step k as a **function of the posterior from the previous time-step $k-1$** :



$$p(\mathbf{x}_k \mid y_{1:k}) = f(p(\mathbf{x}_{k-1} \mid y_{1:k-1}), y_k)$$



Towards Recursive Bayes Filter

$$p(\mathbf{x}_k | y_{1:k}) = \frac{p(y_k | \mathbf{x}_k, y_{1:k-1}) p(\mathbf{x}_k | y_{1:k-1})}{p(y_k | y_{1:k-1})}$$

Assumption 1: *Current measurement is conditionally independent from all previous measurements given x_k .*

$$p(y_k | \mathbf{x}_k, y_{1:k-1}) \equiv p(y_k | \mathbf{x}_k)$$



$$p(\mathbf{x}_k | y_{1:k}) = \frac{p(y_k | \mathbf{x}_k) p(\mathbf{x}_k | y_{1:k-1})}{p(y_k | y_{1:k-1})}$$

Towards Recursive Bayes Filter

- Expand the density $p(\mathbf{x}_k | y_{1:k-1})$:

$$p(\mathbf{x}_k | y_{1:k-1}) = \int p(\mathbf{x}_k | \mathbf{x}_{k-1}, y_{1:k-1}) p(\mathbf{x}_{k-1} | y_{1:k-1}) d\mathbf{x}_{k-1}$$

- Assumption 2: *Current state is conditionally independent from all previous measurements given x_{k-1} .*

$$p(\mathbf{x}_k | \mathbf{x}_{k-1}, y_{1:k-1}) \equiv p(\mathbf{x}_k | \mathbf{x}_{k-1})$$



$$p(\mathbf{x}_k | y_{1:k-1}) = \int p(\mathbf{x}_k | \mathbf{x}_{k-1}) p(\mathbf{x}_{k-1} | y_{1:k-1}) d\mathbf{x}_{k-1}$$

The Bayes Recursive Filter

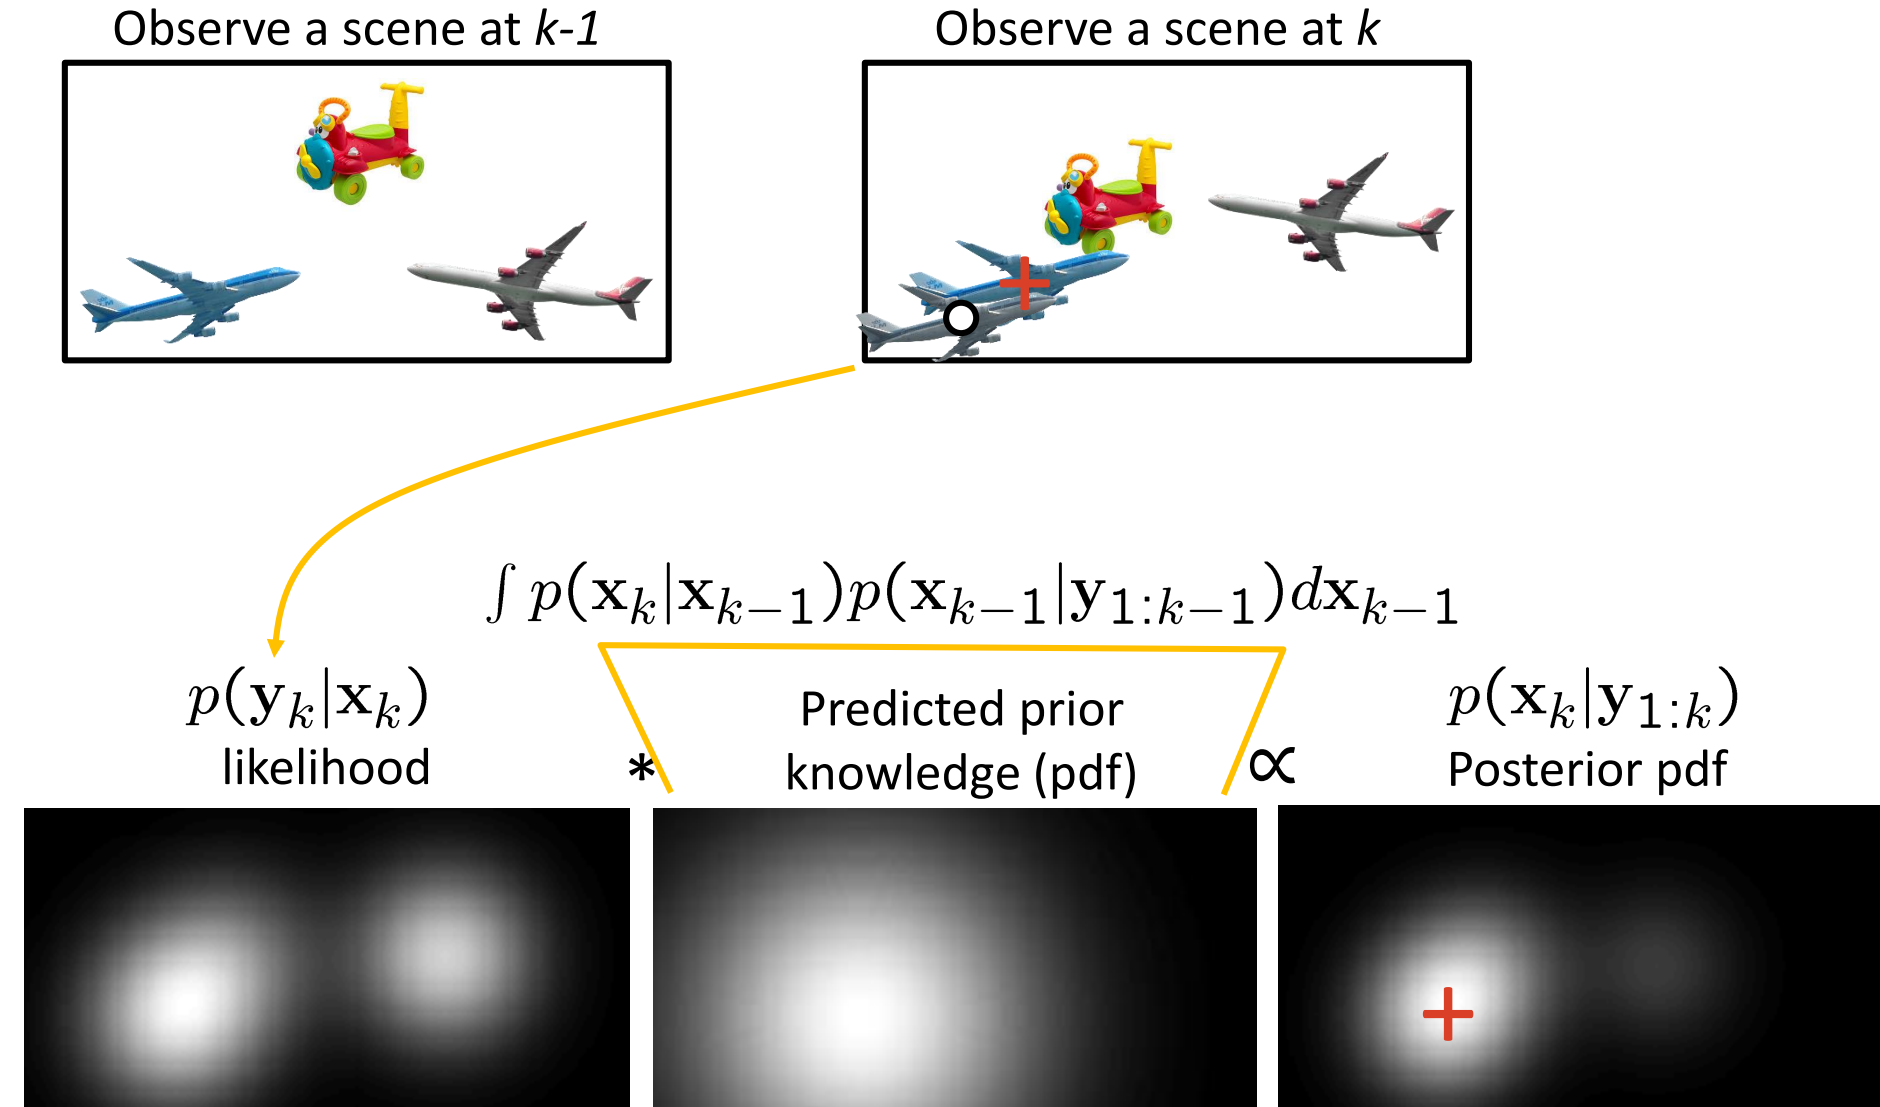
- Putting it all together:

$$p(\mathbf{x}_k | y_{1:k}) = \frac{p(y_k | \mathbf{x}_k) \int p(\mathbf{x}_k | \mathbf{x}_{k-1}) p(\mathbf{x}_{k-1} | y_{1:k-1}) d\mathbf{x}_{k-1}}{p(y_k | y_{1:k-1})}$$

- The denominator does not depend on the state:

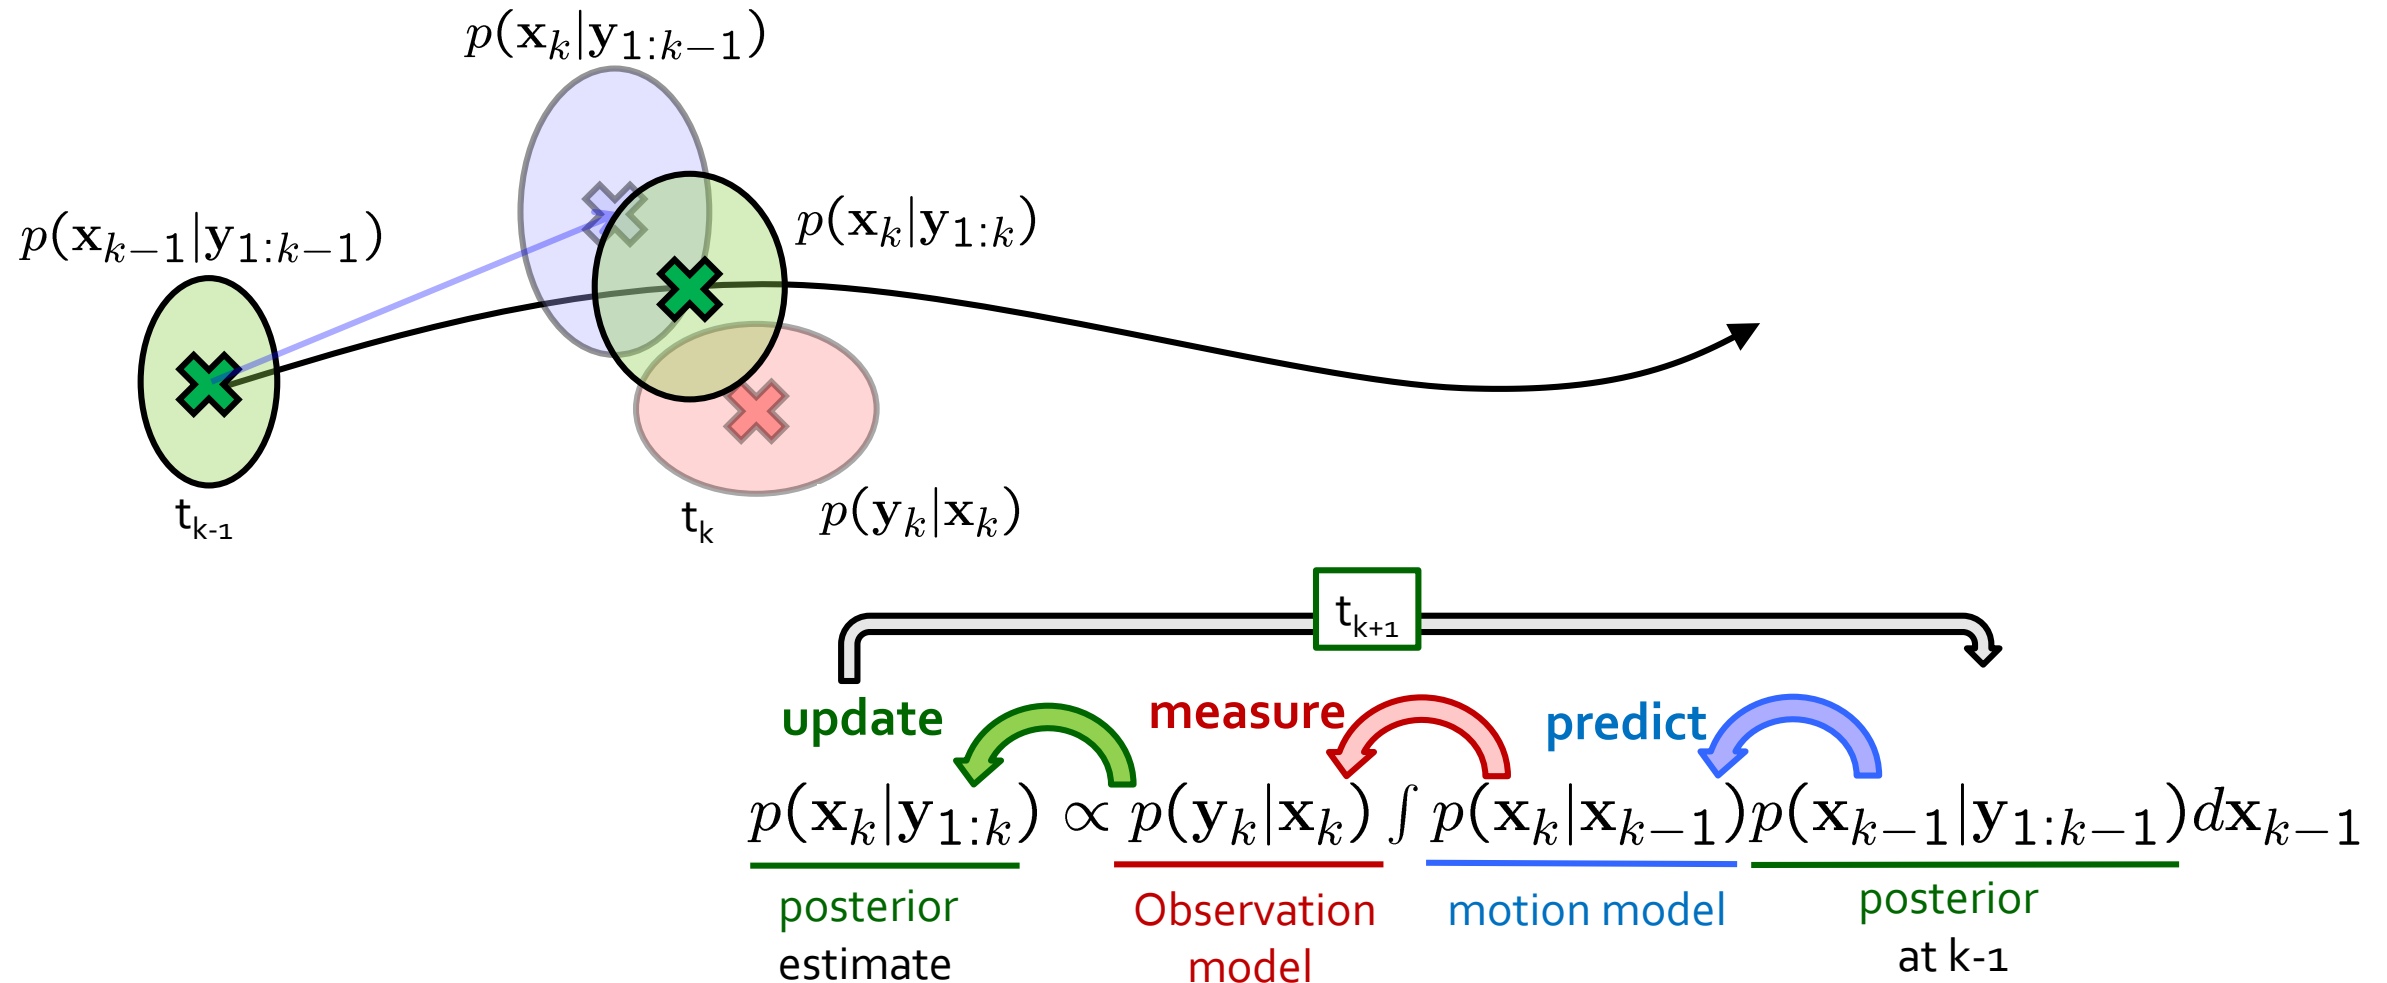
$$\underbrace{p(\mathbf{x}_k | y_{1:k})}_{\text{posterior at } k} \propto \underbrace{p(y_k | \mathbf{x}_k)}_{\text{likelihood (observation model)}} \int \underbrace{p(\mathbf{x}_k | \mathbf{x}_{k-1})}_{\text{motion model (dynamic model)}} \underbrace{p(\mathbf{x}_{k-1} | y_{1:k-1})}_{\text{posterior at } k-1} d\mathbf{x}_{k-1}$$

Recall the airplane tracking example



Recursive Bayesian Filter

- At each time-step estimate the posterior:



Acknowledgement

- Some images and parts of slides have been taken from the following talks:
 - Kevin Smith's "SELECTED TOPICS IN COMPUTER VISION – 2D tracking"