



Advanced computer vision methods

Tracking by Recursive Bayes Filters

Part II: Kalman filter

Matej Kristan

Laboratorij za Umetne Vizualne Spoznavne Sisteme,
Fakulteta za računalništvo in informatiko,
Univerza v Ljubljani

Previously at ACVM...

- Tracking as state estimation

$$\mathbf{x}_k = \begin{bmatrix} x_k \\ y_k \\ \dot{x}_k \\ \dot{y}_k \end{bmatrix}$$

A noisy detector

0.4 1.0
1.0

Observe a scene at $k-1$



Observe a scene at k



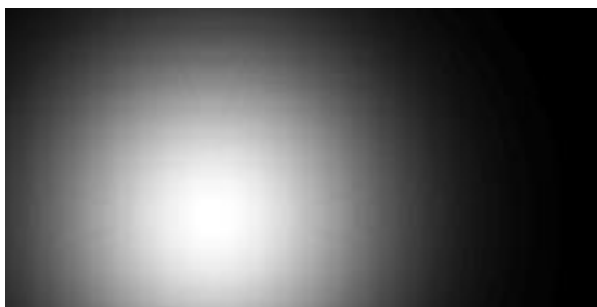
- Encode state info by pdfs

likelihood




*

Predicted prior knowledge (pdf)



\propto

Posterior pdf




Maximum
a posteriori state!

Previously at ACVM ...

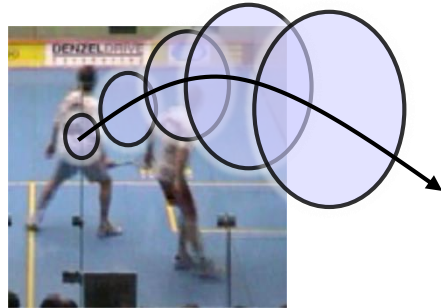
State definition




Observation model



Dynamic model

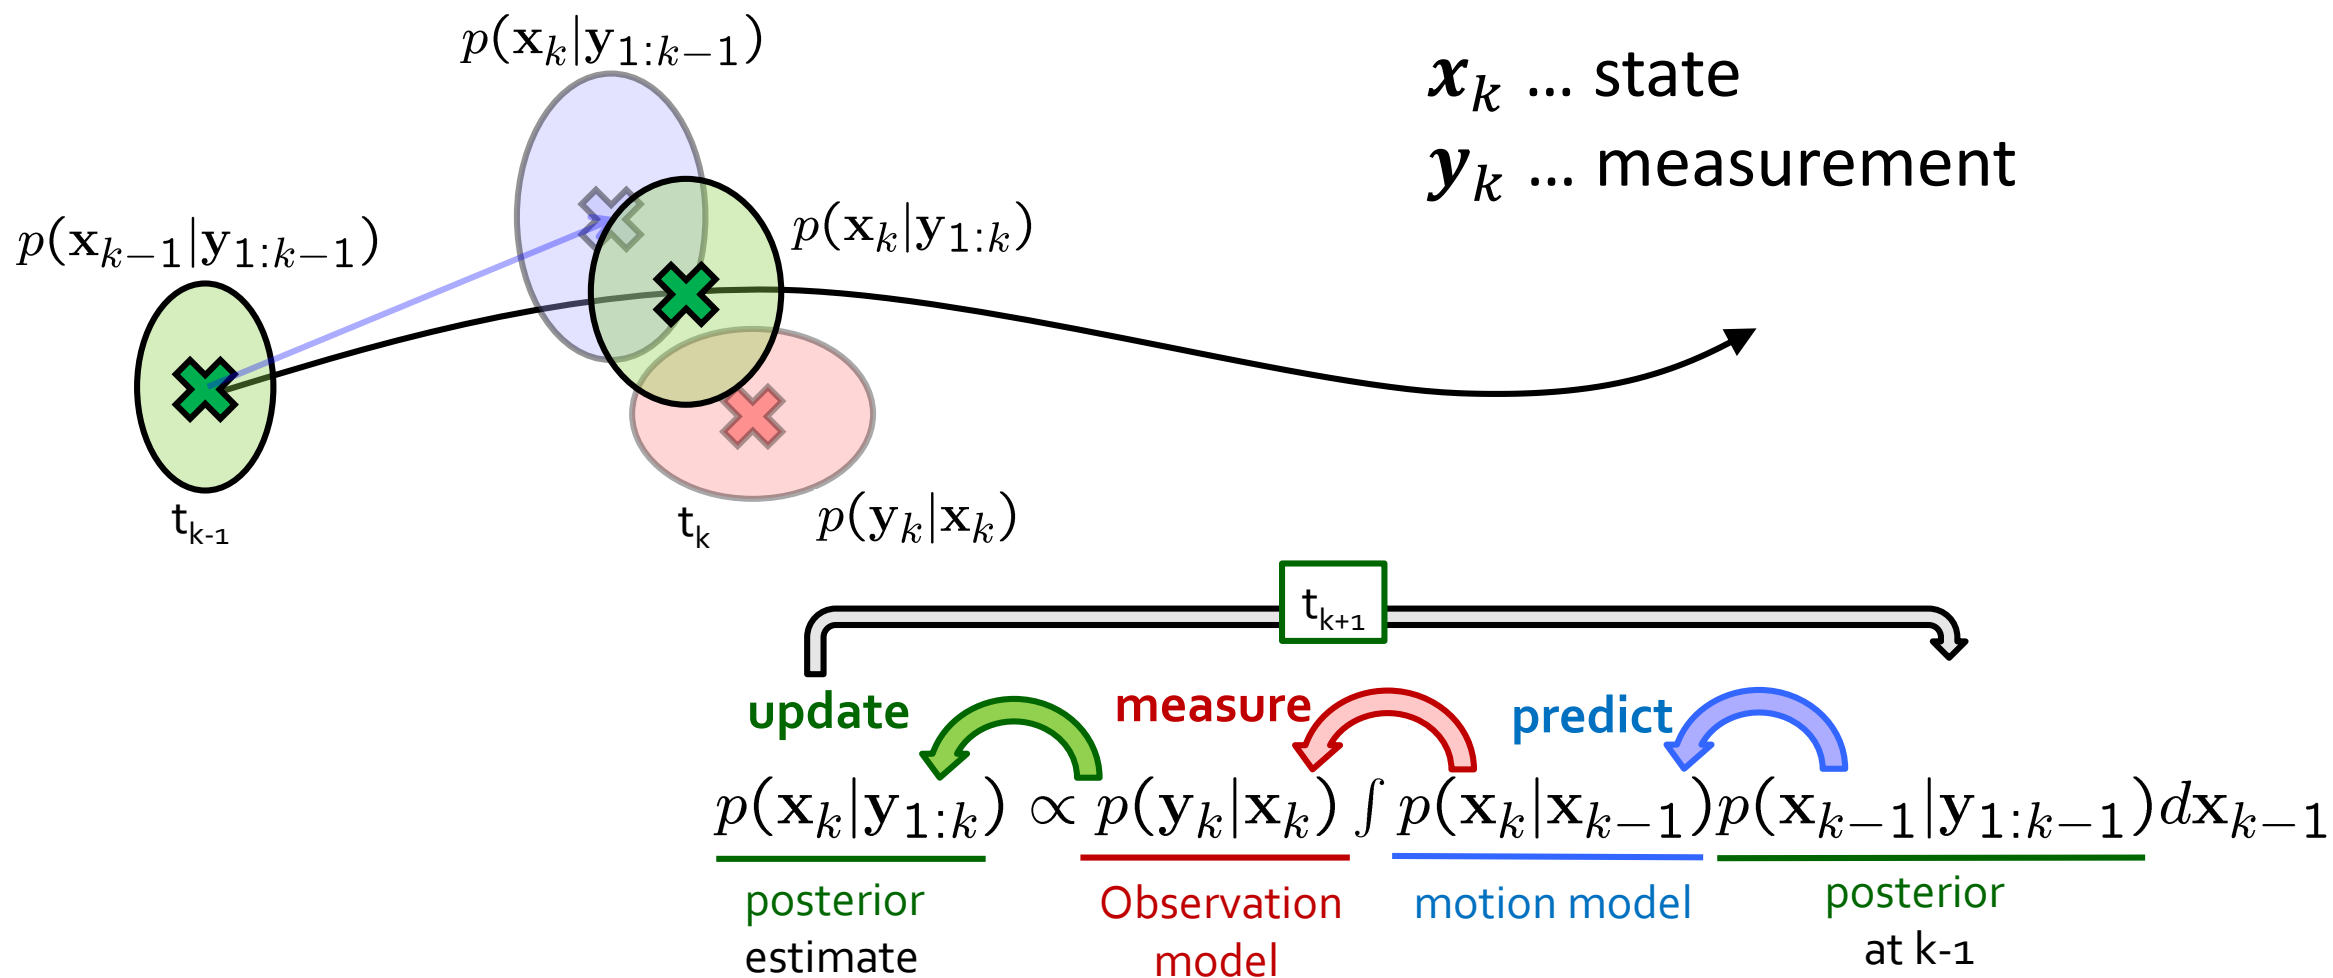


Inference



Previously at ACVM ...

- At each time-step estimate the posterior:



How to model the posterior?

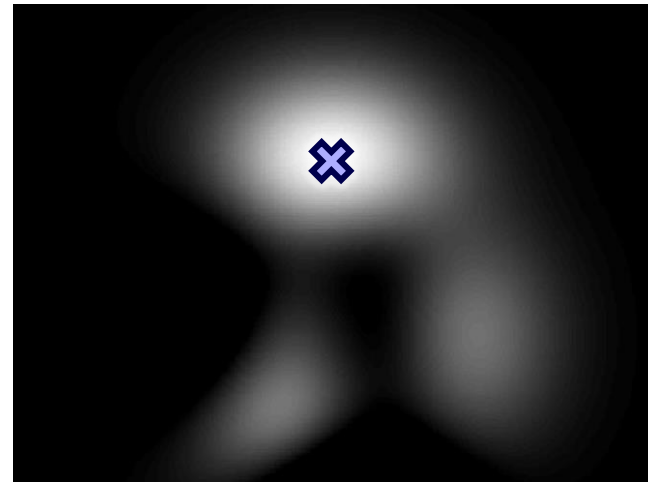
- Face localization

$$p(\mathbf{x}_k | \mathbf{y}_{1:k}) \propto p(\mathbf{y}_k | \mathbf{x}_k) \int p(\mathbf{x}_k | \mathbf{x}_{k-1}) p(\mathbf{x}_{k-1} | \mathbf{y}_{1:k-1}) d\mathbf{x}_{k-1}$$

current input image



current posterior



$$p(\mathbf{x}_k | \mathbf{y}_{1:k})$$

- Indicates **likely and less likely** positions of a tracked face.
- For example, get **most probable position** (maximum a posteriori)

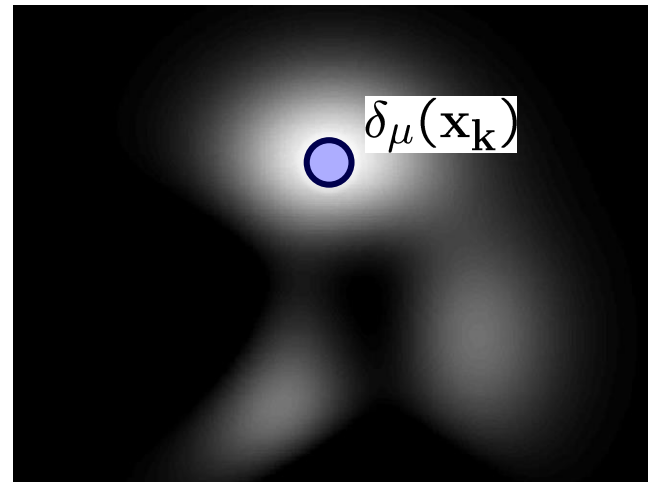
Analytic representations

- A single point : Dirac-delta

current input image



current posterior



$$p(\mathbf{x}_k | \mathbf{y}_{1:k})$$

$$p(\mathbf{x}_k | \mathbf{y}_{1:k}) = \begin{cases} \text{inf} & \text{if } \mathbf{x}_k = \mu \\ 0 & \text{otherwise} \end{cases}$$

$$\int_{-\infty}^{\infty} p(\mathbf{x}_k | \mathbf{y}_{1:k}) d\mathbf{x}_k = 1$$

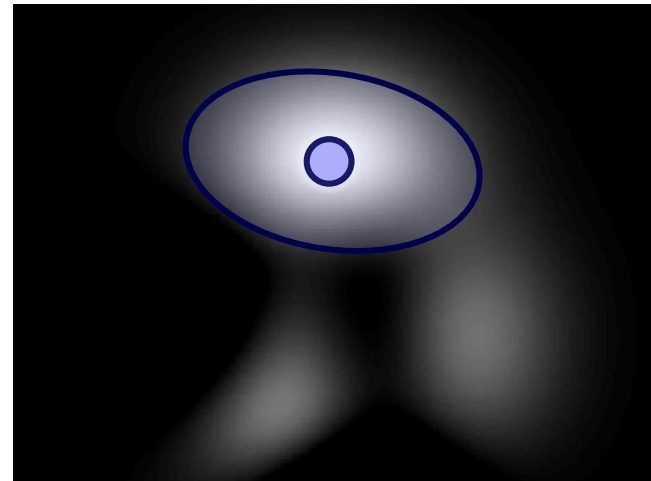
Analytic representations

- A single point + covariance: a Gaussian distribution

current input image



current posterior



$$p(\mathbf{x}_k | \mathbf{y}_{1:k})$$

$$p(\mathbf{x}_k | \mathbf{y}_{1:k}) = \mathcal{N}(\mathbf{x} | \mu; \Sigma)$$

$$\mathcal{N}(\mathbf{x} | \mu; \Sigma) = \frac{1}{\sqrt{(2\pi)^d |\Sigma|}} e^{-\frac{1}{2}(\mathbf{x}_k - \mu)^T \Sigma^{-1} (\mathbf{x}_k - \mu)}$$

What if everything was Gaussian?

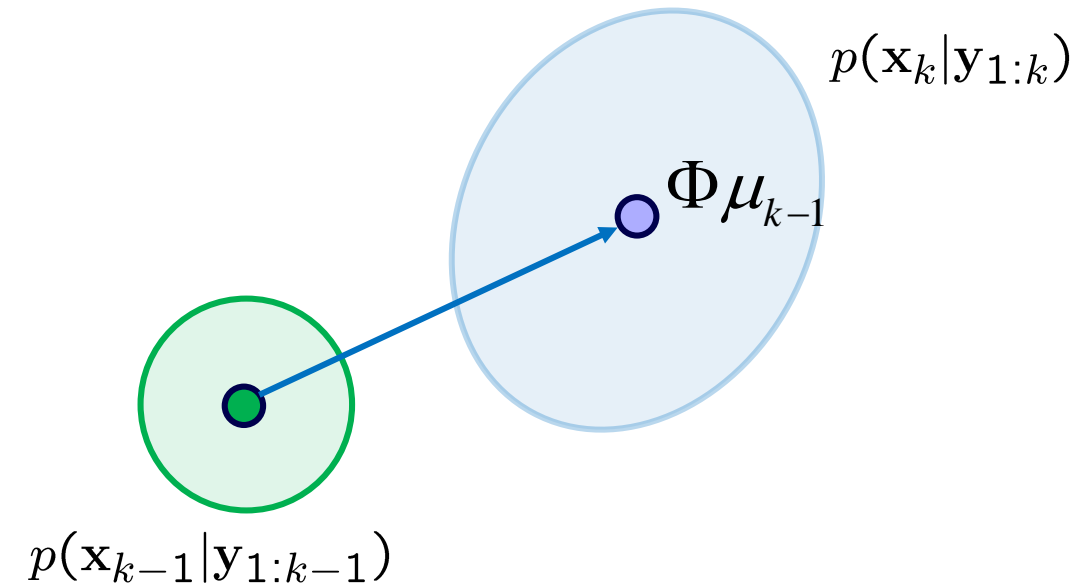
- Assume that:

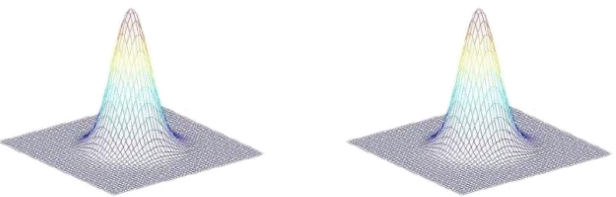
- The **posterior** is a single **Gaussian**.

$$p(\mathbf{x}_k | \mathbf{y}_{1:k}) = \mathcal{N}(\mathbf{x}_k | \mu_k; \mathbf{P}_k)$$

- Dynamic model** is linear with **Gaussian** noise.

$$p(\mathbf{x}_k | \mathbf{x}_{k-1}) = \mathcal{N}(\mathbf{x}_k | \Phi \mathbf{x}_{k-1}; \mathbf{Q}_k)$$



$$\underbrace{p(\mathbf{x}_k | \mathbf{y}_{1:k})}_{\text{posterior estimate}} \propto \underbrace{p(\mathbf{y}_k | \mathbf{x}_k)}_{\text{Observation model}} \int \underbrace{p(\mathbf{x}_k | \mathbf{x}_{k-1})}_{\text{motion model}} \underbrace{p(\mathbf{x}_{k-1} | \mathbf{y}_{1:k-1})}_{\text{posterior at k-1}} d\mathbf{x}_{k-1}$$


What if everything was Gaussian?

- Assume that:

- The **posterior** is a single **Gaussian**.

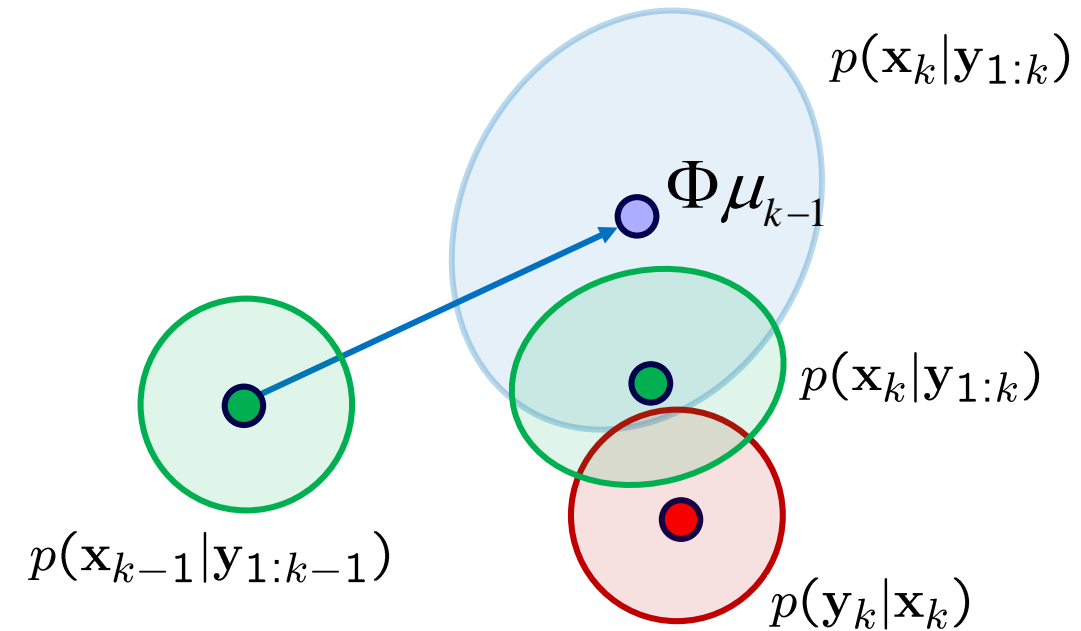
$$p(\mathbf{x}_k | \mathbf{y}_{1:k}) = \mathcal{N}(\mathbf{x}_k | \mu_k; \mathbf{P}_k)$$

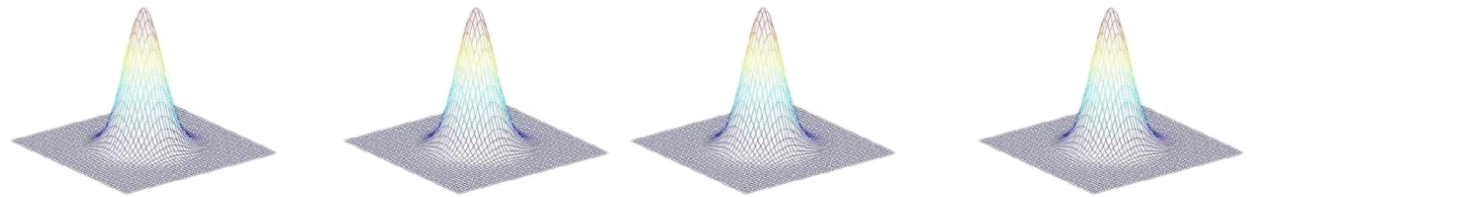
- Dynamic model** is linear with **Gaussian** noise.

$$p(\mathbf{x}_k | \mathbf{x}_{k-1}) = \mathcal{N}(\mathbf{x}_k | \Phi \mathbf{x}_{k-1}; \mathbf{Q}_k)$$

- Observation** model is a **Gaussian**.

$$p(\mathbf{y}_k | \mathbf{x}_k) = \mathcal{N}(\mathbf{y}_k | \mathbf{H} \mathbf{x}_k; \mathbf{R}_k)$$




$$\underbrace{p(\mathbf{x}_k | \mathbf{y}_{1:k})}_{\text{posterior estimate}} \propto \underbrace{p(\mathbf{y}_k | \mathbf{x}_k)}_{\text{Observation model}} \int \underbrace{p(\mathbf{x}_k | \mathbf{x}_{k-1})}_{\text{motion model}} \underbrace{p(\mathbf{x}_{k-1} | \mathbf{y}_{1:k-1})}_{\text{posterior at k-1}} d\mathbf{x}_{k-1}$$

The Kalman filter

- Assume that **all distributions are Gaussians**
- And assume **linear dynamics**
- A **well-known filter** emerges

The Kalman filter!*

- **Originally** presented as a **recursive Least Squares** method, not as a **Recursive Bayes Filter**

Kalman, R. E. 1930-2016. A New Approach to Linear Filtering and Prediction Problems, II Transaction of the ASME—Journal of Basic Engineering, pp. 35-45 (March 1960).



* To be exact: *Stratonovich–Kalman–Bucy filter* (Soviet mathematician Stratonovich proposed a more general case earlier)

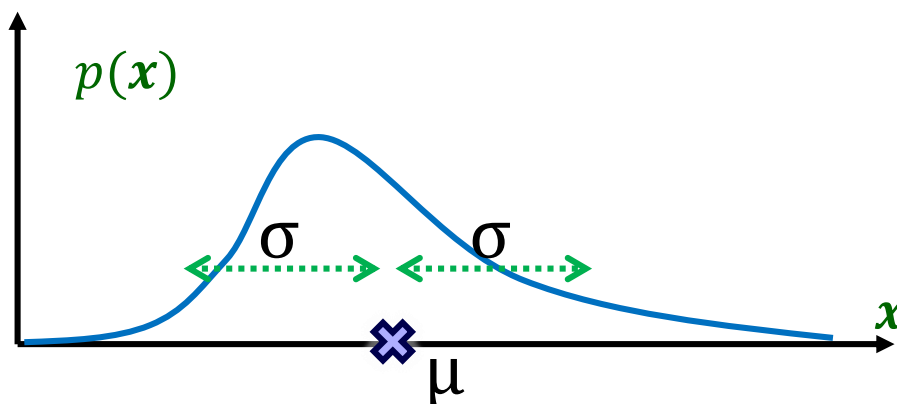
Recall some basic statistics

- Recall expected values
 - Expected value (weighted average of x)

$$\mu = \langle x \rangle_{p(x)} = \int_{-\infty}^{\infty} xp(x) dx \quad \text{For short: } \langle x \rangle = \langle x \rangle_{p(x)}$$

- Variance = expected squared change (*i.e.*, weighted average of $(x - \mu)^2$)

$$\sigma^2 = \langle (x - \mu)^2 \rangle_{p(x)} = \int_{-\infty}^{\infty} (x - \mu)^2 p(x) dx$$



Recall some basic statistics

- Same for vectors...
 - Expected value (weighted average of \mathbf{x})

$$\boldsymbol{\mu} = \langle \mathbf{x} \rangle = \int_{-\infty}^{\infty} \mathbf{x} p(\mathbf{x}) d\mathbf{x}$$

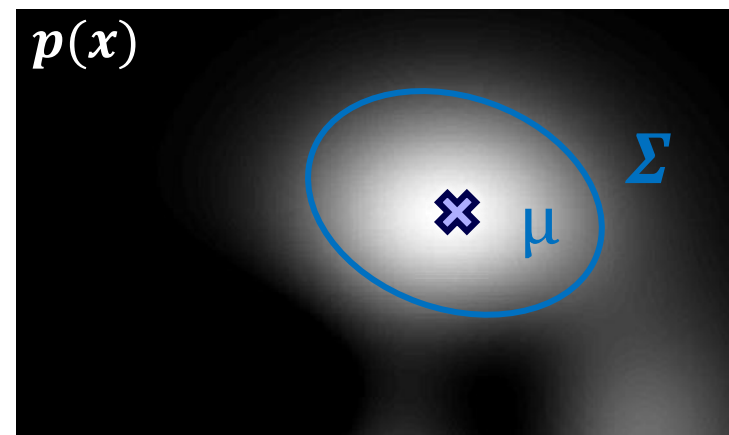
- Variance = expected sq. change,

$$\Delta \mathbf{x} = \mathbf{x} - \boldsymbol{\mu}$$

$$\boldsymbol{\Sigma} = \langle \Delta \mathbf{x} \Delta \mathbf{x}^T \rangle$$

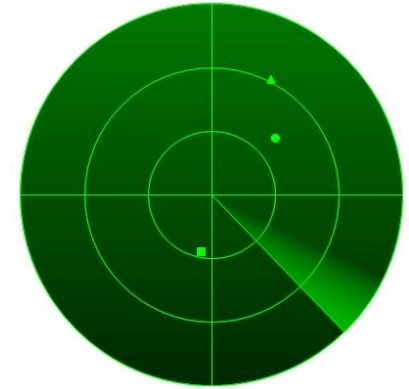
$$= \langle (\mathbf{x} - \boldsymbol{\mu})(\mathbf{x} - \boldsymbol{\mu})^T \rangle$$

$$= \int_{-\infty}^{\infty} (\mathbf{x} - \boldsymbol{\mu})(\mathbf{x} - \boldsymbol{\mu})^T p(\mathbf{x}) d\mathbf{x}$$



A working example

- Track an airplane
- State: position and velocity
- Observe: only position
- Dynamics: Assume a NCV model



$$\mathbf{x}_k = \begin{bmatrix} x_k \\ y_k \\ \dot{x}_k \\ \dot{y}_k \end{bmatrix}$$

$$\mathbf{y}_k = \begin{bmatrix} x_k^{(\text{measured})} \\ y_k^{(\text{measured})} \end{bmatrix}$$

Slightly more formally

- State and observation

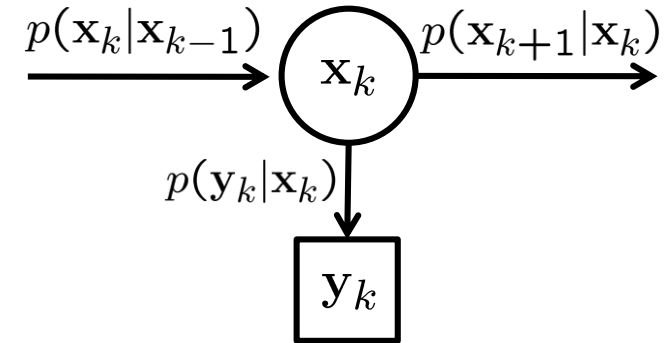
$$\mathbf{x}_k = \begin{bmatrix} x_k \\ y_k \\ \dot{x}_k \\ \dot{y}_k \end{bmatrix} \quad \mathbf{y}_k = \begin{bmatrix} x_k^{(m)} \\ y_k^{(m)} \end{bmatrix}$$

- Dynamic model

$$\mathbf{x}_k = \Phi \mathbf{x}_{k-1} + \underbrace{\mathbf{w}_k}_{\text{noise } \mathbf{Q}_k}$$

- Observation model

$$y_k = \mathbf{H} \mathbf{x}_k + \underbrace{\mathbf{v}_k}_{\text{noise } \mathbf{R}_k}$$



$$\begin{bmatrix} x_k \\ y_k \\ \dot{x}_k \\ \dot{y}_k \end{bmatrix} = \begin{bmatrix} 1 & 0 & \Delta t & 0 \\ 0 & 1 & 0 & \Delta t \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_{k-1} \\ y_{k-1} \\ \dot{x}_{k-1} \\ \dot{y}_{k-1} \end{bmatrix} + \mathbf{w}_k$$

$$\begin{bmatrix} x_k^{(m)} \\ y_k^{(m)} \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \end{bmatrix} \begin{bmatrix} x_k \\ y_k \\ \dot{x}_k \\ \dot{y}_k \end{bmatrix} + \mathbf{v}_k$$

The recursive Bayes filter

- Recall the recursive equation

$$\underbrace{p(\mathbf{x}_k | \mathbf{y}_{1:k})}_{\text{posterior estimate}} \propto \underbrace{p(\mathbf{y}_k | \mathbf{x}_k)}_{\text{Observation model}} \underbrace{\int p(\mathbf{x}_k | \mathbf{x}_{k-1})}_{\text{motion model}} \underbrace{p(\mathbf{x}_{k-1} | \mathbf{y}_{1:k-1})}_{\text{posterior at } k-1} d\mathbf{x}_{k-1}$$

- The next few slides:

1. Solve the integral:

$$p(\mathbf{x}_k | \mathbf{y}_{1:k-1}) = \int p(\mathbf{x}_k | \mathbf{x}_{k-1}) p(\mathbf{x}_{k-1} | \mathbf{y}_{1:k-1}) d\mathbf{x}_{k-1}$$

2. Solve the posterior update:

$$p(\mathbf{x}_k | \mathbf{y}_{1:k}) \propto p(\mathbf{y}_k | \mathbf{x}_k) p(\mathbf{x}_k | \mathbf{y}_{1:k-1})$$

3. We will use a few tricks that apply to Gaussians

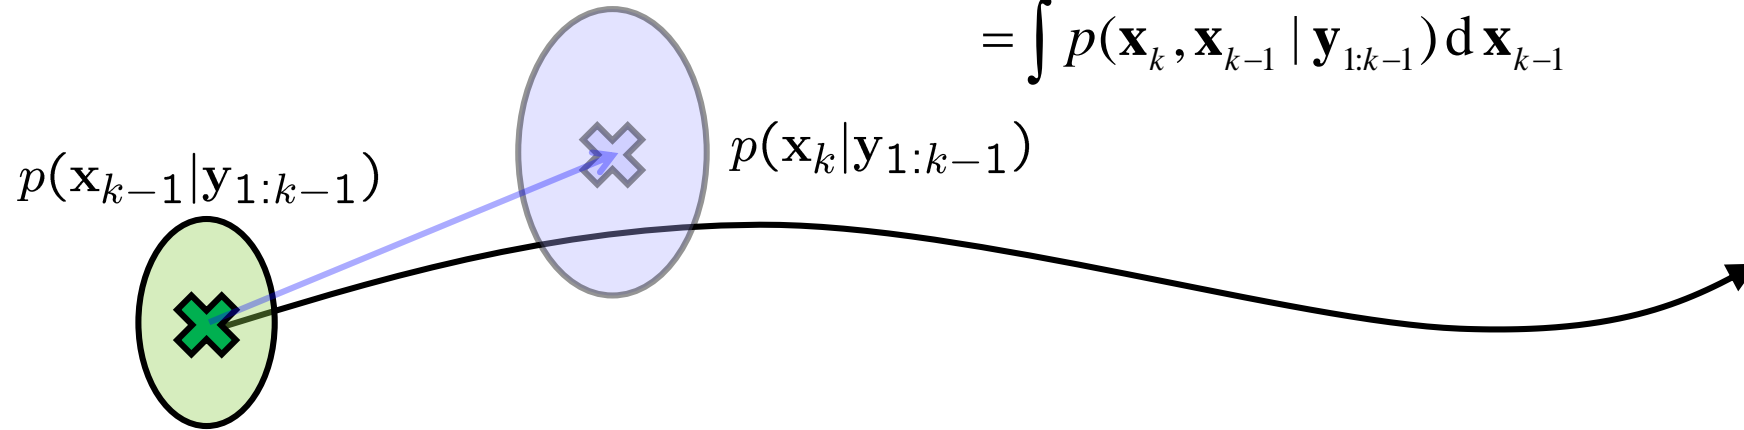
You may want to [review the properties of Gaussian](#), integrals, marginal, etc., in Barber's "Bayesian reasoning and machine learning", Section 8.4.

Solving the prediction

- Solve the integral:

$$p(\mathbf{x}_k | \mathbf{y}_{1:k-1}) = \int p(\mathbf{x}_k | \mathbf{x}_{k-1}) p(\mathbf{x}_{k-1} | \mathbf{y}_{1:k-1}) d\mathbf{x}_{k-1}$$

$$= \int p(\mathbf{x}_k, \mathbf{x}_{k-1} | \mathbf{y}_{1:k-1}) d\mathbf{x}_{k-1}$$

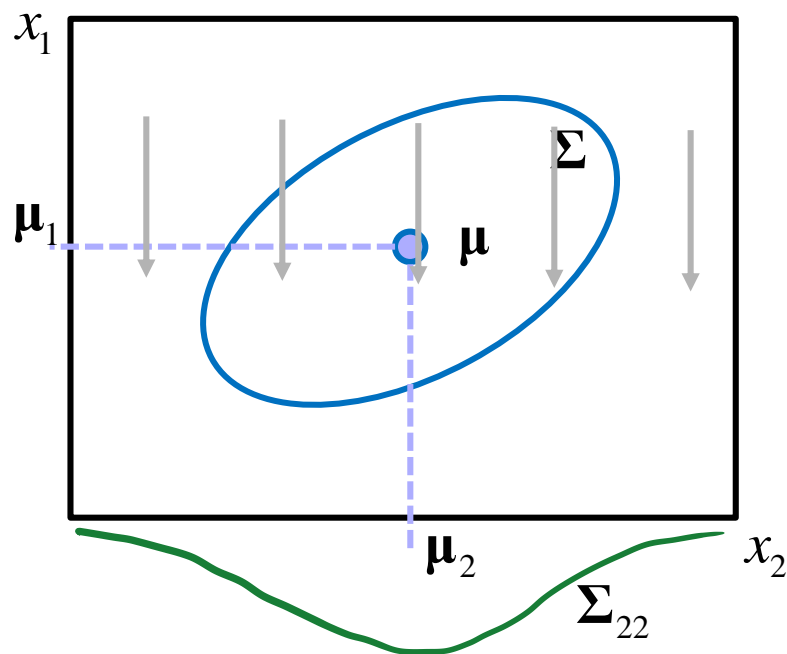


- Since $p(\mathbf{x}_k | \mathbf{x}_{k-1})$ and $p(\mathbf{x}_{k-1} | \mathbf{y}_{1:k-1})$ are both Gaussians, their **product will be a Gaussian** as well.
- Note that we want to **solve the following form**:

$$p(\mathbf{x}_2) = \int p(\mathbf{x}_2 | \mathbf{x}_1) p(\mathbf{x}_1) d\mathbf{x}_1$$

A note on marginalization

- Marginalization:
$$\int p(\mathbf{x}_2 | \mathbf{x}_1) p(\mathbf{x}_1) d\mathbf{x}_1 = \int p(\mathbf{x}_2, \mathbf{x}_1) d\mathbf{x}_1 = p(\mathbf{x}_2)$$



$$p(\mathbf{x}_1, \mathbf{x}_2) = N(\boldsymbol{\mu}, \boldsymbol{\Sigma})$$

$$\boldsymbol{\mu} = \begin{bmatrix} \mu_1 \\ \mu_2 \end{bmatrix}, \boldsymbol{\Sigma} = \begin{bmatrix} \Sigma_{11} & \Sigma_{12} \\ \Sigma_{21} & \Sigma_{22} \end{bmatrix}, \Sigma_{12} = \Sigma_{21}^T$$

$$p(x_2) = \int p(x_1, x_2) dx_1 = N(x_2; \mu_2, \Sigma_{22})$$

- One could solve this integral by “**completing the squares**”¹, but we can take **a shortcut**.

¹See Barber’s “Bayesian reasoning and machine learning”, Section 8.4.1.

Solving the prediction

$$p(\mathbf{x}_2) = \int p(\mathbf{x}_2 | \mathbf{x}_1) p(\mathbf{x}_1) d\mathbf{x}_1$$

- The prior is a Gaussian $p(\mathbf{x}_1) = N(\mathbf{x}_1; \mu_1, \Sigma_1)$
- The **dynamic model** takes the prior and “pushes” it through a **linear model** and adds noise:

$$\mathbf{x}_2 = \Phi \mathbf{x}_1 + \mathbf{W}, \mathbf{W} \sim N(\boldsymbol{\mu} = 0, \mathbf{Q}) \longrightarrow \mathbf{x}_2 \sim N(\mathbf{x}_2; \boldsymbol{\mu}_2, \boldsymbol{\Sigma}_2)$$

$$\boldsymbol{\mu}_2 = \langle \mathbf{x}_2 \rangle = \langle \Phi \mathbf{x}_1 + \mathbf{W} \rangle = \Phi \langle \mathbf{x}_1 \rangle + \langle \mathbf{W} \rangle = \Phi \boldsymbol{\mu}_1$$

$$\Delta \mathbf{x}_2 = \Phi \Delta \mathbf{x}_1 + \Delta \mathbf{W}$$

$$\begin{aligned} \boldsymbol{\Sigma}_2 &= \langle \Delta \mathbf{x}_2 \Delta \mathbf{x}_2^T \rangle = \langle (\Phi \Delta \mathbf{x}_1 + \Delta \mathbf{W})(\Phi \Delta \mathbf{x}_1 + \Delta \mathbf{W})^T \rangle \\ &= \langle \Phi \Delta \mathbf{x}_1 \Delta \mathbf{x}_1^T \Phi^T + \Phi \Delta \mathbf{x}_1 \Delta \mathbf{W}^T + \Delta \mathbf{W} \Delta \mathbf{x}_1^T \Phi^T + \Delta \mathbf{W} \Delta \mathbf{W}^T \rangle \\ &= \Phi \langle \Delta \mathbf{x}_1 \Delta \mathbf{x}_1^T \rangle \Phi^T + \Phi \langle \Delta \mathbf{x}_1 \Delta \mathbf{W}^T \rangle + \langle \Delta \mathbf{W} \Delta \mathbf{x}_1^T \rangle \Phi^T + \langle \Delta \mathbf{W} \Delta \mathbf{W}^T \rangle \\ &= \Phi \boldsymbol{\Sigma}_1 \Phi^T + \mathbf{Q} \end{aligned}$$

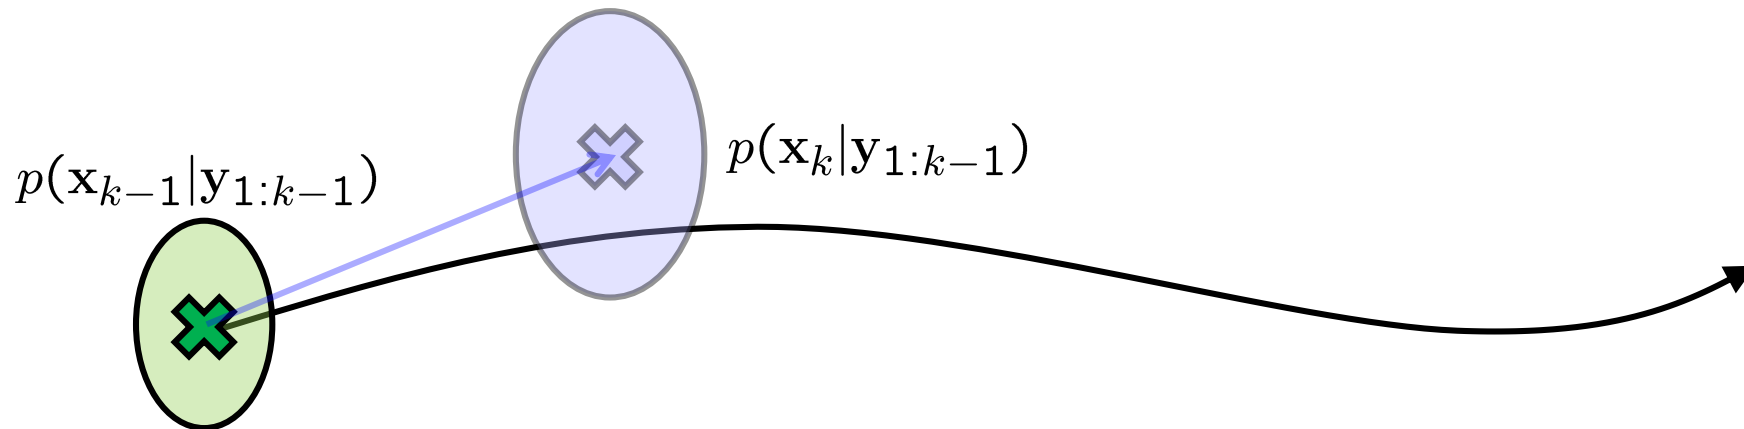
Solving the prediction

- To summarize:

$$p(\mathbf{x}_1) = N(\mu_1, \Sigma_1)$$

$$\mathbf{x}_2 = \Phi \mathbf{x}_1 + \mathbf{W} \quad , \quad \mathbf{W} \sim N(\boldsymbol{\mu} = 0, \mathbf{Q}) \quad , \quad p(\mathbf{x}_2 | \mathbf{x}_1) \sim N(\mathbf{x}_2; \Phi \mathbf{x}_1, \mathbf{Q})$$

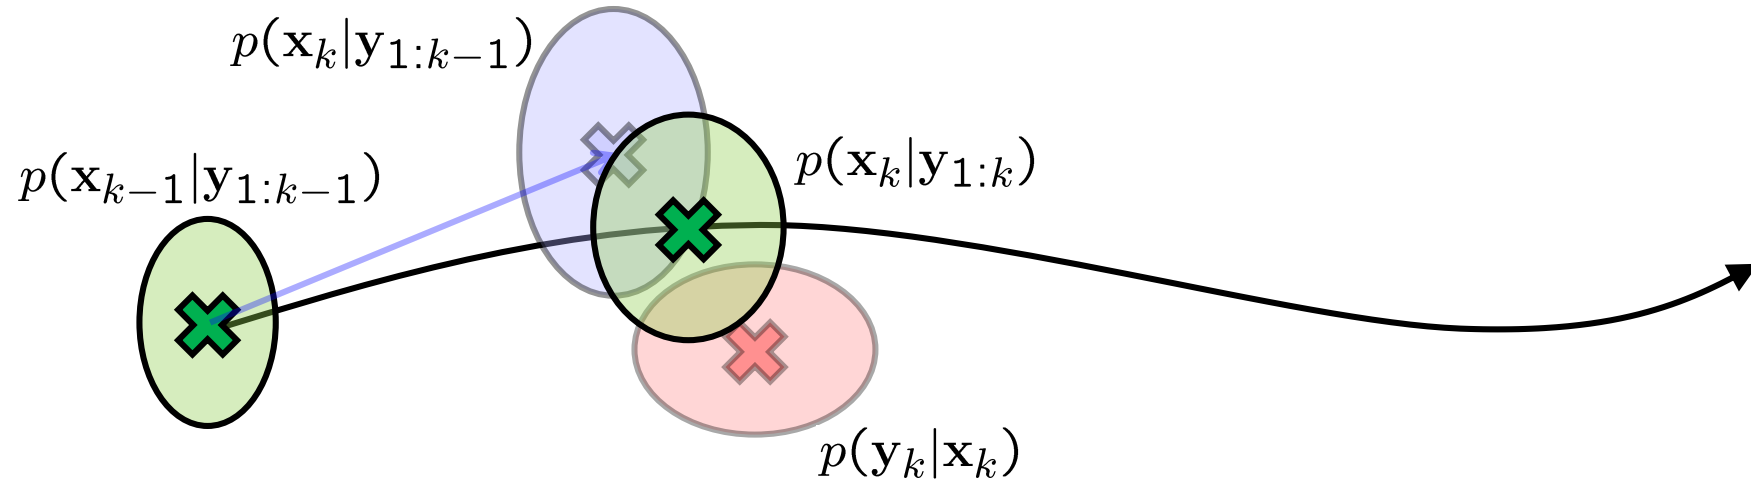
$$p(\mathbf{x}_2) = \int p(\mathbf{x}_2 | \mathbf{x}_1) p(\mathbf{x}_1) d\mathbf{x}_1 = N(\mathbf{x}_2; \Phi \boldsymbol{\mu}_1, \Phi \Sigma_1 \Phi^T + \mathbf{Q})$$



Solving for the posterior update

- We are ultimately after

$$p(\mathbf{x}_K | y_{1:K}) \propto p(\mathbf{y}_K | \mathbf{x}_K) p(\mathbf{x}_K | y_{1:K-1})$$



- We make use of the fact that the product of two Gaussians is a Gaussian as well.
- Note that we are considering the following problem:

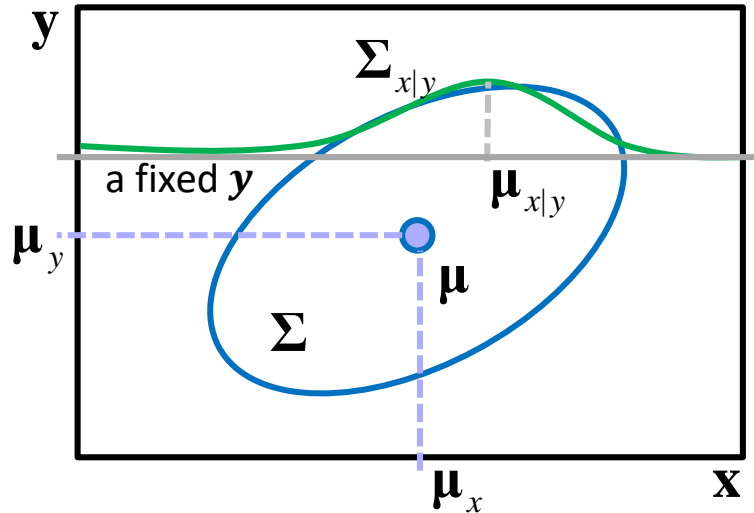
$$p(\mathbf{x} | \mathbf{y}) \propto p(\mathbf{y} | \mathbf{x}) p(\mathbf{x})$$

Solving for the posterior update

- Will take the following shortcut:

$$p(\mathbf{x} | \mathbf{y}) \propto p(\mathbf{y} | \mathbf{x}) p(\mathbf{x})$$

- Compute the joint pdf $p(\mathbf{x}, \mathbf{y})$
- Condition on \mathbf{y} : $p(\mathbf{x} | \mathbf{y})$



$$p(\mathbf{x}, \mathbf{y}) = N(\mu, \Sigma)$$

$$\mu = \begin{bmatrix} \mu_x \\ \mu_y \end{bmatrix}, \Sigma = \begin{bmatrix} \Sigma_{xx} & \Sigma_{xy} \\ \Sigma_{yx} & \Sigma_{yy} \end{bmatrix}, \Sigma_{xy} = \Sigma_{yx}^T$$

$$p(\mathbf{x} | \mathbf{y}) = N(\mathbf{x}; \mu_{x|y}, \Sigma_{x|y})$$

- Established result states that:

$$\mu_{x|y} = \mu_x + \Sigma_{xy} \Sigma_{yy}^{-1} (\mathbf{y} - \mu_y)$$

$$\Sigma_{x|y} = \Sigma_{xx} - \Sigma_{xy} \Sigma_{yy}^{-1} \Sigma_{yx}$$

Solving for the posterior update

- Recall: $p(\mathbf{x} | \mathbf{y}) \propto p(\mathbf{y} | \mathbf{x}) p(\mathbf{x})$

$$p(\mathbf{x}, \mathbf{y}) = N(\boldsymbol{\mu}, \boldsymbol{\Sigma})$$

$$\boldsymbol{\mu} = \begin{bmatrix} \boldsymbol{\mu}_x \\ \boldsymbol{\mu}_y \end{bmatrix}, \boldsymbol{\Sigma} = \begin{bmatrix} \boldsymbol{\Sigma}_{xx} & \boldsymbol{\Sigma}_{xy} \\ \boldsymbol{\Sigma}_{xy} & \boldsymbol{\Sigma}_{yy} \end{bmatrix}$$

$\mathbf{x} \sim N(\boldsymbol{\mu}_x, \boldsymbol{\Sigma}_{xx})$... the prior on \mathbf{x}

$\mathbf{y} = \mathbf{H}\mathbf{x} + \mathbf{V}$, $\mathbf{V} \sim N(\boldsymbol{\mu} = 0, \mathbf{R})$... the observation model

- From the results of conditioning we have:

$$p(\mathbf{x} | \mathbf{y}) = N(\mathbf{x}; \boldsymbol{\mu}_{x|y}, \boldsymbol{\Sigma}_{x|y})$$

$$\left. \begin{aligned} \boldsymbol{\mu}_{x|y} &= \boldsymbol{\mu}_x + \boldsymbol{\Sigma}_{xy} \boldsymbol{\Sigma}_{yy}^{-1} (\mathbf{y} - \boldsymbol{\mu}_y) \\ \boldsymbol{\Sigma}_{x|y} &= \boldsymbol{\Sigma}_{xx} - \boldsymbol{\Sigma}_{xy} \boldsymbol{\Sigma}_{yy}^{-1} \boldsymbol{\Sigma}_{xy}^T \end{aligned} \right\} \Rightarrow$$

The following values are required:

$$\boldsymbol{\mu}_y = ? , \boldsymbol{\Sigma}_{xy} = ? , \boldsymbol{\Sigma}_{yy} = ?$$

Solving for the posterior update

- We are after: $\boldsymbol{\mu}_y = ?$, $\boldsymbol{\Sigma}_{xy} = ?$, $\boldsymbol{\Sigma}_{yy} = ?$

$$\boldsymbol{\mu}_y = \langle \mathbf{y} \rangle = \langle \mathbf{H}\mathbf{x} + \mathbf{V} \rangle = \mathbf{H}\boldsymbol{\mu}_x$$

$$\Delta \mathbf{y} = \mathbf{H}\Delta \mathbf{x} + \Delta \mathbf{V}$$

$$\begin{aligned}\boldsymbol{\Sigma}_{yy} &= \langle \Delta \mathbf{y} \Delta \mathbf{y}^T \rangle = \langle (\mathbf{H}\Delta \mathbf{x} + \Delta \mathbf{V})(\mathbf{H}\Delta \mathbf{x} + \Delta \mathbf{V})^T \rangle \\ &= \mathbf{H}\langle \Delta \mathbf{x} \Delta \mathbf{x}^T \rangle \mathbf{H}^T + 0 + 0 + \langle \Delta \mathbf{V} \Delta \mathbf{V}^T \rangle = \mathbf{H}\boldsymbol{\Sigma}_{xx} \mathbf{H}^T + \mathbf{R}\end{aligned}$$

$$\boldsymbol{\Sigma}_{xy} = \langle \Delta \mathbf{x} \Delta \mathbf{y}^T \rangle = \langle \Delta \mathbf{x} (\mathbf{H}\Delta \mathbf{x} + \Delta \mathbf{V})^T \rangle = \langle \Delta \mathbf{x} \Delta \mathbf{x}^T \rangle \mathbf{H}^T + \langle \Delta \mathbf{x} \Delta \mathbf{V}^T \rangle = \boldsymbol{\Sigma}_{xx} \mathbf{H}^T$$

- And finally: $p(\mathbf{x} | \mathbf{y}) = \mathcal{N}(\mathbf{x}; \boldsymbol{\mu}_{x|y}, \boldsymbol{\Sigma}_{x|y})$

$$\boldsymbol{\mu}_{x|y} = \boldsymbol{\mu}_x + \boldsymbol{\Sigma}_{xx} \mathbf{H}^T (\mathbf{H}\boldsymbol{\Sigma}_{xx} \mathbf{H}^T + \mathbf{R})^{-1} (\mathbf{y} - \boldsymbol{\mu}_y)$$

$$\boldsymbol{\Sigma}_{x|y} = \boldsymbol{\Sigma}_{xx} - \boldsymbol{\Sigma}_{xx} \mathbf{H}^T (\mathbf{H}\boldsymbol{\Sigma}_{xx} \mathbf{H}^T + \mathbf{R})^{-1} \mathbf{H}\boldsymbol{\Sigma}_{xx}$$

$$\begin{aligned}\mathbf{y} &= \mathbf{H}\mathbf{x} + \mathbf{V} \\ \mathbf{V} &\sim \mathcal{N}(\boldsymbol{\mu} = 0, \mathbf{R}) \\ \mathbf{x} &\sim \mathcal{N}(\boldsymbol{\mu}_x, \boldsymbol{\Sigma}_{xx})\end{aligned}$$

Putting it all together: the Kalman filter

External input (if available)

- Dynamic model: $\mathbf{x}_k = \Phi \mathbf{x}_{k-1} + \Gamma \mathbf{u}_k + \mathbf{W}_k$ $\mathbf{W}_k \sim \mathcal{N}(\mu = 0, \mathbf{Q})$
- Observation model: $\mathbf{y}_k = \mathbf{H} \mathbf{x}_k + \mathbf{V}_k$ $\mathbf{V}_k \sim \mathcal{N}(\mu = 0, \mathbf{R})$
- Initial posterior: $p(\mathbf{x}_{k-1} | \mathbf{y}_{1:k-1}) = \mathcal{N}(\mathbf{x}_{k-1}; \mu = \hat{\mathbf{x}}_{k-1}, \Sigma = \mathbf{P}_{k-1})$

1. Prediction: $p(\mathbf{x}_k | \mathbf{y}_{1:k-1}) = \mathcal{N}(\mathbf{x}_k; \mu = \tilde{\mathbf{x}}_k, \Sigma = \tilde{\mathbf{P}}_k)$

$$\tilde{\mathbf{x}}_k = \Phi \hat{\mathbf{x}}_{k-1} + \Gamma \mathbf{u}_k$$

$$\tilde{\mathbf{P}}_k = \Phi \mathbf{P}_{k-1} \Phi^T + \mathbf{Q}$$

2. Update by measurement \mathbf{y}_k :

$$p(\mathbf{x}_k | \mathbf{y}_{1:k}) = \mathcal{N}(\mathbf{x}_k; \mu = \hat{\mathbf{x}}_k, \Sigma = \hat{\mathbf{P}}_k)$$

$$\hat{\mathbf{x}}_k = \tilde{\mathbf{x}}_k + \mathbf{K}(\mathbf{y}_k - \mathbf{H} \tilde{\mathbf{x}}_k)$$

$$\mathbf{P}_k = (\mathbf{I} - \mathbf{K} \mathbf{H}) \tilde{\mathbf{P}}_k$$

This is called the “Kalman gain”:

$$\mathbf{K} = \tilde{\mathbf{P}}_k \mathbf{H}^T (\mathbf{H} \tilde{\mathbf{P}}_k \mathbf{H}^T + \mathbf{R})^{-1}$$

Making more sense of these equations...


- Prediction: $\tilde{\mathbf{x}}_k = \Phi \hat{\mathbf{x}}_{k-1} + \Gamma \mathbf{u}_k$
 $\tilde{\mathbf{P}}_k = \Phi \mathbf{P}_{k-1} \Phi^T + \mathbf{Q}$
- Update: $\mathbf{K} = \tilde{\mathbf{P}}_k \mathbf{H}^T (\mathbf{H} \tilde{\mathbf{P}}_k \mathbf{H}^T + \mathbf{R})^{-1}$
 $\hat{\mathbf{x}}_k = \tilde{\mathbf{x}}_k + \mathbf{K}(\mathbf{y}_k - \mathbf{H} \tilde{\mathbf{x}}_k)$
 $\mathbf{P}_k = (\mathbf{I} - \mathbf{K} \mathbf{H}) \tilde{\mathbf{P}}_k$

Kalman Filter recursion

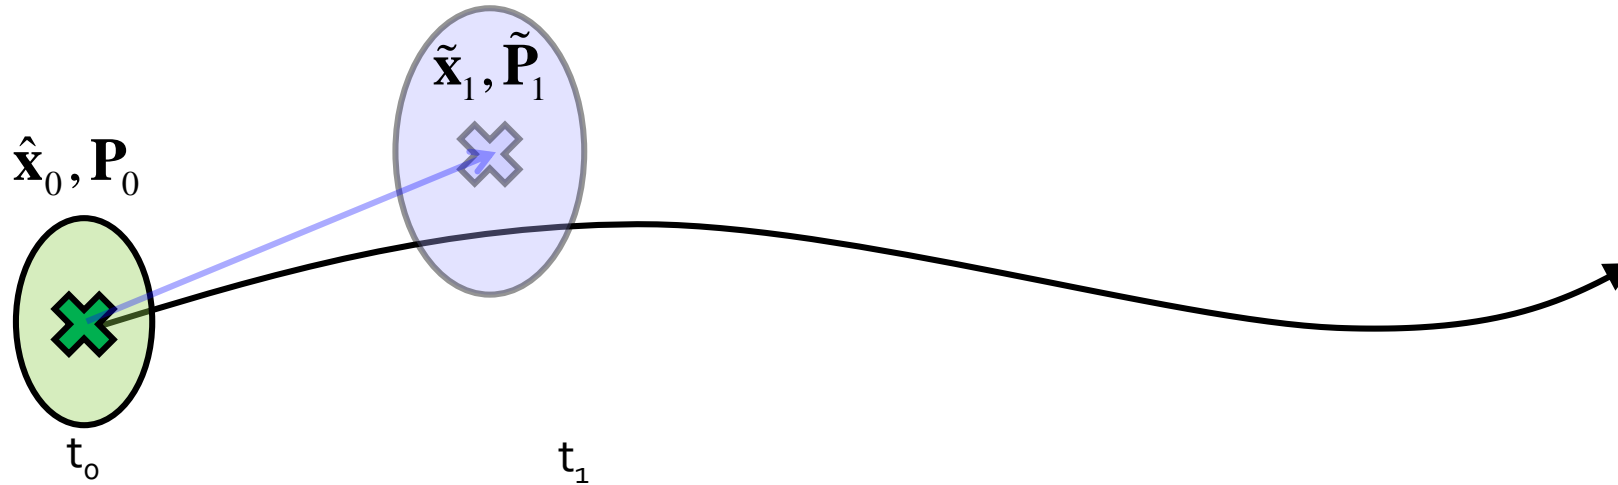


- Initialize

$$\hat{\mathbf{x}}_0 = \begin{bmatrix} x_0 \\ y_0 \\ \dot{x}_0 \\ \dot{y}_0 \end{bmatrix}$$
$$\mathbf{P}_0 = \begin{bmatrix} L & 0 & 0 & 0 \\ 0 & L & 0 & 0 \\ 0 & 0 & L & 0 \\ 0 & 0 & 0 & L \end{bmatrix}$$




Kalman Filter recursion




- Predict:

$$\begin{aligned}\tilde{\mathbf{x}}_1 &= \mathbf{\Phi} \hat{\mathbf{x}}_0 + \Gamma \mathbf{u}_1 \\ \tilde{\mathbf{P}}_1 &= \mathbf{\Phi} \mathbf{P}_0 \mathbf{\Phi}^T + \mathbf{Q}\end{aligned}$$

External input – for most your applications it will be unavailable, i.e., $\mathbf{u}_k = 0$



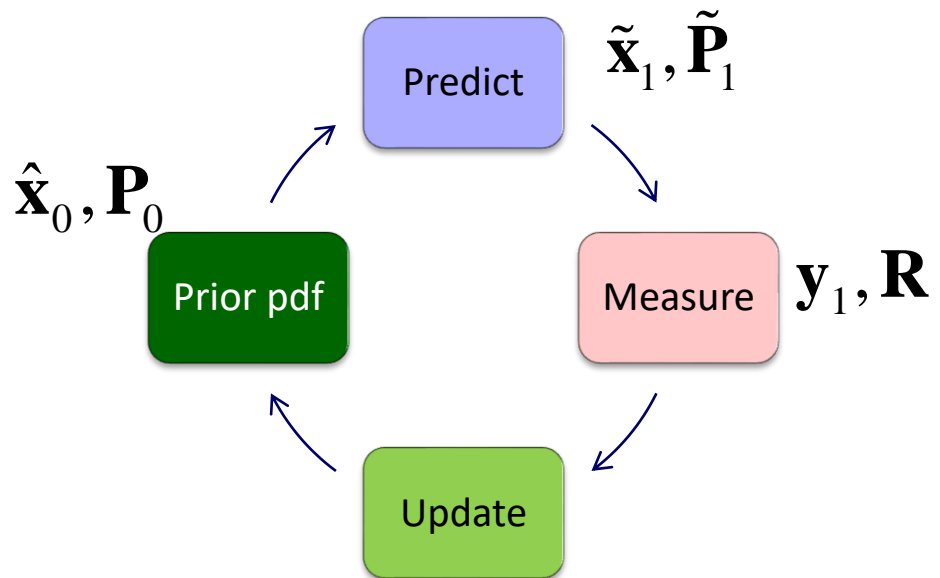
Kalman Filter recursion




- Receive a noisy measurement
 \mathbf{y}_1, \mathbf{R}

Compute the “Kalman gain”:

$$\mathbf{K} = \tilde{\mathbf{P}}_1 \mathbf{H}^T (\mathbf{H} \tilde{\mathbf{P}}_1 \mathbf{H}^T + \mathbf{R})^{-1}$$



Kalman Filter recursion

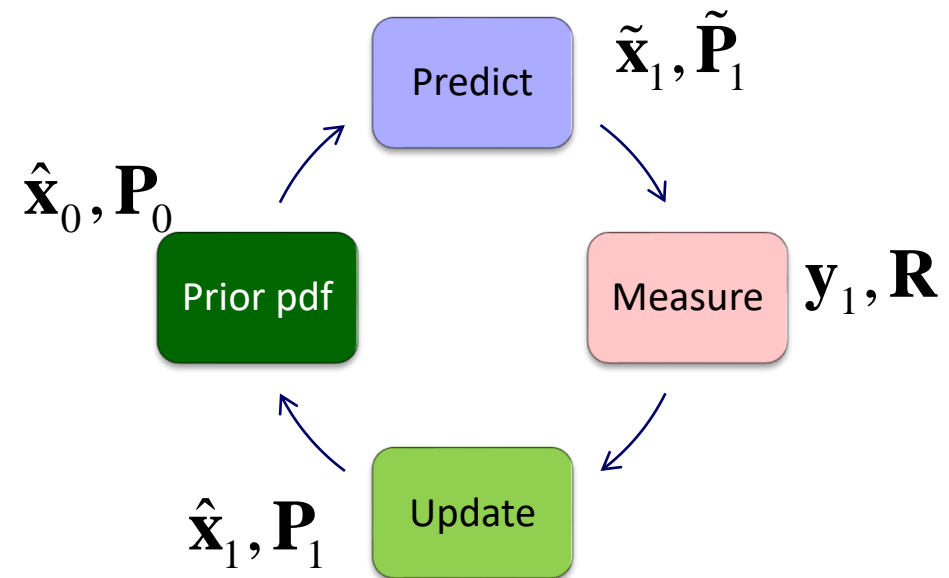


- Update the posterior

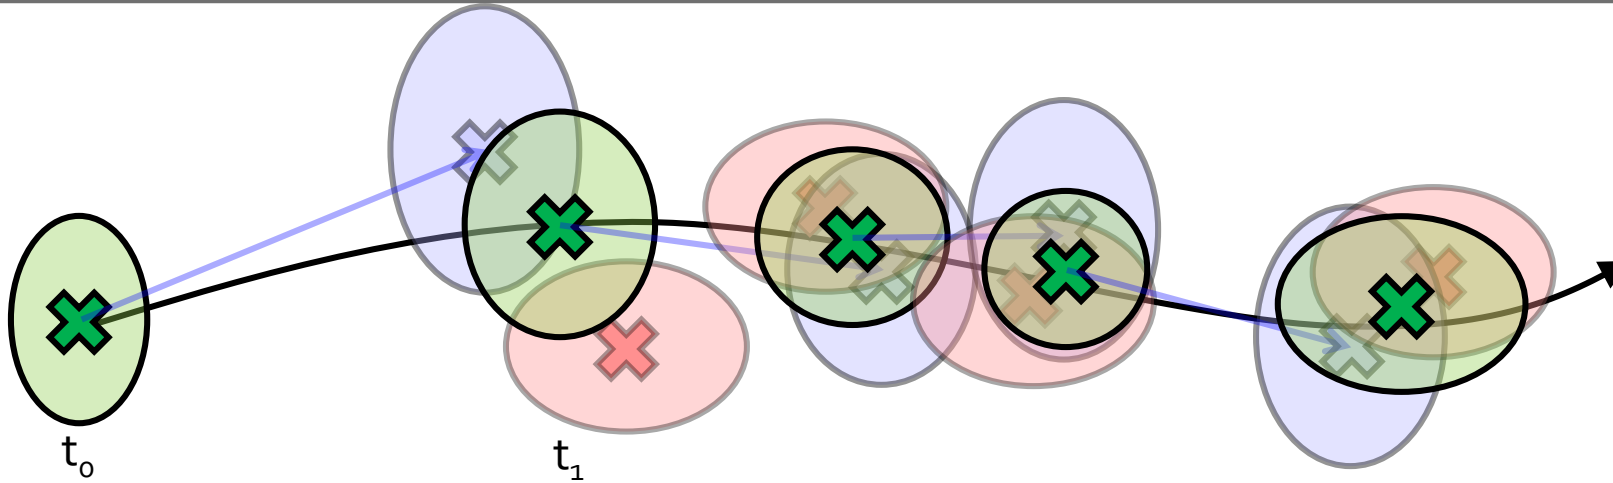
$$\mathbf{K} = \tilde{\mathbf{P}}_1 \mathbf{H}^T (\mathbf{H} \tilde{\mathbf{P}}_1 \mathbf{H}^T + \mathbf{R})^{-1}$$

$$\hat{\mathbf{x}}_1 = \tilde{\mathbf{x}}_1 + \mathbf{K}(\mathbf{y}_1 - \mathbf{H}\tilde{\mathbf{x}}_1)$$

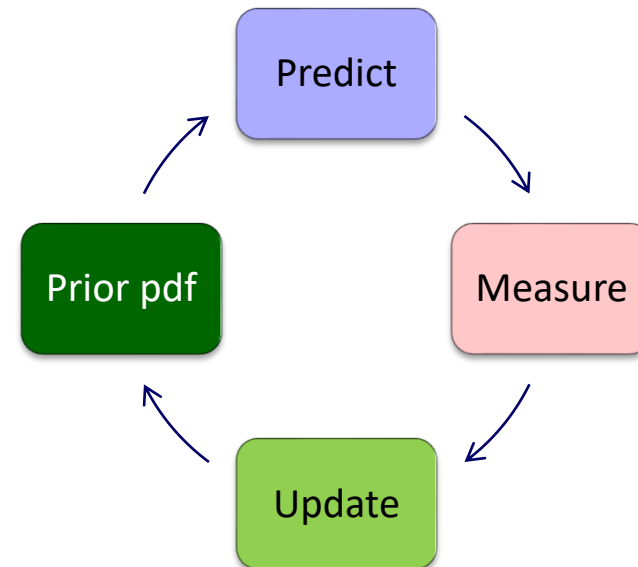
$$\mathbf{P}_1 = (\mathbf{I} - \mathbf{K}\mathbf{H})\tilde{\mathbf{P}}_1$$



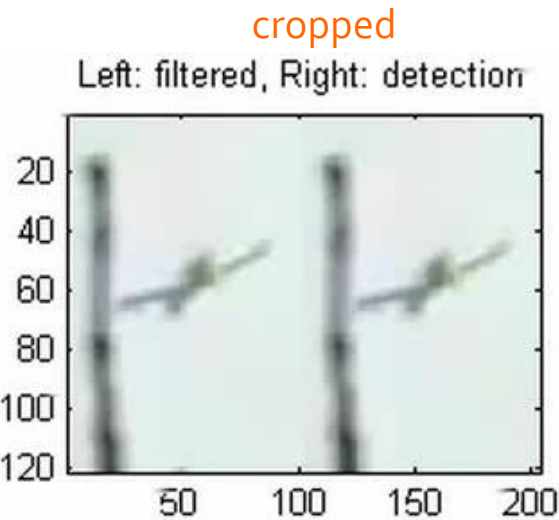
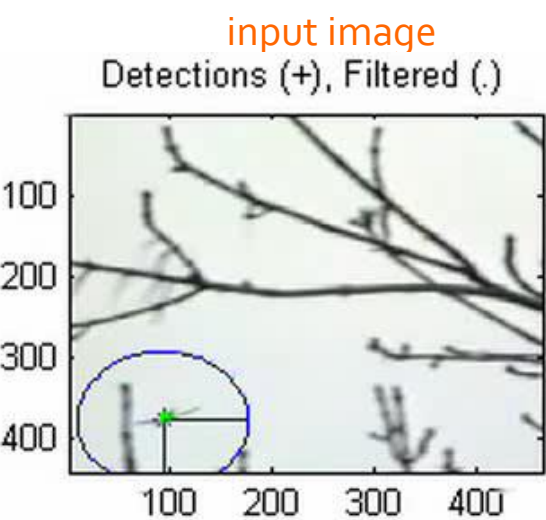
Kalman Filter recursion



1. Prediction from the motion model
2. Receive a noisy measurement
3. Update the posterior



Kalman filter in action




Detections


Estimates








predict




measure



update



Kalman filter: Dynamics

Without velocity (RW): $\mathbf{x}_t = \begin{pmatrix} x \\ y \end{pmatrix}$




Target
is moving

With velocity (NCV): $\mathbf{x}_t = \begin{pmatrix} x \\ y \\ \dot{x} \\ \dot{y} \end{pmatrix}$



Target is **not**
moving



Another example



- detection
- prediction
- update

$$\tilde{\mathbf{x}}_k = \Phi \hat{\mathbf{x}}_{k-1}$$

$$\tilde{\mathbf{P}}_k = \Phi \mathbf{P}_{k-1} \Phi^T + \mathbf{Q}$$

$$\mathbf{K} = \tilde{\mathbf{P}}_k \mathbf{H}^T (\mathbf{H} \tilde{\mathbf{P}}_k \mathbf{H}^T + \mathbf{R})^{-1}$$

$$\hat{\mathbf{x}}_k = \tilde{\mathbf{x}}_k + \mathbf{K}(\mathbf{y}_k - \mathbf{H} \tilde{\mathbf{x}}_k)$$

$$\mathbf{P}_k = (\mathbf{I} - \mathbf{K} \mathbf{H}) \tilde{\mathbf{P}}_k$$

Dynamic model takes over
when target not detected!

Multiple measurements (noise)?


- Assume a NCV model

$$\mathbf{x}_k = \Phi \mathbf{x}_{k-1} + \mathbf{W}_k$$


$$\mathbf{W}_k \sim \mathcal{N}(\mu = 0, \mathbf{Q})$$

- How would you resolve noisy measurements?



$$\mathbf{y}_k = \mathbf{H} \mathbf{x}_k + \mathbf{V}_k$$



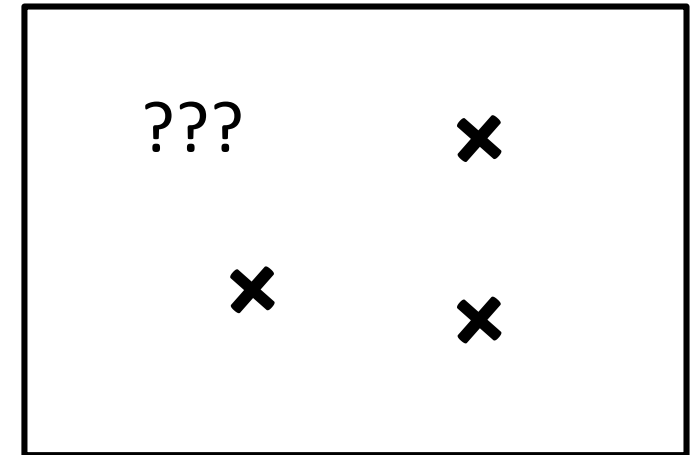
measure



measure



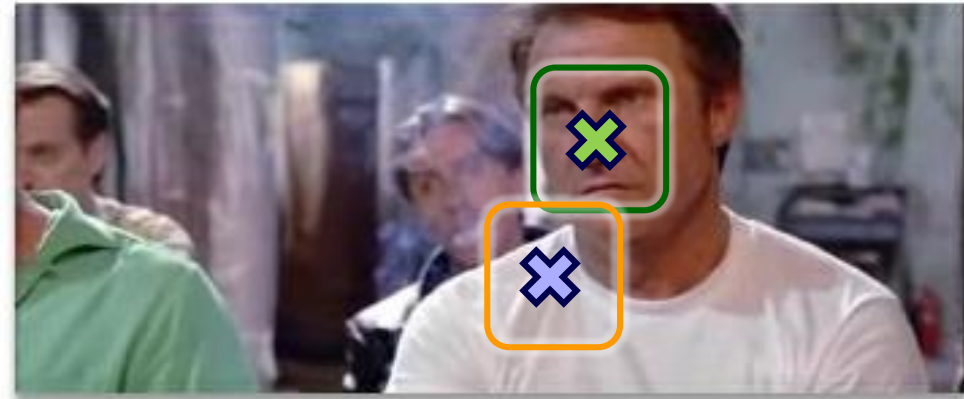
measure



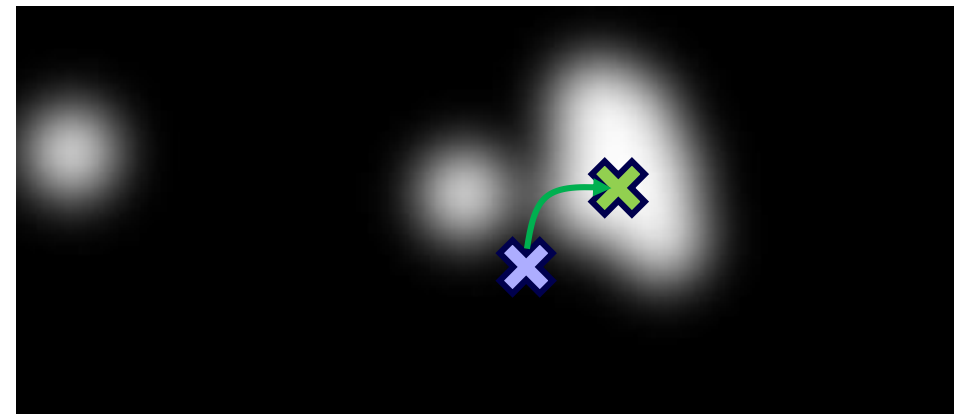
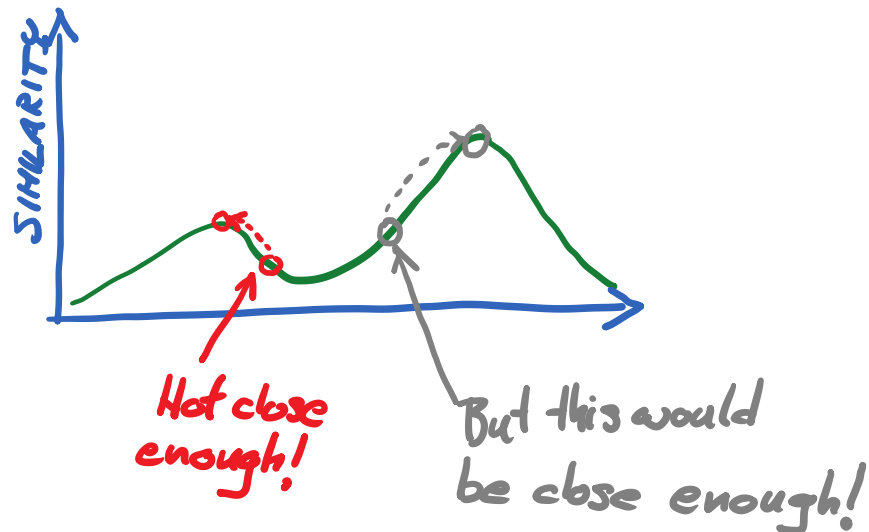
Combine Kalman with local optimization?

- Recall that Mean Shift **converges well** if **initialized close to solution**

Apply Kalman for prediction to better estimate the starting point for MS!



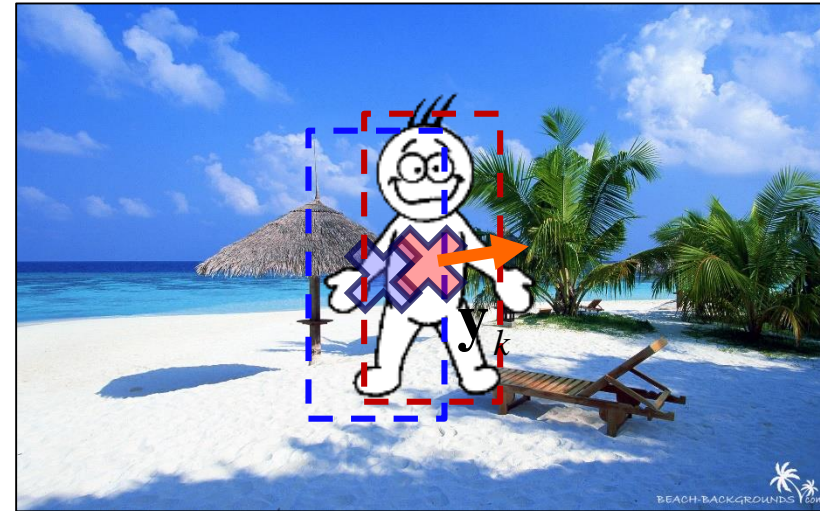
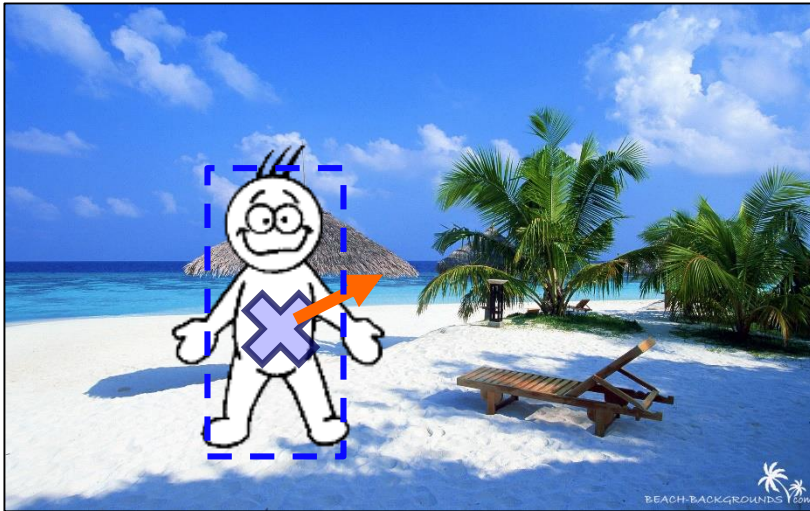
input image



similarity/probability

Improve Mean Shift with the Kalman filter

- Example:
 - Predict initial position from Kalman filter
 - Run Mean Shift to find local optimum – This is measurement \mathbf{y}_k
 - Update Kalman filter by the measurement \mathbf{y}_k
 - Prediction improved for the next time-step



“Possible to improve any local optimization by dynamics in this way”

Setting parameters?

- Usually set only the covariances: Q and R

$$\mathbf{x}_k = \mathbf{\Phi} \mathbf{x}_{k-1} + \mathbf{W}_k \qquad \mathbf{W}_k \sim \mathcal{N}(\mu = 0, \mathbf{Q})$$

$$\mathbf{y}_k = \mathbf{H} \mathbf{x}_k + \mathbf{V}_k \qquad \mathbf{V}_k \sim \mathcal{N}(\mu = 0, \mathbf{R})$$

- By trial and error
- By Expectation Maximization ([see this book](#)¹) on a reference trajectory
- **By rules of thumb** – a good starting point

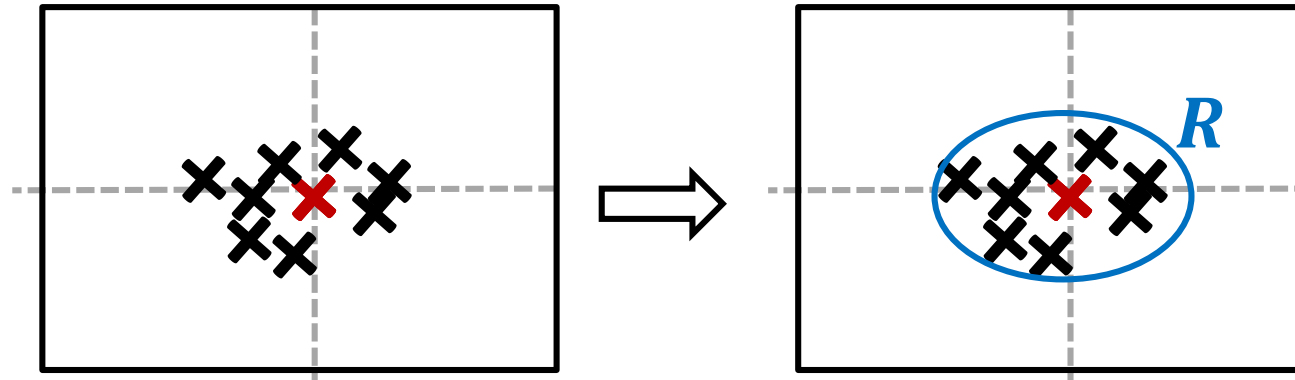
¹Christopher Bishop, Pattern Recognition and Machine Learning, [free online](#)

Setting the measurement cov. R

- Perform detection with your detector on many examples.
- Manually annotate TRUE positions.



- Calculate *changes* from the true position
- Calculate covariance R



Rule of thumb for dynamics noise Q

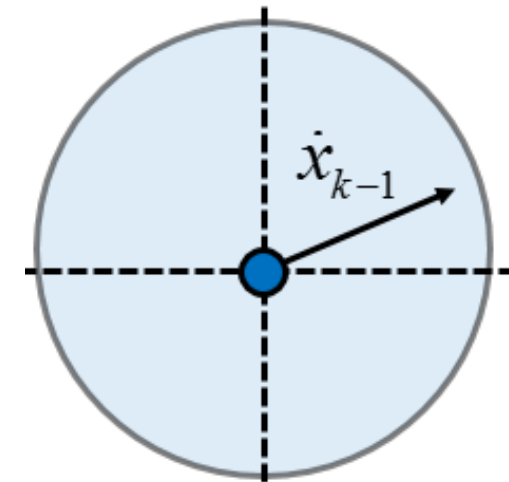
- Example, NCV: $\mathbf{x}_k = \Phi \mathbf{x}_{k-1} + \mathbf{W}$, $\Delta t = 1$

- $\mathbf{W} \sim N(0, \mathbf{Q})$

$$\mathbf{Q} = q_c \begin{bmatrix} \frac{1}{3} & \frac{1}{2} \\ \frac{1}{2} & 1 \end{bmatrix} \quad \Phi = \begin{bmatrix} 1 & 1 \\ 0 & 1 \end{bmatrix}$$

- Assume we know the **expected squared distance** σ_m^2 a target can travel within Δt .
- Target is at origin and *starts moving*, i.e., velocity sampled from noise only:

$$\mathbf{x}_{k-1} = \begin{bmatrix} 0 \\ \dot{x}_{k-1} \end{bmatrix} \quad \dot{x}_{k-1} \sim N(0, q_c)$$



A general approach proposed in: M. Kristan et al. "A Two-Stage Dynamic Model for Visual Tracking". IEEE SMCB, 2010.

Rule of thumb for dynamics noise \mathbf{Q}

- Model specification with NCV: $\mathbf{x}_k = \Phi \mathbf{x}_{k-1} + \mathbf{W}$, $\mathbf{W} \sim N(0, \mathbf{Q})$

$$\Phi = \begin{bmatrix} 1 & 1 \\ 0 & 1 \end{bmatrix} \quad \mathbf{Q} = q_c \begin{bmatrix} \frac{1}{3} & \frac{1}{2} \\ \frac{1}{2} & 1 \end{bmatrix} \quad \mathbf{x}_{k-1} = \begin{bmatrix} 0 \\ \dot{x}_{k-1} \end{bmatrix} \quad \dot{x}_{k-1} \sim N(0, q_c)$$

- The covariance of \mathbf{x}_k (expected sq. change of the state \mathbf{x}_k):

$$\begin{aligned} P &= \begin{bmatrix} \sigma_m^2 & p_{12} \\ p_{12} & p_{22} \end{bmatrix} = \langle (\mathbf{x}_k - 0)(\mathbf{x}_k - 0)^T \rangle = \langle \mathbf{x}_k \mathbf{x}_k^T \rangle = \Phi \langle \mathbf{x}_{k-1} \mathbf{x}_{k-1}^T \rangle \Phi^T + \mathbf{Q} \\ &= \Phi \left\langle \begin{bmatrix} 0 & 0\dot{x}_{k-1} \\ 0\dot{x}_{k-1} & \dot{x}_{k-1}\dot{x}_{k-1} \end{bmatrix} \right\rangle \Phi^T + \mathbf{Q} = \Phi \begin{bmatrix} 0 & 0 \\ 0 & \langle \dot{x}_{k-1} \dot{x}_{k-1} \rangle \end{bmatrix} \Phi^T + \mathbf{Q} \\ &\quad \langle \dot{x}_{k-1} \dot{x}_{k-1} \rangle = q_c \\ &= \begin{bmatrix} q_c & q_c \\ q_c & q_c \end{bmatrix} + \mathbf{Q} = q_c \begin{bmatrix} 1\frac{1}{3} & 1\frac{1}{2} \\ 1\frac{1}{2} & 2 \end{bmatrix} = \begin{bmatrix} \sigma_m^2 & p_{12} \\ p_{12} & p_{22} \end{bmatrix} \longrightarrow q_c = \frac{3}{4} \sigma_m^2 \end{aligned}$$

Rule of thumb for dynamics noise Q

- Example of applying the rule of thumb
- Say: The expected change of target position is 10 pixels.
- Therefore the squared change is approximately:
 $\sigma_m^2 = 10^2$
- Then, applying the rule of thumb, the spectral density is $q_c = \frac{3}{4}10^2$
- So the dynamic model covariance is:

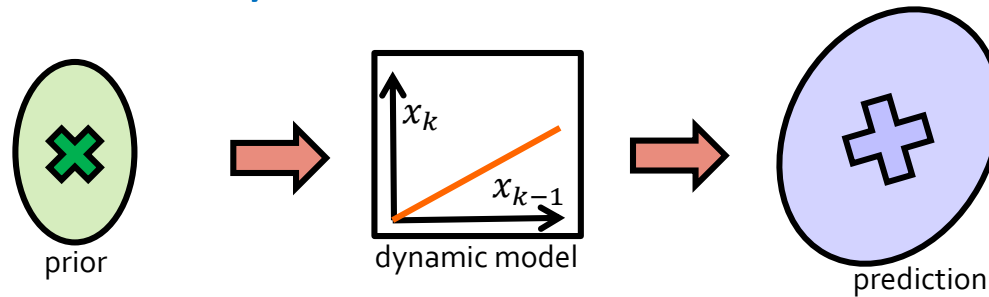
$$Q = q_c \begin{bmatrix} \frac{1}{3} & \frac{1}{2} \\ \frac{1}{2} & 1 \end{bmatrix} = \frac{3}{4}10^2 \begin{bmatrix} \frac{1}{3} & \frac{1}{2} \\ \frac{1}{2} & 1 \end{bmatrix}$$

Consider this an
upper bound!

Beyond the basic Kalman

- Assumes **linear dynamics** with Gaussian noise
- + Simplifies the update equations
- - Cannot account for nonlinear dynamics

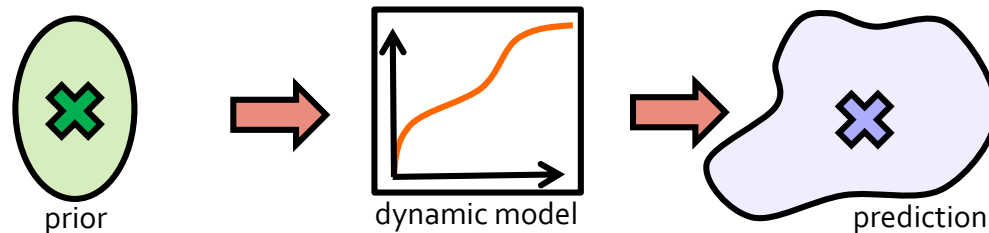
A linear dynamic model:



$$\int p(\mathbf{x}_k | \mathbf{x}_{k-1}) p(\mathbf{x}_{k-1} | \mathbf{y}_{1:k-1}) d\mathbf{x}_{k-1}$$

Integral **easy to solve**

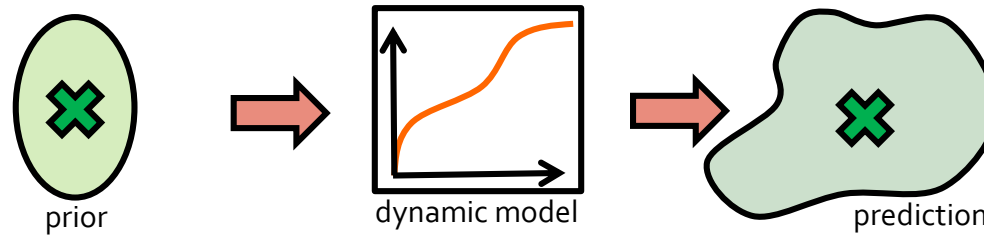
A non-linear dynamic model:



Integral **not easy to solve**
The result is **not a Gaussian**

Handling nonlinear dynamics


- Problem: prediction is no longer a Gaussian!



- Extended Kalman filter:

$$\int p(\mathbf{x}_k | \mathbf{x}_{k-1}) p(\mathbf{x}_{k-1} | \mathbf{y}_{1:k-1}) d\mathbf{x}_{k-1}$$

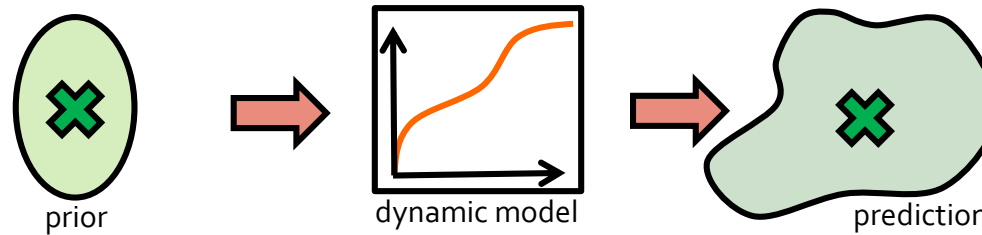
- Linearize the dynamic model at \mathbf{x}_{k-1}



- *Usually does not properly propagate the covariance*

Handling nonlinear dynamics

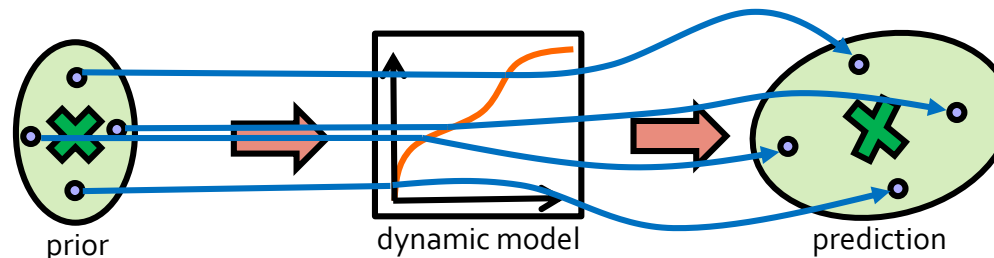
- Prediction no longer Gaussian!



- Unscented Kalman filter:

$$\int p(\mathbf{x}_k | \mathbf{x}_{k-1}) p(\mathbf{x}_{k-1} | \mathbf{y}_{1:k-1}) d\mathbf{x}_{k-1}$$

- Numerically solve the integral using the Unscented transform



A nice summary of the main equations [available here](#).

Wan, E.A.; Van Der Merwe, R., [The unscented Kalman filter for nonlinear estimation](#), Proceedings of the IEEE ASSPCC 2000

See J.D. Prince "Computer vision: models, learning and inference", Section 19.4

References

- Text book Kalman:
 - S. J.D. Prince, “Computer vision: models, learning and inference”, Section 19.2
- For additional info on probability see:
 - S. J.D. Prince, “Computer vision: models, learning and inference”, Chapter 1

Acknowledgement

- Some images and parts of slides have been taken from the following talks:
 - Kevin Smith's "SELECTED TOPICS IN COMPUTER VISION – 2D tracking"